



JOHNS HOPKINS
UNIVERSITY

ARCH Portal

Development Guide

Johns Hopkins University

IT@JHU Research Computing

Advanced Research Computing at Hopkins (ARCH)



Table of Contents

ARCH Portal (CLOACK Stack)	3
High-level design (draft) - Authentication + Scheduler stack	3
Cross-app SSO (ColdFront <-> Helpdesk)	5
start.sh -- Stack Lifecycle Management	6
Bare-metal install -- cloack-deploy RPM	15
Docker Images	16
Partitions & Allocation Tiers	17
Billing & Invoicing System	18
Populating Synthetic Data for Testing	26
10. Slurm CLI Utilities	32
Testing	36
User Guide	37
Multi-Cluster Architecture	38
Monitoring	41
Features	43
Branch Overview	43
Documentation Index	43
Feature Comparison: main vs prod vs dev deployment	43
What prod has over main (69 commits)	50
What dev has over prod (259 commits)	50
Billing-Free Resources	51
Slurm Account Naming Convention	51
UID/GID Ranges	51
Staff Help & Reference Guide	52
{{ section.title }}	52
ColdFront	53
Build	53
Startup (entrypoint.sh)	53
Directory Structure	53
Key Config	53
Environment Variables	54
Port	54
Development	54
Admin Features	54
LDAP Sync Monitoring	55
Custom Configurations	57
Structure	57
How These Customizations Work	57
Key Customizations	58
Deployment Notes	60



ARCH Portal (CLOACK Stack)

The ARCH Portal is built on the CLOACK Stack, a containerized identity and resource management framework for HPC clusters at Johns Hopkins University. It integrates ColdFront, OpenLDAP, Keycloak, and Slurm into a container-based platform for user provisioning, job scheduling, and billing.

Component	Role
Coldfront	HPC resource allocation portal -- orchestrates LDAP/Slurm sync, home dirs, PI scratch links
LDAP / OpenLDAP	Central directory -- users, groups, PIs across JHU and Schmidt (SSCI) domains
OpenLDAP	Single OpenLDAP instance with three naming contexts ('dc=arch', 'ou=jhu', 'ou=schmidt')
Auth (OIDC/SSSD)	OIDC via Keycloak + SSSD for PAM/NSS resolution inside containers
Cluster/Containers	Slurm HPC cluster + Docker Compose orchestration -- builds, networks, and manages all services as a unified stack
Keycloak	Identity broker -- per-realm LDAP federation, Entra ID integration, token issuance

Deployed on the ARCH cluster at Johns Hopkins University, supporting both JHU and Schmidt Sciences (SSCI) r...

Interactive Architecture Diagram -- click any component for details.

High-level design (draft) - Authentication + Scheduler stack

We are using Coldfront as Orchestrator. It making LDAP and Slurm synchronization, as well user's home directory and softlink point to PI scratch folder where they are members.

Use three separate LDAP naming contexts (suffixes) and disjoint uid/gid ranges; create cross-domain admin/group entries (DN-based) that include members from any suffix; configure Open Source Identity and Access Management "Keycloak" realms so each realm points to the correct subtree(s) or adds multiple LDAP providers.

Three suffixes on one OpenLDAP instance:

- Domain A (arch.cluster): ARCH Admins suffix **dc=arch,dc=cluster** UID/GID range starting 1000
- Domain B (jhu.arch.cluster): suffix **ou=jhu,dc=arch,dc=cluster** UID/GID range starting 10000
- Domain C (schmidt.arch.cluster): suffix **ou=schmidt,dc=arch,dc=cluster** UID/GID range starting 30000

Keep UID/GID ranges non-overlapping.

ColdFront as Directory Ground Truth: UID and GID allocation is centralized in the ColdFront database via **Affiliation.next_uid / next_gid** counters (atomic, race-condition free). **UserProfile** stores **ldap_uid**, **ldap_gid**, and **ldap_groups** (readonly, cached from LDAP sync). The Django Admin "Change User" and "UserProfile" pages show a "Directory (LDAP)" fieldset with this information. On ColdFront startup, **sync_ldap** automatically pushes users/groups to LDAP. The Import Users page supports two modes: New Users (CSV, allocate new UIDs) and Migration (LDIF slapcat dump, preserve existing UIDs/GIDs/passwords). New Users mode also has an optional Course code field for disposable class accounts set e.g. en540 and each CSV **uid** becomes **en540-** with a per-course synthetic email (so a student who already has an account gets a separate, UID-distinct identity); the whole cohort is removed at term end via the "Delete Course Accounts" card or **coldfront delete course users course en540 yes** (drops Slurm associations, the user's home dir, LDAP entry, and ColdFront user shared per-PI scratch is never touched). To tear down specific accounts by



name (not a course prefix), use the generic sibling **coldfront delete_users -u dry-run then yes** same four-step teardown, but it refuses system accounts and blocks deleting a PI (which would cascade-delete their projects/allocations) unless **force-pi** is passed. The page also has an Export Users card that streams the directory as CSV (**full_name,email,uid,is_pi,members,member_of**, filterable by scope and affiliation) **members** lists a PI's LDAP group members and **member_of** lists the PI(s) a user belongs to.

Minimal LDAP Bootstrap: **install.sh** creates only the bare minimum (root DN, cn=root, munge/slurm system accounts). OUs (**ou=jhu**, **ou=schmidt**, etc.) and user accounts are not hardcoded they come from ColdFront via:

- Option A: Import Users -> Migration (LDIF upload via UI) replays structural entries + users + groups + creates projects/allocations
- Option B: **./start.sh ldap-reinit** (CLI) restores full slapcat dump
- Option C: **sync_ldap** pushes users/groups from ColdFront DB to LDAP

LDAP admin -> Django privilege mapping. Cluster admin (**cn=admin**) and ColdFront portal admin (**User.is_staff / User.is_superuser**) are decoupled they're different roles.

- Cluster admin (**cn=admin**) is driven by the Django group **Cluster (LDAP admin)** (seeded by migration **0089_admin_django_group**). Staff manage membership at **/admin/auth/group/** same UX as the helpdesk groups. signals. sync helpdesk ldap groups publishes adds/removes into cn=admin,ou=Groups,ou= via m2m changed on auth.User groups: utils/ldap sync. sync cluster admin members() runs an additive reconciliation during every **run_ldap_sync** to self-heal any missed signal. **User.is_staff** no longer affects **cn=admin**.

- ColdFront portal admin is **User.is_staff=True** (sidebar Administration + Django Admin access). **seed_initial_data.sync_admin_group_to_is_staff** still promotes every LDAP **cn=admin** member to **is_staff=True** at boot (and adds them to the **Helpdesk Staff (LDAP helpdesk-staff)** Django group), so LDIF-imported admins keep portal access on first deploy. Going forward the two can be toggled independently.

- Only usernames listed in the **ARCH_SUPERUSERS** env var (comma-separated, seeded into **.env** by **start.sh init** dev default: rdesouz4,jbiondo2,mprotzm2,jasonw) additionally receive **is_superuser=True**. This gates the "Delete User" button and every other permission bypass to an explicit allow-list, so putting a PI in LDAP **cn=admin** does not silently grant full Django Admin power.

- System accounts (**SYSTEM_USERNAMES = admin, coldfront, root**) are never promoted by these helpers, and **utils/ldap_sync.scrub_system_accounts_from_privilege_groups** runs at the tail of every **run_ldap_sync()** to **MODIFY_DELETE** them from **cn=admin**, **cn=helpdesk-admin**, **cn=helpdesk-staff**. The migration LDIF can safely list **coldfront** in **cn=admin**; the scrub neutralizes it on every sync.

Put the cross-domain admin group (e.g., **PI_XX**) as a DN-based group (for example **groupOfNames** or **groupOfUniqueNames**) and include member DNs from any suffix so membership is unambiguous.

Why this is safe / best-practice

- Disjoint numeric ranges prevent collisions for NSS/PAM and NFS.
- DN-based group membership is unambiguous across suffixes and works well with Keycloak group mappers.



```
Minimal LDIF example:
dn: cn=PI_XX,ou=Groups,dc=arch,dc=cluster
objectClass: top
objectClass: groupOfNames
cn: PI_XX
member: uid=alice,ou=People,ou=jhu,dc=arch,dc=cluster
member: uid=bob,ou=People,dc=schmidt,dc=arch,dc=cluster
```

Recommendation: place that group in a shared container (e.g., ou=Groups,ou=jhu,dc=arch,dc=cluster) and conf...

Additional recommendations:

- Do NOT use **memberUid/posixGroup** for cross-suffix admin groups **memberUid** is just a username and can collide across suffixes.
- Ensure the referenced user entries actually exist (add users before groups, or import groups after users when bootstrapping) so DNS resolve correctly.
- Keycloak: configure the LDAP group mapper to search the shared group container and to honor DN-based membership (works well with **groupOfNames** / **groupOfUniqueNames**).
- Configure Keycloak per-realm to point at the appropriate suffix (or add multiple LDAP providers per realm) so each realm sees the right subtree(s).

Cross-app SSO (ColdFront <-> Helpdesk)

Both apps share the same Keycloak realms (**jhu**, **schmidt**). Silent SSO between them works when the second app's sidebar link routes to the realm the user is already authenticated in.

From -> To	Sidebar link	Behaviour
ColdFront <-> Helpdesk (both logged in)	any icon	Silent login on the correct realm
ColdFront (Schmidt) -> Helpdesk	Help icon -> `/schmidt/login/`	Schmidt SSO -> `/tickets/my-tickets/`
ColdFront (JHU) -> Helpdesk	Help icon -> `/jhu/login/`	JHU SSO -> `/dashboard/` (staff)
Helpdesk (Schmidt) -> ColdFront	Portal icon -> `/schmidt/login/`	Schmidt SSO -> home
Helpdesk (JHU) -> ColdFront	Portal icon -> `/jhu/login/`	JHU SSO -> home

URLs available on both apps:

URL	Purpose
/	`home_or_sso` -- authenticated users go to home; anonymous -> `/portal/`
/portal/	Branded landing (JHU + Schmidt + ARCH Helpdesk cards) -- default front door
/login/	Legacy picker (JHU + Schmidt buttons) -- kept for tests
/user/login/	Same picker (alias) -- kept for tests
/jhu/login/	Direct to JHU OIDC
/schmidt/login/	Direct to Schmidt OIDC
/oidc/jhu/authenticate/	mozilla_django_oidc view for JHU
/oidc/schmidt/authenticate/	mozilla_django_oidc view for Schmidt
/oidc/callback/	Realm-aware callback (backend skips wrong realm)

Logout path (**oidc_logout_view**) drops the user back on /, which forwards to **/portal/** when they arrive anonymous. **LOGOUT_REDIRECT_URL** defaults to



/ in coldfront/custom/config/auth.py.

Post-login destinations on helpdesk:

- **is_staff=True** -> **/dashboard/** (ticket management)
- regular user -> **/tickets/my-tickets/** (own tickets only)

Routing logic:

- ColdFront sidebar **Help** link uses **UserProfile.affiliation.code** ->

context_processors.py

appends **/jhu/login/** or **/schmidt/login/** to **SiteConfiguration.help_url**.

- Helpdesk sidebar **Portal** link uses **session['helpdesk_oidc_realm']**

(or falls back to **ssci-*** username prefix**)

helpdesk_extensions/context_processors.py

appends the same shortcut to **COLDFRONT_URL**.

Cross-realm token exchange: both apps' OIDC backends

(**JHU*Backend**, **Schmidt*Backend**) check **session[SESSION_REALM_KEY]**

and return **None** when the code came from a different realm. Without

this guard, the first backend in **AUTHENTICATION_BACKENDS** would try

to exchange a Schmidt code at the JHU token endpoint and raise

HTTPError 400 before the Schmidt backend got a chance to handle it.

start.sh -- Stack Lifecycle Management

Use **start.sh** to manage the full CLOACK stack. It wraps **docker compose** with idempotency checks, host bootstrapping, and service scoping.

Commands

Command	Description
<code>./start.sh init</code>	Create host user/group, prepare dirs, write <code>.env</code> (only if <code>[dev]</code>)
<code>./start.sh start</code>	Build (if needed) and start all managed services (auto-builds missing images)
<code>./start.sh restart</code>	Stop, rebuild all images (<code>--no-cache</code>), then start (preserves data)
<code>./start.sh stop</code>	Stop managed services
<code>./start.sh destroy</code>	Stop and remove managed containers (keeps volumes/data)



<code>./start.sh rebuild [dev]</code>	Stop, rebuild all images from scratch (<code>--no-cache</code>), then start						
<code>./start.sh populate dev [--partiti</code>	Import departments, setup billing, generate <code>`SlurmJob`</code> (dev only). Flags: <code>--skip-billing`</code> , <code>--with-usage`</code>						
<code>./start.sh [p1,p2,...] ldap-reinit [dev]</code>	If <code>`migration/ldap-*.cat`</code> exists: restore LDAP from backup (passwords patched from <code>`.env`</code>). Otherwise: reinitialize via <code>`install.sh`</code> . Both sync passwords and re-seed						
<code>./start.sh h cleanup</code>	Deprecated unused <code>`arch/*`</code> images not referenced by any container						
<code>[dev]it.sh h prune-volumes</code>	Stop services and delete OpenLDAP named volumes (data loss -- use with care)						
<code>[dev]it.sh h prune-sl</code>	Stop Slurm containers and delete all Slurm volumes (munge key reset)						
<code>./start.sh [dev]` prune-i</code>	Remove all Slurm Docker images (<code>`arch/cn-rockylinux`</code> , <code>`arch/slurmd`</code> , etc.)						
<code>./start.sh h clean-all`</code>	Full cleanup: stop, destroy, remove arch volumes/images/networks only (safe for co-hosted services). In dev mode also removes						
<code>./start.sh h log [SERVI</code>	Triggers for data services (default: <code>`coldfront`</code>)						
<code>./start.sh h</code>	Push updated <code>`slurm.conf`</code> to all nodes (<code>`scontrol reconfigure`</code>)						
<code>./start.sh h test [dev] [MODUL</code>	Run Django unit tests (default: all arch_sync tests). Example: <code>./start.sh test coldfront.plugins.arch_sync.test</code>						
<code>./start.sh h helpdesk {start\</code>	<code>stop_billing`</code>	restart\	status\	destroy\	log\	populate}`	Manage the Helpdesk service. <code>`populate`</code> creates sample queues and

tickets
(idempotent)



<code>`.start.sh mta {create}`</code>	<code>start\</code>	<code>stop\</code>	<code>status\</code>	<code>log}`</code>	Manage the inbound-email MTA (<code>`arch-mta` postfix, edge profile</code>). Receives SMTP via the NPM <code>`:25` stream -> Maildir -> Helpdesk tickets. `--from-ghcr` pulls the published image</code>
<code>`.start.sh help`</code>	Show usage				

Quick start

```
# First-time bootstrap: prepare host dirs and generate .env
./start.sh init dev      # dev
./start.sh init         # prod

# Start the stack
./start.sh start dev     # dev (docker-compose-dev.yml, no /opt/mprov mounts)
./start.sh start         # prod (docker-compose.yml + docker-compose-prod.yml)

# Dry-run: preview what would happen (no changes)
./start.sh start dev --dry-run

# Force a full rebuild + redeploy
FORCE_BUILD=1 FORCE_REDEPLOY=1 ./start.sh start dev

# Restart: stop + rebuild + start (preserves data)
./start.sh restart dev

# Rebuild all images from scratch (--no-cache) and restart
./start.sh rebuild dev

# Reinitialize OpenLDAP and re-sync users into ColdFront
./start.sh ldap-reinit dev # dev
./start.sh ldap-reinit    # prod

# Default: departments + billing rates/configs + SlurmJob
# Naming convention: low/med/high (CPU by tier), a100/l40s/h100/h200 (GPU by hardware)
./start.sh populate dev --partitions low,med,high,a100

# Skip billing rates/AllocationBilling; keep SlurmJob
./start.sh populate dev --partitions low,med,high,a100 --skip-billing

# Include SlurmAccountUsagePeriod (billing/usage pages)
./start.sh populate dev --partitions low,med,high,a100 --with-usage

# Stop / remove containers (keeps volumes)
./start.sh stop dev
./start.sh destroy dev
```

`./start.sh start -- internal flow`

Each **start** runs these steps in order:

1. Start **postgres** and **openldap** first (before all other services)
2. Wait for **postgres** to become healthy + **init-databases.sh** to complete
3. Sync DB role passwords (**ALTER USER admin/coldfront/keycloak**) ensures .env passwords match postgres, even after init regenerated secrets
4. Start all remaining services (**compose_up_scoped**)
5. Post-compose sync re-sync postgres + LDAP passwords (handles container recreation), restart unhealthy app containers if needed
6. Wait for all services to become healthy



Note: LDAP migration from **migration/ldap-*.cat** is not done during **start**. Use **./start.sh ldap-reinit [dev]** after the stack is healthy (see Deploy flows below).

Generate an LDIF dump from the source cluster:

```
docker exec openldap slapcat > migration/ldap-$(date +%F).cat
```

Copy **migration/ldap-*.cat** to the target repo and run **./start.sh ldap-reinit**. The restore preserves:

- UIDs and GIDs
- Password hashes (SSHA)
- PI groups and memberships
- Affiliation counters are auto-calibrated

Native docker compose diagnostics

The script intentionally omits status/log/stats commands use **docker compose** directly:

```
docker compose ps          # service status + health
docker compose logs -f SERVICE # tail logs for a service
docker compose stats       # live resource usage
```

dev vs production

The optional **dev** argument switches the compose file. All commands accept **[dev]**.

Setting	Production	Dev
Compose files	`docker-compose.yml` + `docker-compose-prod.yml`	`docker-compose-dev.yml`
Slurm host mounts	Yes (`docker-compose-prod.yml`)	No
`PLUGIN_SLURM`	`1`	`0` (default)
`SLURM_NOOP`	`0`	`1` (default)
`SEED_ON_RESTART`	`true` (default)	`true` (default)

Compose file overlay **docker-compose-prod.yml** is automatically appended by **start.sh** when running in production mode. It adds bare-metal Slurm bind mounts to **coldfront** and **qcluster**. Docker Compose merges the files: list fields (**volumes**, **ports**) are appended; scalar fields (**image**, **restart**) are overridden by the overlay. **docker-compose-dev.yml** completely redefines those services so it never inherits the prod mounts.

SEED_ON_RESTART controls whether **seed initial data.py** syncs LDAP users into ColdFront on every container start. Set **SEED_ON_RESTART=false** in **.env** to skip on routine restarts. **./start.sh ldap-reinit [dev]** always forces a full sync regardless of this setting.

Environment variables

./start.sh init writes a **.env** file (skipped if already present). Key variables:



```
# Compose file override (alternative to `start dev`)
COMPOSE_FILE_PATH=docker-compose-dev.yml

# Force a full image rebuild
FORCE_BUILD=1

# Force redeployment even if services are already healthy
FORCE_REDEPLOY=1

# Slurm (leave empty on dev; set real paths on production)
SLURM_CONF=/opt/mprov/cloack/slurm/current/etc/slurm.conf
SLURM_MPROV=/opt/mprov/cloack/slurm/current/bin
PLUGIN_SLURM=1
SLURM_NOOP=0
```

Other auto-detected variables written by **init**:

Variable	Purpose	Auto-detected
<code>`DOCKER_SOCKET`</code>	Docker socket path for container management	macOS: <code>`~/docker/run/docker.sock`</code> , Linux: <code>`/var/run/docker.sock`</code>
<code>`ARCH_HOST_PROJECT_ROOT`</code>	Host project root for bind mounts	<code>`\$(pwd)`</code>

T`

JWT token for slurmrestd

start.sh slurm start generates a 30-day JWT token and writes it to **`$(BASE_ETC)/slurm/slurm_jwt.token`**. This file is bind-mounted into coldfront and qcluster containers. The **SlurmClient** reads the token for REST API authentication. A crontab entry (installed by **init**) renews it daily at 3am.

In dev mode, if no **`docker-compose-slurm-*.yml`** overlay exists, **slurm start** auto-generates a default **`docker-compose-slurm-arch-dev.yml`** with the standard 4-node cluster (login01, cn-01, cn-02, cn-03) and default partitions (premium, scavenger). In production, the overlay must be deployed from the ColdFront UI.

Note: JWT auth requires Slurm 23.x+. Docker dev currently uses Slurm 22.05 which has a known JWT bug [upgrade to 25.11 is planned](#).

Compose project isolation

Project	Compose file	Services
<code>`arch`</code>	<code>`docker-compose-dev.yml`</code> (or <code>`docker-compose.yml`</code> + <code>`docker-compose-slurm-arch-dev.yml`</code>)	postgres, openldap, coldfront, keycloak, qcluster,
<code>`arch-slurm`</code>	<code>`docker-compose-slurm-arch-dev.yml`</code>	slurmdbd, slurmrestd, slurm-db, cn-*, login*

Slurm containers are isolated in their own project so **stop/restart** from CLI or UI doesn't affect main services.

Security: password separation

The **.env** file contains several password variables. Two are especially important to keep separate:

Variable	Purpose	Used by
<code>`LDAP_ADMIN_PASS`</code>	Bind DN password (<code>`cn=admin,dc=arch,dc=cluster`</code>) -- full LDAP admin access	ColdFront LDAP bind, Keycloak LDAP federation, <code>`ldapmodify`</code> operations



<code>`JHU_ADMIN_PASS WORD`</code>	User login password (<code>`uid=admin`</code>) -- ColdFront web login	Staff browser login to ColdFront admin
--	---	--

Why they must stay separate: if both use the same password, anyone who knows the ColdFront admin login also has full LDAP admin access (create/delete users, modify ACLs, read all attributes). In production this is a security risk.

Other password variables: **LDAP_ROOT_PASS** (LDAP `cn=root` used internally by `replace_ldif.sh` for LDIF imports), **POSTGRESQL_PASSWORD**, **COLDFRONT_DB_PASSWORD**, **KEYCLOAK_DB_PASSWORD**.

Break-glass **admin** account: after the passwordless rollout (**auth-passwordless-model**), the local Django **admin** user is the only password-backed entry point left in the stack every other path goes through Keycloak / OIDC. Its credential handling (strong random password, enterprise vault, quarterly rotation, audit logging via **admin login alert task**) is documented in **docs/break-glass-admin.md**. Quarterly rotation is automated by **scripts/rotate-admin-password.sh**: it generates a 32-char password via **openssl rand**, applies it inside the container via **manage.py changepassword**, appends an entry to the gitignored audit log, and prints the new value once so it can be pasted into the team vault. Run with **help** for usage / flags.

Deploy flows

update vs deploy -- quick reference

Command	When	Effect
<code>./start.sh update`</code>	**Day-to-day in prod**	Git pull + restart. ~15s coldfront downtime, no data lost. Hot-reload via bind mounts covers Python code changes.
<code>./start.sh deploy dev -y`</code>	Initial bootstrap, full reset for dev/test, or fresh install on a new machine	**Destructive by design.** Wipes containers/volumes/BASE_DATA/BASE_ETC and rebuilds everything. ~5-10 min. Scavenger-only test-jobs (billing manual). Seeds 4 login-ready test users (<code>`ext-staff`/`ext-user`/`ssci-staff`/`ssci-user`</code> , LDAP password <code>`123`</code>).
<code>./start.sh deploy dev -y --paid-account`</code>	Full-coverage validation including paid partitions (low/med/high/premium/a100)	Same as above PLUS seeds paid SlurmAccounts via <code>`billing populate --paid-allocs`</code> before test-jobs. ~8-12 min. Does NOT write billing output rows.
<code>./start.sh rebuild [dev]`</code>	When <code>`requirements.txt`</code> or <code>`Dockerfile`</code> change	Rebuild images without wiping DB/LDAP/volumes.

What update preserves (no data loss)

Data	Where it lives	Touched by `update`?
Tickets (helpdesk)	<code>`\${BASE_DATA}/helpdesk/...`</code> + PostgreSQL volume	? No
Allocations / Projects / Users (ColdFront)	PostgreSQL volume (DB <code>`coldfront`</code>)	? No -- <code>`migrate`</code> applies schema changes, data intact
LDAP Users	<code>`\${BASE_DATA}/openldap/ldap/`</code> + <code>`slapd.d/`</code>	? No
Keycloak realms/users	PostgreSQL (DB <code>`keycloak`</code>)	? No
Slurm jobs history	PostgreSQL (DB <code>`slurm_jobs`</code>) + MariaDB (slurmdbd)	? No
Home/Scratch filesystem	<code>`\${HOME_DATA}`</code> / <code>`\${SCRATCH_DATA}`</code>	? No
Slurm config	<code>`\${BASE_DATA}/slurm/...`</code>	? No
JWT tokens	<code>`\${BASE_ETC}/slurm/slurm_jwt.token`</code>	? No

update only restarts the application containers (coldfront, qcluster, helpdesk). All named volumes and bind mounts stay untouched the restart is a stop/start, not a recreate. Only the Python process is restarted to load the new code from the



bind mount.

Single exception: if a Django migration intentionally drops data (**RunPython(drop_old_table)**), it runs. Review the diff before large updates. In contrast, **deploy dev -y** wipes all of the above that's why the name "deploy" is for initial setup, not routine updates.

Is .env preserved across commands?

Command	`.env` preserved?
<code>./start.sh</code>	? Yes, 100% -- the file is never touched (only <code>`git pull`</code> + <code>`docker restart`</code>).
<code>update.sh start [dev]</code>	? Yes -- <code>`write_env()`</code> only appends missing keys ([start.sh:367-369](start.sh#L367-L369)). Existing values are intact.
<code>./start.sh init [dev]</code>	?? Regenerated -- but the previous <code>`.env`</code> is auto-backed up to <code>`.env.YYYYMMDDHHMMSS`</code> first ([start.sh:973-977](start.sh#L973-L977)). Every <code>`deploy`</code> runs <code>`init`</code> , so you always have the backup.
<code>./start.sh deploy [dev] -y`</code>	?? Same as <code>`init`</code> -- <code>`.env`</code> is regenerated; <code>`.env.<timestamp>`</code> saved.
<code>./start.sh clean-all`</code>	? Not touched -- wipes <code>`\${BASE_DATA}``\${BASE_ETC}`</code> but leaves <code>`.env`</code> at the repo root (it's gitignored).

For custom secrets (Entra/Azure credentials, Trello API keys, OIDC client secrets), use separate files:

- **entra.env** (gitignored, chmod 600) auto-sourced by **init** before **write env**, so values there override any defaults generated.
- **trello.env** (gitignored, chmod 600) same pattern.

This way your secrets survive every **init** / **deploy** the generated **.env** picks them up automatically. The **.env** itself is considered regenerable; secrets sensitive enough to survive a wipe belong in **entra.env** / **trello.env**.

Automatic pre-update backup. **./start.sh update** now snapshots postgres + openldap to **\${BASE_DATA}/backup/** before running migrations:

```

${BASE_DATA}/backup/postgres_YYYY-MM-DD_HHMMSS.sql    # pg_dumpall (all DBs)
${BASE_DATA}/backup/ldap_YYYY-MM-DD_HHMMSS.ldif      # slapcat -n 1

```

Retention is manual delete older files as needed. To restore a bad update:

```

# Postgres
docker exec -i postgres psql -U postgres < ${BASE_DATA}/backup/postgres_*.sql

# OpenLDAP (offline restore)
docker stop openldap
docker exec openldap slapadd -n 1 -l /var/backups/ldap_*.ldif
docker start openldap

```

A. Fresh deploy (no migration)

Full ordered sequence:



```
./start.sh clean-all          # wipe arch volumes/images
./start.sh init dev           # write .env, prep dirs
./start.sh start dev          # bring up coldfront + ldap + keycloak ...
./start.sh slurm create arch-dev # generate Slurm overlay
./start.sh slurm start arch-dev # start slurmd + slurmd + slurmd + ...
./start.sh ldap-reinit dev     # seed minimal LDAP + sync to ColdFront
./start.sh populate dev --partitions a100,l40s,h100,h200 # import depts + seed partitions (no bi...)
./start.sh helpdesk start      # build + start helpdesk (+ helpdesk-qc...)
./start.sh helpdesk populate   # sample queues + tickets
./start.sh ai pull             # download local LLM weights (sha256-ve...)
./start.sh ai status           # AI engine health (local triage -- see ...)
```

Add **with-billing** to **populate** if you want fake **BillingRate/AllocationBilling** records.

About the **dev** token: it is an **environment modifier** (not a cluster name). When present as the first arg after the subcommand, it switches the compose overlay from **docker-compose.yml** (prod) to **docker-compose-dev.yml**. **clean-all** and **slurm** ignore it (for **slurm create**, the second arg is the cluster name). **helpdesk** respects it. **./start.sh helpdesk dev** uses the dev compose file (needed so the container's compose hash matches the rest of the dev stack; otherwise **./start.sh start dev** flags helpdesk as an orphan).

B. Deploy + migrate from production

Option 1: CLI

```
./start.sh clean-all dev
./start.sh init dev
./start.sh start dev
./start.sh ldap-reinit dev      # restore LDAP from migration/ldap-*.cat
./start.sh slurm create arch-dev
./start.sh slurm start arch-dev
```

Option 2: UI (after **init + start**)

1. Import Users -> Migration upload slapcat .cat file (creates LDAP structure + users + groups + projects)
2. Admin -> Clusters -> Add -> Deploy create and start Slurm cluster
3. Sync Slurm from cluster detail page (creates Slurm accounts)

C. Reinitialize LDAP (no migration backup)

```
# Runs install.sh to create minimal LDAP structure (root + munge + slurm)
./start.sh ldap-reinit dev
```

How **ldap-reinit** works:

- If **migration/ldap-*.cat** exists -> restores from backup via **slapadd**, patches all passwords (**cn=root**, **cn=readonlyroot**, **uid=admin**) with SSHA hashes from **.env**, dedupes LDIF by **mail** before **slapadd** (keeps the posixAccount with lowest **uidNumber** per email, drops matching **cn=** PI groups, scrubs dangling **member:/memberUid:** refs), syncs **olcRootPW** (config DB), seeds ColdFront
- If no cat file -> removes init marker, restarts openldap (re-runs **install.sh**), syncs passwords, seeds ColdFront

The mail-dedup mirrors the Import Users -> Migration preview UI (coldfront/custom/plugins/arch_sync/staff_views.py **ImportUsersView**): the UI shows a dual-list of Available vs. Chosen entries with per-row conflict flags and blocks



commits that still carry duplicates (admin can override via "Confirm anyway"). The CLI path is non-interactive and applies the same rule deterministically so CI/scripts never fail **ModelDuplicateException** at Keycloak sync time. Keycloak realms **jhu** and **schmidt** keep **duplicateEmailsAllowed=false** as a hard-asserted invariant (keycloak/configure-keycloak.sh **configure realm themes**).

D. LDAP cleanup (after migration restore)

Migration dumps may contain test accounts not in **install.sh**. Use **ldap-cleanup.sh** to find and remove them:

```
# Dry-run -- compare migration dump against install.sh, save extras to scripts/ldap-extras.txt
./scripts/ldap-cleanup.sh --dump migration/ldap-2026.04.03.10.cat

# Apply on live LDAP (deletes extra entries)
./scripts/ldap-cleanup.sh --apply
```

OpenLDAP diagnostics (manual):

```
# Verify LDAP entries
source .env
docker exec openldap ldapsearch -H ldapi:/// -x \
  -D "cn=admin,dc=arch,dc=cluster" -w "${LDAP_ADMIN_PASS}" \
  -b "dc=arch,dc=cluster" "(uid=*)" uid | grep "^uid:"

# Force LDAP -> ColdFront re-sync
./start.sh ldap-reinit dev
```

Deploy pipeline (4 stages)

The CLOACK stack moves through 4 environments along a single pipeline. Each stage has a clear responsibility; tagged GHCR images flow under explicit operator control. Full strategy in **docs/deploy-pipeline-4-stages.md**.

[1] dev laptop	[2] dev VM Proxmox	[3] edulogin bare-metal	[4] portal bare-metal
macOS arm64	Linux amd64	Linux amd64	Linux amd64
build local	build local	pull GHCR :dev	pull GHCR :prod
git push code	git pull code	git pull code	git pull code
(no image push)	docker push :dev	docker push :prod	(no push)
	docker push :sha-XXX	docker push :rc-<UTC>	
		(retag, not rebuild)	

Tag scheme

Tag	Created by	Mutable?	Purpose
`:sha-<git-short>`	Stage 2 (publish)	**Immutable**	Snapshot tied to a git commit. "What code is in this image?"
`:dev`	Stage 2 (publish)	Mutable head	Latest Stage 2 build. What edulogin pulls.
`:rc-YYYY-MM-DD-HHMM`	Stage 3 (promote)	**Immutable**	Validated + promoted at this exact UTC moment. Rollback target.
`:prod`	Stage 3 (promote)	Mutable head	Latest validated build. What portal pulls.



Rollback is a retag (**:prod** <- previous **:rc-*****), **not a rebuild** same digest, new label. The original **:rc-***** **is untouched and serves as immutable** "what was in prod between X and Y" record.

Subcommands

Subcommand	Stage	Status
<code>./start.sh publish [--tag <tag>] [--check]</code>	2	**Shipped (dormant)** -- refuses dirty working tree; needs GHCR auth to
<code>./start.sh promote [--from <tag>] [--check] [--rollback <rc-tag>]</code>	3	**Shipped (dormant)** -- refuses until <code>`docker-compose-ghcr.yml`</code> lands
<code>./start.sh verify [--from-ghcr [dev prod]]</code>	3 & 4	**Shipped (dormant)** -- pulls cloack images from GHCR; refuses if <code>`docker-compose-ghcr.yml`</code> is missing

The **promote** subcommand pulls **:dev**, runs **scripts/verify-deploy.sh** as pre-flight gate, then retags + pushes **:prod** + **:rc-** atomically (all source images must pull successfully before any push). See **./start.sh promote help** for full options.

Phase 2 multi-FQDN compat (2026-06-10): **KC_HOSTNAME_STRICT** + **KC_HOSTNAME_BACKCHANNEL_DYNAMIC** env vars on Keycloak, **HELPDESK_SCRIPT_NAME** for helpdesk subpath mount, per-realm **KEYCLOAK_HOSTNAME_JHU** / **_SCHMIDT** driving **configure_realm_frontend_urls** in **configure-keycloak.sh**. All defaults preserve single-FQDN dev behaviour. See **docs/deploy-pipeline-4-stages.md** Phase 2 for the full operator env-var glossary, and **docs/roles/admin.md** for the runbook.

Role-based User Guide: five role guides ship at **docs/roles/** pick the one that matches your role (User / PI / Staff / Superuser / Admin) and follow the 5-section template (scope, quick start, reference matrices, troubleshooting, where to go next). **docs/CLOACK_User_Guide.pdf** is the printable / offline mirror.

Bare-metal install -- cloack-deploy RPM

For the full bare-metal deployment guide (5 playbooks, what each delivers end-to-end, what is intentionally NOT covered, single source-of-truth conventions, version pin policy), see



slurm/ANSIBLE.md.

The production artefact is the **cloak-deploy** RPM (Rocky 9). It ships the Ansible tree plus a **cloak-deploy** CLI wrapper, all under **/opt/cloack/**, with **/usr/bin/cloack-deploy** as a PATH symlink so cron and systemd find it without touching **/etc/profile.d/**.

```
# 1. Install on a deploy workstation (Rocky 9)
dnf install cloak-deploy

# 2. Point at your ARCH Portal and get a DRF token
export CLOACK_API_URL=https://coldfront.jhu.edu
export CLOACK_API_TOKEN=<token from /api/v1/auth/token/>

# 3. Fetch the per-cluster inventory + run the read-only pre-flight
cloak-deploy fetch <cluster>          # -> /opt/cloack/etc/inventory.d/<cluster>/
cloak-deploy verify <cluster>         # ansible-playbook slurm-verify (no changes)
cloak-deploy prod-sim <cluster>       # ansible-playbook --check --diff
cloak-deploy server <cluster>         # real: slurmctld/slurmdbd/slurmrestd
cloak-deploy client <cluster>         # real: slurmd on compute + login nodes
```

The bundle endpoint lives at

GET /api/v1/clusters//ansible-bundle.tar.gz (DRF token auth;

superuser or staff with the cluster's **Resource** in their

StaffResourceAssignment scope). Build the RPM locally with

make rpm (requires **rpm-build**) or via

.github/workflows/rpm.yml on tag push.

Spec + CLI wrapper: packaging/rpm/.

Source-tarball assembler: scripts/build_rpm_source.py.

Docker Images

All images use the **arch/** namespace. Build order matters for Slurm images (base first).

Component Versions

Component	Version
ColdFront	1.1.7
Django	4.2
Python	3.10
Slurm	22.05.9
PostgreSQL	16 (supported until Nov 2028; v17 available but no significant benefits for Django/ColdFront)
Keycloak	26.5
MariaDB	10.11
OpenLDAP	1.5.0 (osixia)



Rocky Linux	9 (Slurm base)
Munge	bundled with Slurm

Core Services

Image	Dockerfile	Base	Description
`arch/coldfront`	`coldfront/Dockerfile`	`python:3.10-slim`	Django portal + django-q worker. Runs ColdFront with custom plugins
`arch/openldap`	`openldap/Dockerfile`	`osixia/openldap:1.5.0`	OpenLDAP server with multi-arch support (amd64/arm64)
`arch/keycloak`	`keycloak/Dockerfile`	`keycloak:latest`	Keycloak OIDC with custom themes and LDAP providers
`arch/keycloak-config`	`keycloak/Dockerfile.conf`	`debian:bookworm-slim`	One-shot: configures Keycloak realms and LDAP federation
`arch/postgres`	`postgres/Dockerfile`	`postgres:16`	PostgreSQL with multi-database init (coldfront, keycloak, slurm_jobs)
`arch/phpldapadmin`	`phpldapadmin/Dockerfile`	`osixia/phpldapadmin:lates`	phpLDAPadmin web UI for LDAP browsing

t`

Slurm Services (build order: base -> munge -> slurmdbd -> slurmctld -> slurmd -> slurmrestd)

Image	Dockerfile	Base	Description
`arch/cn-rockylinux`	`slurm/base/Dockerfile`	`rockylinux:9`	Base image: Rocky 9 + Slurm 25.11 + Munge + SSSD + SPANK plugins (myspank, spank_reemit_cost, pyxis)
`arch/slurm-munge`	`slurm/munge/Dockerfile`	`arch/cn-rockylinux`	One-shot: generates shared munge key in `munge-key` volume
`arch/slurmdbd`	`slurm/slurmdbd/Dockerfile`	`arch/cn-rockylinux`	Slurm accounting daemon (connects to MariaDB)
`arch/slurmctld`	`slurm/slurmctld/Dockerfile`	`arch/cn-rockylinux`	Slurm controller daemon
`arch/slurmd`	`slurm/slurmd/Dockerfile`	`arch/cn-rockylinux`	Slurm worker/login node daemon -- also ships NVIDIA enroot runtime for `--container-image`
`arch/slurmrestd`	`slurm/slurmrestd/Dockerfile`	`arch/cn-rockylinux`	Slurm REST API daemon
`arch/mariadb`	`slurm/mariadb/Dockerfile`	`mariadb:10.11`	MariaDB for Slurm accounting (port 3307)

Other

Image	Dockerfile	Base	Description
`arch/ondemand`	`ondemand/Dockerfile`	`rockylinux:9`	Open OnDemand web portal

Partitions & Allocation Tiers

ARCH standardizes partition and QoS naming across all clusters so users submit jobs the same way regardless of which cluster they target.

Partition naming convention

Kind	Names	Notes
------	-------	-------



CPU (by billing tier)	`low`, `med`, `high`, `premium`	Differ in `TRESBillingWeights` (cost per core-hour). Tier code on `BillingRate` joins to the partition. `premium` spans all CPU nodes with a 7-day wall and excludes `interactive` QoS.
GPU (by hardware)	`a100`, `l40s`, `h100`, `h200`	Submit to the partition matching the GPU type the job needs. `kind=gpu`, `hardware=<model>`. Default `a100` uses `gpu-01` with `Gres=gpu:2` and `TRESBillingWeights="GRES/gpu=2.0"`.

QoS

Name	Purpose
`scavenger`	Free, **preemptible** . Long max wall, no cost. May be killed if a paid job needs the resources.
`interactive`	Higher priority (1000), short max wall (4h), capped resources per user. For tuning, debugging, notebooks.
`normal` / `premium`	Standard paid QoS for batch work -- retained for backward compat where allocations require them.

Submission examples

```

sbatch --partition=med --qos=normal job.sh           # paid CPU work
sbatch --partition=scavenger --qos=scavenger run.sh  # free, preemptible
srun   --partition=a100 --qos=interactive --pty bash  # 4h interactive on A100

```

Seeding (per cluster, idempotent)

```
./start.sh populate dev --partitions a100,l40s,h100,h200,low,med,high,scavenger
```

SlurmPartition records are created per **SlurmCluster** from the **partitions** CSV. **SlurmQOS** records (**scavenger** / **interactive** / **normal** / **premium**) are seeded automatically by **slurm sync** and do not require a CLI flag. Multi-cluster federation: each **SlurmCluster** gets its own copy.

Billing & Invoicing System

ARCH includes a comprehensive billing system built into ColdFront to track resource usage (compute, GPU, storage, backup) and manage invoicing for JHU and Schmidt Sciences allocations.

Key Features

Two-Track Billing Model

The system supports three billing types:

Type	For	How It Works
Pay Per Use	PIs without hardware credit pool	Full rate charged per hour (compute + O&M bundled)
Credit Holder	PIs who purchased hardware	Credits cover compute hours only; O&M (power/cooling) charged separately in cash
No Charge	Default, scavenger, and Schmidt allocations	Usage tracked but not billed (\$0 total)

Important: Credits do NOT pay O&M bills. The system UI clearly separates credit consumption from cash charges to prevent confusion.



Billing eligibility is controlled by **AllocationBilling.is_billable** (set automatically by signal):

- **is_billable=True** -> allocation is charged (Pay Per Use or Credit Holder)
- **is_billable=False** -> allocation is tracked with **billing_type='no_charge'** and **payment_status='no_charge'**

Allocations with **is_billable=False** still appear in billing reports showing real usage (hours, rate) but with \$0 total and "No Charge" label. Staff can toggle **is_billable** via Django Admin. The billing report includes a "Billable only" checkbox (default ON) to filter the view.

Staff Billing Report

Superusers and billing staff can:

- View filterable billing report with all usage and charges
- Filter by: PI, project, user, school, department, affiliation, billing type, date range
- See summary metrics: total jobs, CPU hours, GPU hours, average hours/job, total cost
- Download CSV export for accounting integration
- Access at: **/billing/report/**

PI Usage & Billing Dashboard

PIs can:

- See credit balance and estimated runway (days until exhaustion)
- View separate O&M charges in current billing period (30-day rolling window)
- Track recent usage by period
- Download printable invoices with IOS/cost center number
- Access at: **/billing/my-usage/**

Rate Management

Superusers can:

- Define rates for compute (CPU/GPU), storage, and backup services
- Set different rates per cluster and partition
- Create rate history with effective date ranges
- Rates are stored separately for:
 - **cpu_full / gpu_full** full rates for pay-per-use allocations
 - **cpu_om / gpu_om** O&M only rates for credit holders

Hardware Credit Pools

Superusers track hardware purchases:

- Record hardware value in dollars
- Convert to compute credit hours (e.g., "4x H200 GPU" -> "87,600 GPU-hours")
- Track credit consumption and remaining balance
- Set expiry dates if credits sunset

IOS Numbers & Department Tracking



Allocations can have:

- IOS Number cost center / account number for JHU service center
- Department organizational affiliation (e.g., "Computer Science", "Biology")
- Billing Type pay-per-use, credit holder, or no charge
- Provisioned Storage/Backup manual entry for non-compute charges

Usage Guide

Setting Up Billing

1. Enable billing collection (environment variable):

```
export BILLING_COLLECT_ENABLED=1
```

2. Create billing rates (Django Admin or CLI):

```
# Via Admin: /admin/arch_sync/billingrate/add/
# Create rates for cpu_full, gpu_full, cpu_om, gpu_om, storage, backup
```

3. Import departments (optional, for organizational grouping):

```
python manage.py import_departments departments.csv
# CSV format: name,code,school,affiliation
```

4. Assign PIs to departments (optional):

```
python manage.py assign_pi_departments pi_departments.csv
# CSV format: username,department_name
```

5. Set allocations as billable:

- Open Allocation Detail
- Scroll to "Billing Info" card
- Set IOS number and billing type (staff/superuser only edit form is not shown to PIs or managers)
- Save

Automatic creation on approval: When an allocation transitions to Active (approved by admin), an `Allocatio...



Visibility rules for Billing Info card (`allocation\_detail.html`):

The card is shown when any of these conditions is true:

- `request.user.is\_superuser` or `can\_approve\_deny` -- superusers and billing staff
- `request.user == allocation.project.pi` -- the PI who owns the project
- `is\_allowed\_to\_update\_project` -- project managers (`ProjectPermission.UPDATE`)

Role	Can see IO number & billing type?	Can edit (IO / billing type)?
Superuser / Staff (`is_superuser` or `can_approve_deny`)	Yes	Yes
PI (`request.user == allocation.project.pi`)	Yes (read-only)	No
Project Manager (`ProjectPermission.UPDATE`)	Yes (read-only)	No
Regular allocation user	No	No

The edit form (IO Number input + Update Billing button) is wrapped in its own `{% if request.user.is\_superuser or can\_approve\_deny %}` guard and is never shown to PIs or managers.

Collecting Billing Data

Billing usage data is stored in two models:

Model	Source	Purpose
`SlurmJob`	`collect_sacct` (from `sacct`) -- runs every 15 min	Per-job records -- source of truth for Job Metrics dashboard
`SlurmAccountUsagePeriod`	`sync_job_usage` -- manual only , fired by Compute Billing on <code>/billing/report/</code>	Per-period usage snapshots -- source of truth for billing

Job data collection (**collect_sacct**, every 15 min):

```
# Collect last 7 days of per-job data from sacct
coldfront collect_sacct --current-week

# Collect current month
coldfront collect_sacct --current-month

# Import from a pipe-delimited sacct dump file (for dev/testing)
coldfront collect_sacct --from-file coldfront/templates/jobs_usage.csv

# Dry-run (parse and log, no DB writes)
coldfront collect_sacct --dry-run
```

Billing aggregation manual only. **SlurmAccountUsagePeriod** rows (the "Usage & Charges" table on `/billing/report/`) are never written by a scheduled task or by the `./start.sh deploy` seeding flow. Staff must press **Compute Billing** on `/billing/report/`, which internally runs:

```
# Recompute the is_billable classification on all AllocationBilling rows
coldfront recalculate_billing --all

# Aggregate SlurmJob -> SlurmAccountUsagePeriod (per allocation x day)
coldfront sync_job_usage
```

Rationale: staff wanted a predictable, explicit compute step so the ledger never advances on its own (see **project_billing_manual_only** memory). The 15-min **sync_slurm_all** task keeps the Core Usage (Hours) gauge and per-user snapshots fresh, but it leaves the billing ledger untouched.



Offline recompute / audit. For finance handoff, historical recompute, or air-gapped review, **scripts/billing_offline.py** produces the same **PI,IO,Total** summary from a Slurm **sacct** export, a finance-supplied **accounts.csv** (**[pi,]account,io[,share,storage_tb,backup_tb]**), and a copy of **slurm/conf/rates.lua** pure stdlib, zero ColdFront dependency. Multi-IO splits use **share percent** mirroring **AllocationCostCenter**. The optional **pi** column lets the script work with project-based account naming (e.g. Discovery's **pmcmood-projects**) where PI cannot be derived from the account name see **scripts/README_billing.md** and **scripts/README_billing_DISCOVERY.md**.

Printable Invoices

PIs can view and print invoices from their dashboard:

- Access: **/billing/my-invoice/**
- Browser print (Ctrl+P or Cmd+P) -> Save as PDF
- Invoice shows two sections:

1. Compute Credits Used (hardware-backed hours, NOT a cash charge)

2. O&M Charges Due (real dollars owed for power/cooling)

- Includes IOS number for payment routing

Example Data & Quick Test

We include example data to test the billing system (see **BILLING_SETUP_GUIDE.md** for full details):

Departments included (12 total):

- Whiting School of Engineering: CS, ECE, Applied Physics, ME, BME
- Krieger School of Arts & Sciences: Biology, Chemistry, Physics
- School of Medicine: Neuroscience, MBG
- Schmidt Science: Structural Biology, Computational Biology

PIs included (10 total):

- rdesouz4 (Computer Science)
- jsmith5 (Biology), mchen2 (Applied Physics), agarcia7 (ECE)
- bpatel9 (Chemistry), kwilson3 (Neuroscience)
- dgonzalez1 (Structural Biology), hjones8 (Computational Biology)
- Irodriguez4 (Physics), ekim6 (Mechanical Engineering)

Example rates created (6 total):

- CPU Full Rate: \$0.42/hr | CPU O&M Only: \$0.08/hr
- GPU Full Rate: \$3.50/hr | GPU O&M Only: \$0.65/hr
- Storage: \$0.012/TB/month | Backup: \$0.006/TB/month

To populate example data:



```
# Import departments (12)
python manage.py import_departments examples_departments.csv

# Assign PIs to departments (10)
python manage.py assign_pi_departments examples_pi_departments.csv

# Create billing rates (6)
python manage.py shell < populate_example_rates.py

# Then in ColdFront UI:
# 1. Go to each PI's allocation
# 2. Scroll to "Billing Info" card
# 3. Set IOS Number (e.g., 1010500000) and Billing Type
# 4. Click Update Billing

# Collect usage
python manage.py collect_billing --current-period

# View reports
# Staff: /billing/report/
# PI: /billing/my-usage/
```

See **BILLING_SETUP_GUIDE.md** for detailed instructions and troubleshooting.

Staff Workflows

View all usage:

```
Staff menu -> Billing Report
```

Manage rates:

```
Staff menu -> Rate Management -> Add Rate
```

Track hardware credits:

```
Staff menu -> Hardware Credits -> Add Credit Pool
```

Set allocation billing info:

```
Allocation Detail -> Billing Info card -> Set Billing
```

Job Usage Metrics Dashboard

An XDMoD-style analytics dashboard for job usage, available at **/billing/jobs/**:

- Group by: User, Account, Partition, QOS, State, Department (staff only), School (staff only)
- Sortable columns: CPU Hours, GPU Hours, Wall Hours, CPU Hrs/Job, Avg Wait, Jobs, Users, Avg CPUs, % of CPU Total
- Summary cards: Total CPU/Wall hours, job count, active users, avg wait/wall per job; GPU Hours card only appears when **total_gpu > 0**
- CSV export and Raw Jobs detail view with pagination
- Access control data is scoped by role:



Role	Jobs visible
Staff / Superuser	All jobs across all accounts
PI (`userprofile.is_pi`)	Their own jobs + all jobs under accounts in their allocations
Regular user	Only their own jobs (`user=username`)

Data Model

Model	Purpose
`BillingRate`	Rate catalog with effective dates (CPU/GPU/storage/backup)
`HardwareCreditPool`	Hardware purchases converted to credit hours
`AllocationBilling`	Billing config per allocation (IOS, type, storage, backup)
`SlurmJob`	Per-job record from sacct -- source of truth for Job Metrics dashboard
`SlurmAccountUsagePeriod`	Per-period usage snapshot (CPU, GPU, job counts) -- billing source of truth
`Department`	Organizational hierarchy for grouping PIs

Extended models:

- **UserProfile** now has **department** FK for filtering by school/department

Configuration

Relevant environment variables:

```
# Enable billing data collection from Slurm
BILLING_COLLECT_ENABLED=1

# Invoice feature (upstream ColdFront)
INVOICE_ENABLED=1
INVOICE_DEFAULT_STATUS=Payment Pending

# Slurm sync (required for billing data)
SLURM_BIN=/opt/mprow/cloack/slurm/current/bin
SLURM_NOOP=0
```

URL Endpoints

Endpoint	Access	Purpose
`/billing/report/`	Staff/Superuser	Filterable usage report
`/billing/report/csv/`	Staff/Superuser	CSV export
`/billing/rates/`	Staff/Superuser	Rate catalog
`/billing/rates/create/`	Superuser	Add rate
`/billing/rates/<id>/edit/`	Superuser	Edit rate
`/billing/credits/`	Staff/Superuser	Hardware credit pools
`/billing/credits/create/`	Staff/Superuser	Add credit pool
`/billing/my-usage/`	PI	Usage dashboard
`/billing/my-invoice/<allocation-id>/`	PI	Printable invoice
`/billing/jobs/`	PI/Staff	Job Usage Metrics dashboard (group-by analytics)
`/billing/jobs/detail/`	PI/Staff	Raw job list with filters and pagination
`/billing/jobs/csv/`	PI/Staff	CSV export of job metrics



<code>`/allocation/<id>/'</code>	All users	Allocation Detail -- Billing Info card visible to superuser, staff, PI, and project managers (read-only for PI/manager; edit form for staff/superuser only)
--	-----------	---

Examples

Example 1: Credit Holder Invoice

PI Ricardo purchased 4x H200 GPUs -> granted 87,600 GPU-hours of credit.

In a billing period:

- GPU usage: 5,000 hours
- O&M rate: \$0.30/hour
- O&M charge: \$1,500 (real dollars owed)
- Credits consumed: 5,000 hours
- Invoice shows:
 - Section 1 (Credits): 5,000 GPU-hours consumed (no charge, from credit pool)
 - Section 2 (O&M): \$1,500 due in cash

Example 2: Pay-Per-Use Allocation

PI Jsmith does not have a hardware purchase, uses pay-per-use rate.

In a billing period:

- CPU usage: 10,000 hours
- Full CPU rate: \$0.05/hour
- Total charge: \$500 (compute + O&M bundled)
- Invoice shows:
 - Single line: 10,000 CPU-hours @ \$0.05/hour = \$500

Example 3: Filter by Department

Staff runs billing report filtered by department "Computer Science":

- See all allocations in that department
- View usage and costs grouped by PI and project
- Export to CSV for departmental billing

Migration & First-Time Setup

1. Apply database migrations:

```
python manage.py migrate
```

2. Create initial rates via Django Admin:

- Log in as superuser
- Go to `/admin/arch_sync/billingrate/add/`



- Add rates for each service type and effective period

3. Mark allocations as billable:

- Open each allocation in Allocation Detail
- Click "Set Billing" in the Billing Info card
- Enter IOS number and select billing type

4. Collect baseline data:

```
python manage.py collect_billing --current-period
```

5. Monitor via:

- Staff -> Billing Report (usage overview)
- Staff -> Hardware Credits (credit consumption)
- PI -> My Usage & Billing (personal dashboard)

Troubleshooting

No usage data appearing?

- Verify **BILLING_COLLECT_ENABLED=1** is set
- Check Slurm **sreport** binary is accessible: **which sreport**
- Run **collect_billing dry-run** to debug Slurm queries
- Check allocation is marked "Active" status

Allocation not in billing report?

- Verify **AllocationBilling.is_billable=True** is set (via Allocation Detail -> Billing Info)
- Ensure allocation has usage data (run **collect_billing**)

Wrong rates showing?

- Verify rate **effective_from** date is on or before the billing period
- Check rate **effective_to** date is after the period (or NULL for ongoing rates)
- Confirm cluster/partition match your allocation

Populating Synthetic Data for Testing

For development and testing, use **./start.sh populate dev** to seed departments, billing structure, and SLURM usage data in one command.

Models

Model	What	Page
`Affiliation`	Organizational domain (JHU, Schmidt) -- managed via Admin, FK on UserProfile/Department/School. Includes `next_uid`/`next_gid` counters for centralized UID/GID allocation	Admin
`BillingRate`	Rate catalog: cost per CPU/GPU/storage/backup per partition (e.g. `\$0.35/hr CPU`, `\$1.20/hr GPU`)	`/billing/rates/`
`AllocationBilling`	Per-allocation config: IO number, billing type, is_billable flag, storage/backup TB, department	`/allocation/<id>/`



`SlurmAccountUsagePeriod`	Monthly aggregated usage: CPU hours, GPU hours, job count per account + payment status	`/billing/report/` `/billing/my-usage/`
`SlurmJob`	Individual job: user, partition, CPU/GPU used, duration, state	`/jobs/metrics/`

Quick Start

```
# Default: departments only (billing data SKIPPED)
./start.sh populate dev --partitions a100,l40s,h100,h200,low,med,high,scavenger

# Generate fake BillingRate / AllocationBilling / credit pools too
./start.sh populate dev --partitions a100,l40s,h100,h200,low,med,high,scavenger --with-billing
```

What populate does (in order)

Step	Command	Creates
1	Copy `examples_*.csv`	`departments.csv`, `pi_departments.csv` (skip if already present)
2	`import_departments`	`Department` records from `departments.csv`
3	`populate_fake_billing_data --no-usage`	`BillingRate`, `AllocationBilling`, `HardwareCreditPool` -- **only with `--with-billing`**

Defining billing manually (Admin)

For real (non-fake) data, use Django Admin (/admin/arch_sync/). The full pricing pipeline:

1. *BillingRate* -- /admin/arch_sync/billingrate/

Catalog of rates with effective-date history. One row per rate type per cluster/partition.

Field	Purpose
`cluster`	Apply to one cluster (blank = all clusters)
`partition_name`	Apply to one partition (blank = all partitions)
`rate_type`	`cpu_full`, `cpu_om`, `gpu_full`, `gpu_om`, `storage`, `backup`
`rate_per_unit`	Decimal \$ per SU (CPU), GPU-hour, or TB-month
`effective_from` / `effective_to`	Date range -- leave `effective_to` blank for current rate

Two-track pricing: *****_full for pay-per-use PIs**, *****_om** (O&M only, lower) for hardware-credit holders.

2. *AllocationBilling* -- /admin/arch_sync/allocationbilling/ (one per Allocation)

Per-allocation billing config. Created/updated only when staff runs Compute Billing on </billing/report/> or the **coldfront recalculate_billing** management command no signal auto-creates these rows.

Field	Purpose
`io_number`	JHU service-center cost code (no IO -> `is_billable=False`)
`department`	FK to Department -- drives the Slurm account hierarchy
`billing_type`	`pay_per_use` / `credit_holder` / `no_charge`
`is_billable`	Single source of truth -- `False` excludes from charges (still tracked)
`storage_tb` / `backup_tb`	Provisioned storage/backup in TB (multiplied by `storage`/`backup` rate per month)
`qos`	FK to SlurmQOS -- overrides cluster default (e.g. `interactive`)



3. HardwareCreditPool -- /admin/arch_sync/hardwarecreditpool/

When a PI buys hardware, the dollar value is converted into prepaid compute credits. Credits cover compute hours only; O&M is billed in cash via *****_om rates**.

Field	Purpose
`allocation`	FK to the Allocation receiving the credits
`description`	Free text (e.g. "4x H200 GPU purchase FY2026")
`hardware_value`	Dollar value of the hardware
`total_credits`	Compute hours granted (e.g. 87600 GPU-hours)
`credit_type`	`cpu` / `gpu` / `storage` / `memory` (matches `BillingRate.rate_type` prefix)
`granted_date` / `expiry_date`	Validity window

credits_used is gated by **allocation.billing.billing_type == 'credit_holder'**: a pool on a pay-per-use or no_charge allocation reports 0.0 even when **SlurmAccountUsagePeriod** rows exist. Staff flipping **billing_type** away from credit_holder after the fact does not double-bill historical usage.

4. SlurmPartition.tres_billing_weights -- /admin/arch_sync/slurmpartition/

Per-partition multiplier injected as **TRESBillingWeights=** in **slurm.conf**. Combined with **BillingRate** to compute the actual charge. Example: **CPU=1.0,Mem=0.25G,GRES/gpu=2.0** -> 1 GPU-hour = 2 billing units.

5. SlurmAccountUsagePeriod -- /billing/report/, /billing/my-usage/

Per-day aggregation (allocation x account x date) of CPU hours, GPU hours, and job count. Drives invoices and PI dashboards. Not auto-populated. Rows appear only after staff presses Compute Billing on **/billing/report/**, which runs **coldfront sync_job_usage** internally.

Conversion formulas

From	To	Formula
Dollars	Slurm `GrpTRESMins=billing=`	`\$ x 1000 x 60` (1 SU = \$0.001)
Slurm SU	Dollars	`SU x 0.001`
Hours	Dollars	`hours x rate_per_unit` (looked up by `rate_type`)

Helpers: **coldfront/custom/plugins/arch_sync/billing_queries.py** exports **dollars_to_grptres_mins**, **grptres_mins_to_dollars**, **hours_to_dollars**.

Individual commands (via docker exec)

Import departments only:



```
# From default CSV
docker exec -it coldfront python -m django import_departments \
/opt/arch/coldfront/templates/departments.csv

# From custom CSV (copy file first)
docker cp /path/to/my_departments.csv coldfront:/tmp/departments.csv
docker exec -it coldfront python -m django import_departments /tmp/departments.csv --dry-run
docker exec -it coldfront python -m django import_departments /tmp/departments.csv
```

Regenerate SlurmJob records only:

```
docker exec -it -e SLURM_PARTITIONS="premium,scavenger,gpu" coldfront bash -c "
python -m django shell < /opt/arch/coldfront/templates/generate_sacct_dump.py && \
python -m django collect_sacct --clear --from-file /opt/arch/coldfront/templates/sacct_dump.txt
"
```

Clear all SlurmJob records:

```
docker exec -it coldfront python -m django collect_sacct --clear --from-file /dev/null
```

Assign departments to PIs:

```
# CSV format: username,department_names (semicolon-separated)
docker exec -it coldfront python -m django assign_pi_departments \
/opt/arch/coldfront/templates/pi_departments.csv
```

Expected Results

After populating, you should see:

In Billing Report (**/billing/report/**):

- 4,200+ rows of usage data
- 700+ allocations
- Spanning 6 months of history
- Total cost across all allocations: \$100,000+
- Filterable by: PI, department, school, affiliation, billing type

In PI Dashboard (**/billing/my-usage/**):

- Credit balance for hardware credit holders
- O&M charges owed for all allocations
- Usage trends over 6 months
- Estimated runway (days until credits expire)

In Allocations:

- Each allocation shows IOS number and billing type
- Department assignment visible
- Billing history available



CSV File Formats

slurm_usage.csv

```
allocation_id,account,start_date,end_date,cpu_hours,gpu_hours,job_count,partition_name
1000,alloc_1000,2025-12-01,2025-12-31,45000.5,12000.75,1560,gpu
1001,alloc_1001,2025-12-01,2025-12-31,21000.5,0.0,450,cpu
```

pi_departments.csv

```
username,department_names
jbiondo2,Computer Science;Applied Physics
rdesouz4,Computer Science;Biomedical Engineering
ssci-huan,Data Science;Computational Biology
```

departments.csv (for import_departments command)

```
name,code,school,school_code,affiliation
"Computer Science","CS","Whiting School of Engineering","wse","jhu"
"Data Science","DS","Schmidt Sciences Institute","schmidt","schmidt"
```

Billing Report Files (Generated by generate_billing_report command)

The billing system generates three complementary CSV reports plus a formatted PDF saved to **/opt/arch/coldfront/report/**:

1. billing_YYYY-MM_detail.csv Full breakdown per allocation

```
PI,Project,Allocation ID,Department,IO Number,Month,Jobs,CPU Hours,GPU Hours,CPU Cost,GPU Cost,Storage (TB)...
arochab1,arochab1,1299,Computer Science,NO_IO,2026-03,10,3532.70,0.00,176.63,0.00,0.00,0.00,176.63
bblanch6,bblanch6,1300,Computer Science,NO_IO,2026-03,28,5404.41,0.00,270.22,0.00,0.00,0.00,270.22
```

Use case: Detailed allocation-level tracking for finance/compliance audits.

2. billing_YYYY-MM_summary.csv Simplified format (PI, IO, Total ONLY)

```
PI,IO,Total
arochab1,NO_IO,$176.63
bblanch6,NO_IO,$270.22
jperafa1,NO_IO,$874.96
ssci-adityag,NO_IO,$1740.32
```

Use case: Quick summary for executive reports, dashboard displays, or simplified invoicing.

3. billing_YYYY-MM_by_io.csv Aggregated by IO number with all metrics



```
PI,IO Number,Jobs,CPU Hours,GPU Hours,Node Hours,CPU Cost,GPU Cost,Storage (TB),Storage Cost,Backup (TB),Ba...
arochab1,NO_IO,10,3532.70,0.00,180.60,$176.63,$0.00,0.00,$0.00,0.00,$0.00,$176.63
bblanch6,NO_IO,28,5404.41,0.00,289.80,$270.22,$0.00,0.00,$0.00,0.00,$0.00,$270.22
```

Use case: IO-based billing analysis, cost allocation by infrastructure component.

4. `billing_YYYY-MM_report.pdf` Formatted PDF (PI+IO aggregate)

Landscape Letter PDF with org header, summary line (Jobs / CPU hr / GPU hr / Total), and one row per PI+IO with Jobs, CPU Hours, GPU Hours, and Total. Generated by the same **generate_billing_report** command (and by the Compute Billing button on `/billing/report/`). Served via `/billing/report/download/pdf/?month=YYYY-MM`.

Use case: Static archival, email attachment, finance hand-off same layout as the dynamic PDF button on the Billing Report page but pre-generated.

Generation Commands:

```
# Generate all three CSVs + PDF for a specific month
python manage.py generate_billing_report --month 2026-03 --exclude-admin-pi

# With custom rates
python manage.py generate_billing_report --month 2026-03 \
  --cpu-rate 0.05 \
  --gpu-rate 0.50 \
  --storage-rate 10.0 \
  --backup-rate 5.0

# Generate text invoices (one TXT file per allocation)
python manage.py generate_text_invoices --month 2026-03

# Generate text invoices AND convert to PDF
python manage.py generate_text_invoices --month 2026-03 --to-pdf
```

Customization

To adjust data generation, edit `coldfront/templates/generate_large_usage.py`:

```
# Change allocation type distribution
allocation_types = [
    {'name': 'gpu_intensive', 'cpu': (40000, 100000), 'gpu': (20000, 60000), ...},
    {'name': 'cpu_only', 'cpu': (20000, 50000), 'gpu': (0, 0), ...},
    # Add more types...
]

# Change date ranges
for month_offset in range(6): # 6 months -> change to 12 for 1 year
```

Then regenerate:

```
python3 coldfront/templates/generate_large_usage.py
```

Troubleshooting



"No usage data appearing in report?"

- Verify allocations are marked as billable (Allocation Detail -> Billing Info -> check box)
- Run: **docker compose -f docker-compose-dev.yml exec coldfront python manage.py shell**

```
from coldfront.plugins.arch_sync.models import SlurmAccountUsagePeriod
SlurmAccountUsagePeriod.objects.count() # should be > 1000
```

"Import says 'allocation not found'"

- Verify allocations were created by `populate_fake_billing_data`
- Check allocation IDs in CSV match database: **SELECT id FROM allocation LIMIT 5;**

"Charts/totals look wrong?"

- Verify rates exist and have correct effective dates
- Check that allocations have **billingis_billable=True**

10. Slurm CLI Utilities

These utilities are installed in the base Slurm image (**arch/cn-rockylinux**) at **/usr/local/bin/** and available on all nodes. All ARCH tools support **-M** for multi-cluster queries.

stree -- Account Hierarchy

```
# Show your own account tree (auto-detects $USER)
stree

# Show your tree with user associations
stree --full

# Show subtree (admin or own accounts only)
stree -r jhu

# Full hierarchy from root (admin only)
stree -r root --full
```

Note: Regular users can only view accounts they belong to. ``stree -r root`` requires admin privileges (root/...

Run **stree -r root** on a login node to see the live hierarchy. Example output:



```

root
|-- jhu    [Fairshare=20]
|   |-- wse    (school)
|   |   |-- cs    (department)
|   |   |   |-- rdesouz4    [100000]  <- PI default (QOS: scavenger)
|   |   |   |-- rdesouz4
|   |   |   |-- rdesouz4_proj1    [100000]  <- custom project (QOS: normal,premium)
|   |   |   |   |-- rdesouz4_proj1_A    [100000]  <- subproject
|   |   |   |   |-- rdesouz4, arochab1
|   |   |   |   |-- rdesouz4_proj1_B    [100000]
|   |   |   |   |-- rdesouz4
|   |   |   |-- arochab1    [100000]
|   |   |   |-- arochab1
|   |-- it    (school)
|   |-- arch    (department)
|   |-- bblanch6    [100000]
|   |-- bblanch6
+-- schmidt    [Fairshare=80]
    |-- ss    (Schmidt Sciences)
    |   |-- ssci-jdoe    [100000]  <- always free (ssci-* prefix)
    |   |-- ssci-jdoe
    |-- ssci-gomes    [100000]
    |-- ssci-gomes

[N] = Core Usage (Hours) / Fairshare weight
Use stree --full to include users

```

Troubleshooting:

- PIs flat under root (not under school/dept)? -> PI project has no Department set. Fix: assign a department in ColdFront Project Detail, then run **sync_slurm**
- Hierarchy not updating? -> Run: **docker exec coldfront coldfront sync_slurm**
- Need to reparent an account manually? -> **sacctmgr modify account where name=ACCT set parent=NEWPARENT -i** (requires **where** keyword)

interact -- Interactive Sessions

```

# List available partitions + your accounts & QOS
interact -l

# Start a session (1 core, 1 hour, scavenger)
interact -p scavenger -n 1

# 4 cores, 16GB, premium with specific account
interact -p premium -n 4 -m 16G -a mypi_proj1

# GPU session
interact -p gpu -g 1 -n 4 -m 32G -t 2:00:00

```

Built-in checks: validates walltime against partition limit, auto-detects GPU partitions, shows weekly cap budget remaining before submitting.

sbalance -- Usage Report



```
# Show usage for all your accounts
sbalance

# Show weekly spend cap status (cap vs usage in $)
sbalance --week

# Show usage for a specific PI or user
sbalance -p rdesouz4
sbalance -u rdesouz4

# Hide accounts with zero usage
sbalance --hide-zero
```

Troubleshooting:

- All zeros? -> No jobs have run yet, or **collect_sacct** hasn't imported them
- Missing accounts? -> Run **sync_slurm** to ensure ColdFront allocations are synced to Slurm

sqos -- QOS Definitions

```
# Show all QOS definitions
sqos all

# Show a specific QOS
sqos name=normal

# Show QOS for a specific account or user
sqos account=rdesouz4
sqos rdesouz4
```

Old name `slurm qos` still works as a symlink.

Other Tools

Command	Purpose
`sqme`	Show your active jobs (formatted `squeue`)
`seff <jobid>`	Show CPU and memory efficiency for a completed job
`sfeatures`	Show node features, GRES, and partition info
`sbrief`	Job state summary (count by Running, Pending, etc.)
`sassoc`	Show your associations (account, QOS, fairshare)
`sassoc --clusters`	List all registered Slurm clusters
`sadmin priority <jobid>`	Set high priority (admin only)
`sadmin addtime <jobid> <time>`	Add time to running job (admin only)

sacctmgr -- Common Debugging Commands



```
# List all associations for an account
sacctmgr -n -P list assoc where account=ACCOUNT format=Cluster,Account,User,Partition,QOS

# List all accounts with parents (use ParentID/ID -- ParentName is unreliable)
sacctmgr -n -P show assoc where user= format=Account,ParentName,ParentID,ID,Share

# List accounts under a specific parent
sacctmgr -n list account where parent=schmidt format=Account,Description

# Show account details
sacctmgr show account name=rdesouz4 withassoc format=Account,User,QOS,GrpTRESMins

# Reparent an account (requires 'where' keyword!)
sacctmgr modify account where name=rdesouz4 set parent=nodep -i
```

Note: In Slurm 25.11, `ParentName` is reported empty for every account whose parent is not `root` -- even when `ParentID` is correct (e.g. `noschool->jhu` shows `ParentName=''` but `ParentID = jhu's ID`). The hierarchy is still intact and the scheduler enforces it; only the display string is blank. Build/verify trees from `ParentID->ID`, not `ParentName` (this is what `stree`'s `build_tree` does). Do not "heal" empty `ParentName` by reparenting in sync -- it churns a modify per account every run. `sacctmgr list account format=ParentName` is likewise empty; use `show assoc where user=`.

start.sh slurm -- Cluster Management

Option A: CLI

```
./start.sh slurm create          # Generate compose overlay + DB record
./start.sh slurm start          # Build images (if needed) + compose up
./start.sh slurm status         # Lists all clusters from DB (container + bare-metal), then Docker ps +...
./start.sh slurm stop           # Stop Slurm services
./start.sh slurm restart        # Restart (slurmdbd -> slurmctld -> nodes)
./start.sh slurm destroy        # Stop + remove containers, volumes, overlay
./start.sh slurm test-jobs      # Seed real sbatch jobs for every userxaccountxpartition
./start.sh slurm ansible [name|all] # Regenerate ansible/clusters/<name>/ bundle from DB
./start.sh slurm prod-sim [name]  # Bare-metal simulation: ansible-playbook --check --diff against per-cl...
```

Option B: UI (Admin -> Clusters)

1. Add Cluster [register cluster with name, type, and deployment model preset](#)
2. Build Images [compile Slurm base + service images](#)
3. Deploy [generate slurm.conf + compose overlay + start containers](#)
4. Stop / Restart [manage running cluster from detail page](#)

Design: Docker exec vs local binaries

All Slurm CLI commands (**sacctmgr**, **sinfo**, **sreport**, etc.) run via **docker exec {cluster}-slurmctld** rather than local binaries:

	Docker exec	Local binaries + volumes
Simplicity	1 function, no extra volumes	3 init containers + 2 volumes
Multi-cluster	Routes to correct slurmctld	Local binary can't resolve cluster hostname
Dependencies	Zero -- only Docker SDK	Needs munge key, slurm libs, Idconfig



Maintenance	None	Slurm upgrade requires rebuild of init container
-----------------	------	--

Bare-metal (prod) uses local binaries directly handled by the **else** branch in **slurm_command()**.

Troubleshooting:

- Nodes showing "down" after restart? -> Wait 10 seconds (auto-resume runs in background). If still down: **scontrol update nodename=cn-01 state=resume**
- "No docker-compose-slurm-*.yaml found"? -> Deploy the cluster from the ColdFront UI first (Clusters -> Build Images -> Deploy)
- Init containers showing "Exited (0)" in red? -> Normal. **munge-init** is a one-shot container that generates the shared munge key. It consumes zero CPU/memory after exit.

Testing

Commands

```
# Run all tests (108 tests)
./start.sh test

# Full stack integration test (16 phases: containers, LDAP, Slurm, jobs, billing, QOS, sync)
./start.sh test deploy

# List available test modules with counts
./start.sh test list

# Show test help with examples
./start.sh test --help

# Run a specific test module
./start.sh test coldfront.plugins.arch_sync.tests.test_billing

# Run a specific test class
./start.sh test coldfront.plugins.arch_sync.tests.test_models.CoreHoursResetSettingsTest

# Run a specific test method
./start.sh test coldfront.plugins.arch_sync.tests.test_signals.DepartmentSignalTest.test_billing_dept_propa...
```

Available Test Modules

Module	Type	What It Tests
`arch_sync.tests.test_billing`	SimpleTestCase	`dollars_to_grptres_mins`, `grptres_mins_to_dollars`, `hours_to_dollars`, `SYSTEM_USERNAMES`
`arch_sync.tests.test_models`	TestCase	`CoreHoursResetSettings.compute_next_reset()` (6 periods), `PaymentStatus` ordering
`arch_sync.tests.test_slurm_sync`	TestCase	`SlurmNode/SlurmPartition.slurm_conf_line` (AllowQos, TRESBillingWeights), weekly cap conversion
`arch_sync.tests.test_signals`	TestCase	Department propagation signals, AllocationBilling auto-creation
`arch_sync.tests.test_ldap`	SimpleTestCase	DN parsing, SSHA hash, tree config, UID/GID allocation, email comparison
`core.user.tests`	TestCase	User registration, PI request flow

Full module path prefix: **coldfront.plugins.** (for arch_sync) or **coldfront.** (for core).

Example Output



```
$ ./start.sh test
[+] Running Django tests: coldfront.plugins.arch_sync.tests

v Applying contenttypes.0001_initial
v Applying auth.0001_initial
...
v test_decimal (coldfront.plugins.arch_sync.tests.test_billing.DollarsToGrpTRESminsTest)
v test_one_dollar (coldfront.plugins.arch_sync.tests.test_billing.DollarsToGrpTRESminsTest)
v test_dept_propagates_to_default_project (...) -- Setting billing dept propagates to Project.department.
v test_simple_dn (coldfront.plugins.arch_sync.tests.test_ldap.UidFromDnTest)
...

-----
Ran 108 tests in 0.729s

OK
```

Note: **SimpleTestCase** tests (test_billing, test_ldap) run without a database. **TestCase** tests require CREATEDB permission (granted automatically by `./start.sh test`).

User Guide

Regular Users

Users join existing projects and use allocations managed by PIs.

Action	How
View your projects	Projects (sidebar)
View your allocations	Click a project, see Allocations section
Join a project	PI adds you via Add Users on the project page
Check Slurm account	Allocation detail shows `slurm_account_name`
Submit a job (web)	Jobs > Submit -- choose partition + account
View job history	Jobs > Metrics -- filter by your username
Request more hours	Allocation detail > Request Change

Principal Investigators (PIs)

PIs own projects, manage members, and control allocations.

Action	How
Create a project	Projects > Create Project -- set title, abbreviation, department, IO Number
Request storage/backup	Check Storage/Backup boxes on project create/update form
Add users	Project detail > Add Users -- search by username, select allocations
Set weekly spend cap	Allocation detail > Set Weekly Cap (\$)
View billing	My Usage & Billing (sidebar)
Request allocation change	Allocation detail > Request Change (e.g. more Core Hours)

Staff Tools

Tool	Purpose
Manage Users	Toggle PI status, assign cluster/affiliation scope, impersonate users



Assign Staff Scope	Set cluster scope (manage clusters) and affiliation scope (manage users/PIs)
Staff User Profile	Edit name/email, set LDAP password, toggle shell access (`/bin/bash` vs `/bin/false`)
Promote to PI	Creates default project + allocation + Slurm account automatically
Elevated approvals	Grant via Admin > User profiles > "Elevate to superuser for approvals" -- allows scoped staff to approve/deny without full superuser
Staff Help	In-app reference at /staff/help/ with downloadable PDF
Service Health	Real-time status of all services at /user/service-health/

Delegated Staff Administration

Staff members can be given two independent scopes via the "Assign Staff Scope" page:

Scope	What It Controls
Cluster Scope	View cluster status, drain/resume nodes for assigned clusters
Affiliation Scope	See/manage users, PIs, projects, allocations of assigned affiliations

Multiple clusters and affiliations can be assigned (checkboxes). Superusers always see everything. Scoped staff see a simplified Manage Users page with only their affiliation's users. Deploy/stop/restart operations remain superuser-only.

Best Practices

- Keep projects tidy [archive inactive projects](#)
- Use descriptive abbreviations [they become Slurm account names](#)
- Plan ahead for renewals [submit extension requests before allocations expire](#)
- Set weekly spend caps on paid projects to prevent overspending

Multi-Cluster Architecture

CLOACK supports two deployment models for scaling beyond a single cluster. Both use the same ColdFront codebase. [the **SlurmCluster** model and **-M** flag routing are already implemented.](#)

Docker Multi-Cluster: Each cluster gets its own prefixed containers (**{name}-slurmctld**, **{name}-cn-01**) and volumes (**{name}_slurm_state**). Shared services (**slurmdbd**, **slurmrestd**, **slurm-db**) run as single instances supporting federation. Multiple clusters can run simultaneously.

Topology matrix (**cluster_type** x **core_services**): **SlurmCluster** carries two independent enums that together describe all four federation layouts. **cluster_type** says where **slurmctld** + compute nodes live; **core_services** says where **slurmdbd** + **slurmrestd** live. For dedicated-core clusters, **slurmdbd_host** (host:port) is used by the Service Health TCP probe.

#	cluster_type	core_services	Scenario
A	container	shared_container	Dev docker default
B	bare_metal	dedicated	Prod bare-metal (own slurmdbd/slurmrestd)
C	container	dedicated	Docker compute + external DBD
D	bare_metal	shared_container	Hybrid: bare-metal compute + shared Docker core


The Admin > Slurm Clusters page ships a Hybrid preset that pre-fills topology D (**bare_metal** + **shared_container**,



slurmrestd_url=http://slurmrestd:6820). Existing records migrated via **0086_slurmcluster_core_services** default to **shared_container** for container clusters and **dedicated** for bare-metal.

Model A -- Single Cluster (Recommended for Start)

One slurmctld manages all nodes. Different hardware is separated by partitions.



Benefit	Detail
Simple operations	One slurmctld to manage, monitor, and upgrade
Unified scheduling	Scheduler sees all resources, optimal placement
Easy configuration	One slurm.conf, one set of partitions
Scales to ~10K nodes	Single slurmctld handles 1000 nodes easily

Drawback	Detail
Single point of failure	slurmctld crash affects all users
Maintenance downtime	Upgrade/restart impacts all jobs
Mixed policies	All hardware shares the same scheduling policies

Model B -- Federation (When You Need Isolation)

Multiple slurmctld instances sharing one slurmdbd. Each cluster has its own partitions and scheduling policies. Users submit cross-cluster with **sbatch -M cluster_name**.



```

+-----+
| slurmdbd | <- shared accounting DB
| slurmrestd | <- single REST API for all clusters
L-----+-----+
      +-----+-----+
      ?           ?           ?
+-----+ +-----+ +-----+
|slurmctld| |slurmctld| |slurmctld|
|"arch-cpu"| |"arch-gpu"| |"arch-dev"|
L-----+-----+ L-----+-----+ L-----+-----+
      ?           ?           ?
+-----+ +-----+ +-----+
|600 CPU | | 200 GPU | | 50 dev |
| nodes  | | nodes  | | nodes  |
L-----+-----+ L-----+-----+ L-----+-----+
      ?           ?           ?
      L-----+-----+
              +-----+-----+
              | login01 | <- shared login node
              L-----+-----+

```

Benefit	Detail
Fault isolation	arch-gpu crash does not affect arch-cpu jobs
Independent maintenance	Upgrade GPU cluster without touching CPU cluster
Optimized scheduling	Each cluster tuned for its hardware (GPU policies differ from CPU)
Dev isolation	Users test on dev cluster without impacting production
Scalable beyond 10K nodes	Each cluster handles up to ~10K nodes independently
Unified billing	Shared slurmdbd = single accounting database across all clusters
Cross-cluster submission	`sbatch -M arch-gpu` from any login node

Drawback	Detail
Operational complexity	Multiple slurmctld instances to manage
Configuration overhead	Separate slurm.conf per cluster
Resource fragmentation	Scheduler in arch-cpu cannot use idle arch-gpu nodes

Comparison

Aspect	Single Cluster	Federation
Scheduling	1 scheduler for all jobs	Dedicated scheduler per cluster
Scalability	~10K nodes per slurmctld	~10K nodes **per cluster**
Failure blast radius	All users affected	Only users of that cluster
Maintenance	Downtime affects everyone	Upgrade one cluster at a time
Partitions	Shared namespace	Isolated per cluster (no name conflicts)
Hardware	Mixed or homogeneous	Dedicated per cluster type
Billing	1 accounting DB	1 accounting DB (shared slurmdbd)
Cross-cluster jobs	N/A	`sbatch -M cluster_name` from shared login
CLI tools	`sacct`, `sbalance`, `stree`, `sqos`	Same commands with `-M` flag

ColdFront Integration

The **SlurmCluster** model supports both architectures. The **root_account** field drives which deployment model is used. no code changes needed:

`root_account`	Deployment Model	Affiliations	Example
`root` / `arch` / other	Model A/C	Both JHU + Schmidt	Single cluster or hardware split



`jhu`	Model B	JHU only	Federation -- JHU-dedicated cluster
`schmidt`	Model B	Schmidt only	Federation -- Schmidt-dedicated cluster

```
# Model A: single cluster (both affiliations)
SlurmCluster(name="arch", root_account="root", slurmrestd_url="http://slurmrestd:6820")

# Model B: federation -- one cluster per affiliation
SlurmCluster(name="jhu-cluster", root_account="jhu", slurmrestd_url="http://slurmrestd:6820")
SlurmCluster(name="schmidt-cluster", root_account="schmidt", slurmrestd_url="http://slurmrestd:6820")

# Model C: hardware split (both affiliations per cluster)
SlurmCluster(name="arch-cpu", root_account="arch", slurmrestd_url="http://slurmrestd:6820")
SlurmCluster(name="arch-gpu", root_account="arch", slurmrestd_url="http://restd-2:6820")
```

The **slurm_sync** respects **root_account** automatically:

- Model A/C (**root_account** != jhu/schmidt): creates both **jhu** and **schmidt** root accounts under **cluster_root**, syncs all allocations
- Model B (**root_account** = jhu): creates only **jhu** root, filters out Schmidt PI allocations
- Model B (**root_account** = schmidt): creates only **schmidt** root, filters out JHU PI allocations

Admin presets in the Django Admin pre-fill the correct **root_account** for each deployment model.

Real-time Account Provisioning

When an allocation transitions to Active (staff approval), the system immediately creates the Slurm account hierarchy via REST API no waiting for the periodic sync. This is best-effort: if the REST API is unavailable, the periodic sync (every 3 hours) acts as a safety net.

The provisioning respects Model B affiliation filtering: a JHU PI allocation approved on a Schmidt-only cluster is silently skipped.

Recommendation for JHU (1000 nodes, 10K users): Start with Model A (single cluster with partitions). Migrate to Model B when hardware diversity or operational isolation justifies the complexity. The migration is infrastructure-only add another slurmctld, create a new **SlurmCluster** in Django Admin with the appropriate **root_account**, no code changes needed.

Monitoring

External Grafana + Prometheus at <https://dash.arch.jhu.edu/> scrape our hosts in pull mode over the JHU LAN. Exporters and scrape configs live under **monitoring/**:

- **monitoring/README.md** entry point, daemon x layout matrix
- **monitoring/catalog.md** canonical exporter inventory (image, port, DSN shape)
- **monitoring/docker/** sidecars for the Docker side (Layout A + D, and the **cf-vm** of B/C)
- **monitoring/ansible/roles/monitoring/** bare-metal exporters (Layout B + C)
- **monitoring/scrape_configs/layout_{a,b,c,d}.yml** per-layout target snippets to hand to the **dash** operator

Covered daemons: coldfront + qcluster + helpdesk (django-prometheus + custom view), postgres (:5432), mariadb



(Slurm acct DB on the non-standard **:3307**), openldap, keycloak (built-in **/metrics**), slurmctld/slurmdbd/slurmrestd (single exporter via REST + rotating JWT), munge (textfile collector), cli_filter (login nodes only), plus node_exporter + cadvisor everywhere. Compute nodes (**cn-XX**) are intentionally out of scope.

The **slurm_exporter** reuses the rotating token that **start.sh** already keeps fresh at **\${BASE_ETC}/slurm/slurm_jwt.token** (cron 03:00 daily, lifespan 30 days) no manual token generation. See [memory project monitoring dash](#) for the full handoff details.



Features

CLOACK = ColdFront + LDAP + OIDC + ARCH + Cluster/Containers + Keycloak

Containerized identity and resource management for HPC clusters at Johns Hopkins University.

Branch Overview

Branch	Version	Purpose	Commits ahead of main
`main`	v0.1	Baseline -- upstream ColdFront + OIDC + LDAP scaffolding	--
`prod`	v0.2	Current production -- LDAP/Slurm sync, staff views, basic signals	69
`dev`	v0.3	Development -- full billing, weekly cap, cluster management, tests	322

Repository: https://github.com/jhu-arch/arch_jhu_schmidt

Documentation Index

File	Description
README.md	ARCH Portal -- stack overview, components, architecture
CLOACK.md	Feature comparison across branches (main vs prod vs dev)
ColdFront	
coldfront/README.md	ColdFront -- HPC resource management portal
coldfront/custom/README.md	Custom configurations (ARCH JHU DSAI Schmidt)
coldfront/BILLING_SETUP_GUIDE.md	Billing system setup guide
coldfront/BILLING_DATA_IMPORT.md	Billing system data import guide
coldfront/BILLING_AND_STORAGE_CHANGES.md	Billing info & storage/backup implementation summary
coldfront/INVOICE_SETTINGS.md	Invoice customization settings
coldfront/SCHOOL_DEPARTMENT_TRACKING.md	School & department tracking for job usage
coldfront/RECOLLECT_HISTORICAL_JOBS.md	Re-collect historical jobs with school & department
coldfront/init/README.md	Initialization scripts (seed data, LDAP import)
Slurm	
slurm/README.md	Slurm -- job scheduler stack, cost enforcement, CLI utilities
slurm/SLURM_RESTART_GUIDE.md	Slurm service restart guide (correct order)
Infrastructure	
openldap/README.md	OpenLDAP -- user & group directory
keycloak/README.md	Keycloak -- identity & access management (OIDC)
keycloak/ldap-providers/README.md	Keycloak LDAP provider snippets and import instructions
postgres/README.md	PostgreSQL -- database server
base/README.md	Base -- shared configuration (UID/GID ranges, paths)
phpldapadmin/README.md	phpLDAPadmin -- LDAP web UI
ondemand/README.md	Open OnDemand -- web portal for HPC
scripts/README.md	Utility scripts (wait-for.sh, etc.)



Feature Comparison: main vs prod vs dev deployment

Billing & Cost Management

Feature	main	prod	dev	Notes
BillingRate catalog (6 rate types)	--	--	Yes	cpu_full, cpu_om, gpu_full, gpu_om, storage, backup
AllocationBilling per allocation	--	--	Yes	IO number, department, billing_type (Pay Per Use / Credit Holder / No Charge)
HardwareCreditPool	--	--	Yes	Track hardware purchases converted to compute credits; credits offset compute, not O&M
SlurmAccountUsagePeriod	--	--	Yes	Per-account per-period usage -- billing source of truth
PaymentStatus (configurable)	--	--	Yes	Admin-managed status codes (pending, paid, overdue, no_charge) without code changes
InvoiceSettings (singleton)	--	--	Yes	Customizable invoice branding, credit/O&M explanation text
Billing Report (staff view)	--	--	Yes	Filterable by PI, project, department, school, month, billing type
PI Billing Summary	--	--	Yes	PI-facing "My Usage & Billing" dashboard
Rate Management UI	--	--	Yes	CRUD for BillingRate with effective date versioning
Hardware Credits UI	--	--	Yes	CRUD for HardwareCreditPool with credit tracking
Unit conversion helpers	--	--	Yes	`dollars_to_grptres_mins()`, `grptres_mins_to_dollars()`, `hours_to_dollars()`
collect_billing command	--	--	Yes	Populate SlurmAccountUsagePeriod from sreport + sacct

Weekly Spend Cap

Feature	main	prod	dev	Notes
Dollar-based weekly cap	--	--	Yes	AllocationAttribute `Weekly Cap (\$)`
Account-level GrpTRESMins	--	--	Yes	Cap applied directly on Slurm account (not per-QOS)
Weekly cap usage (live)	--	--	Yes	On-demand via `GET /api/v1/allocations/<pk>/weekly-usage/`; AJAX gauge; debounced AllocationAttribute UPSERT
Cap-to-TRES conversion	--	--	Yes	`\$30/week` -> `GrpTRESMins=billing=1800000` (formula: `\$ * 1000 * 60`)
Weekly reset via cron	--	--	Yes	`check_and_run_core_hours_reset` resets GrpTRESMins periodically

How it works:

Project	Cap	Slurm GrpTRESMins
`arochab1_f01_a`	\$30/week	`billing=1800000`
`arochab1_f01_b`	\$20/week	`billing=1200000`
`bgermro1`	No cap	(not set = unlimited)

Instead of creating per-allocation QOS (**wk_{id}**), the cap is enforced at the account association level:

```
sacctmgr modify account where name=arochab1_f01_a set GrpTRESMins=billing=1800000
```

Storage & Backup

Feature	main	prod	dev	Notes
---------	------	------	-----	-------



Project storage/backup fields	--	--	Yes	`storage_enabled`, `storage_tb`, `backup_enabled`, `backup_tb` on Project model
Storage/backup on project	--	--	Yes	Create + Update forms with validation
AllocationBilling storage/backup	--	--	Yes	`storage_tb`, `backup_tb` fields for billing
Storage/backup billing rates	--	--	Yes	BillingRate types: `storage` (TB/month), `backup` (TB/month)

Core Hours Reset

Feature	main	prod	dev	Notes
CoreHoursResetSettings	--	--	Yes	Singleton model with 6 period options
Period options	--	--	Yes	weekly, monthly, quarterly, 4-month, semester, annual
`compute_next_reset()`	--	--	Yes	Auto-calculates next reset date based on period
Per-allocation override	--	--	Yes	AllocationAttribute `Next Reset` overrides global schedule
Reset management UI	--	--	Yes	Staff page to configure schedule and trigger manual resets
Next Reset on allocation gauge	--	--	Yes	Shows next reset date below Core Usage gauge

Slurm Integration

Feature	main	prod	dev	Notes
Basic slurm_sync (account + users)	--	Yes	Yes	Create Slurm accounts, add users to accounts
SlurmUserUsageSnapshot	--	Yes	Yes	Per-user usage snapshots from sacct
SlurmCluster / SlurmNode / SlurmPartition models	--	--	Yes	Full cluster topology in Django
Dynamic slurm.conf generation	--	--	Yes	NodeName/PartitionName generated from UI models
TRES billing weights per partition	--	--	Yes	`TRESBillingWeights` field on SlurmPartition
AllowQos per partition	--	--	Yes	`allow_qos` field generates `AllowQos=...` in slurm.conf
SlurmQOS model	--	--	Yes	QOS definitions per cluster; replaces hardcoded bootstrap
QOS (via) location	--	--	Yes	AllocationBilling.qos FK overrides default QOS assignment
export_qos_config	--	--	Yes	Generates `qos_config.lua` for cli_filter.lua from DB
Slurm account hierarchy	--	Flat	Yes	`root -> school -> dept -> PI -> project` (auto-managed)
Orphan account cleanup	--	--	Yes	`sync_slurm` removes empty intermediate school/dept accounts
Slurm REST API client	--	--	Yes	`slurm_client.py` for slurmrestd integration
Docker-based cluster management	--	--	Yes	`cluster_manager.py` -- build/deploy/restart/stop via UI
sacct collection (every 15 min)	--	--	Yes	SlurmJob model + `collect_sacct` command
scontrol reconfigure via UI	--	--	Yes	"Sync Accounts" button triggers `scontrol reconfigure`
Defensive scontrol reconfigure post-sync	--	--	Yes	`_refresh_ctld_cache` runs `scontrol reconfigure` at the end of each `_run_sync_for_cluster` -- prevents stale assoc cache rejecting submissions
TRES types configuration	--	--	Yes	`SlurmCluster.tres_types` -- configurable per cluster
ClusterReservation model	--	--	Yes	PI-requested node reservation; status pending/active/expired/denied/cancelled; admin-only (CRUD UI pending)
JobTemplate model	--	--	Yes	Reusable sbatch script with `{placeholders}`; per-owner uniqueness; `is_public` shares with lab
DeploymentLog model	--	--	Yes	Audit trail for deploy/stop/restart/drain/resume/reserve/ansible events
`slurm prod-sim` CLI	--	--	Yes	`./start.sh slurm prod-sim [name]` runs ansible-playbook `--check --diff` against per-cluster inventory (read-only pre-flight)



`slurm status` from DB	--	--	Yes	Lists all clusters (container + bare-metal) from SlurmCluster with type/root/nodes/parts/REST
------------------------	----	----	-----	---

Department & School Hierarchy

Feature	main	prod	dev	Notes
School model	--	--	Yes	Extracted from Department; FK relationship
Department with school linkage	--	--	Yes	Department -> School -> Slurm hierarchy
AllocationBilling as hierarchy source	--	--	Yes	Billing dept drives Slurm account placement, not Project.department
Department propagation signals	--	--	Yes	Billing dept auto-propagates to Project + PI UserProfile
CSV department import	--	--	Yes	Staff UI for bulk department import
Department user membership	--	--	Yes	FilteredSelectMultiple on Department admin + auto-propagation
Shares on School admin	--	--	Yes	FairShare weights configurable per school

User Management

Feature	main	prod	dev	Notes
Staff scoped views	--	Partial	Yes	prod: basic scoping; dev: full resource-based scoping
AssignStaffResourcesView	--	--	Yes	Superuser assigns cluster resources to staff
CSV user import	--	--	Yes	Bulk user creation with LDAP provisioning
PI SubProject scheme	--	--	Yes	Numeric (A, B, C) or alpha sequence per PI
Activate project (staff)	--	--	Yes	Staff can activate New-status projects
Manage Users view	--	--	Yes	Staff user management across projects
Extra Core Hours attribute	--	--	Yes	Additive core hours on top of base allocation
System account exclusions	--	--	Yes	admin, coldfront, root excluded from all views/sync/billing

Authentication & Identity

Feature	main	prod	dev	Notes
Multi-realm Keycloak	Partial	Yes	Yes	JhuOIDCBackend + SchmidtOIDCBackend with separate realms
LDAP sync (daily + on-dem)	Partial	Yes	Yes	prod: basic sync; dev: full bidirectional with home dirs
Auto-user creation on OIDC	--	Yes	Yes	Signal-based user provisioning with auto-affiliation detection
Directory ground truth (UID/GID)	--	--	Yes	Centralized UID/GID allocation, cached on UserProfile, visible in Admin
LDAP migration import	--	--	Yes	Import slapcat dumps preserving UIDs/GIDs/passwords via Import Users UI
PI promotion workflow	Partial	Partial	Yes	dev: enhanced with department propagation



Entra ID integration	--	--	Yes	`JHU_ENTRA_ENABLED` flag for Azure AD federation
Olds-a pp SSO (ColdFront <-> Helpdes	--	--	Yes	Sidebar links route to correct realm; `home_or_sso` view + realm-skip guard in OIDC backends. Edge/subdomain mode: helpdesk auto-SSO entry (`home_or_sso`/`LOGIN_URL` -> host-aware `realm_dispatch_login`) and the portal ARCH Helpdesk card are per-affiliation (`helpdesk.<base>` -> matching realm), not JHU-hardcoded
Portal landing	--	--	Yes	`/` redirects anonymous users to `/portal/` (JHU + Schmidt + ARCH Helpdesk cards); legacy `/user/login/` picker kept for tests; logout lands on `/`
page usernam e normaliz ation	--	--	Yes	Both backends strip `@domain` from claims and dedupe JHU alias domains (`jhu.edu` <-> `jhmi.edu` <-> `jh.edu`); `normalize_oidc_usernames` Django command collapses legacy `user@domain` rows on every `start.sh helpdesk start`
(JHU+ bates k) -> Keycloa k master	--	--	Yes	OIDC brokering from JHU realm into master; composite role `limited-admin` (45 sub-roles across master/jhu/schmidt `*-realm` clients, no super-admin/create-realm/impersonation); `admin-only-gate` IdP mapper gates on inbound JHU `realm_access.roles` containing `admin`; non-admin JHU users get zero master roles -> Forbidden; unconditional amber policy banner on the admin console
SSHole OTP via Keycloa	--	--	Yes	Per-login-node `otp_mode` (disabled/mixed/device_only/ropc_only); `mixed`: ssci-*/ext-*/JHU admins (cn=admin group) via ROPC terminal (password + OTP), JHED via Device Flow browser; 2-prompt PAM via `pam_kc_ssh.so`; Keycloak Direct Grant + OTP Form automated in `configure-keycloak.sh`
Fake GPU support	--	--	Yes	`SlurmNode.gres` -> `gres.conf` + `AccountingStorageTRES` + `TRESBillingWeights`
(pre, en v secrets	--	--	Yes	Gitignored file auto-sourced by `start.sh init` -- persists Azure secrets across `clean-all`
pattern v + Trello	--	--	Yes	Read-only Kanban/Cards/Lists/Gantt at `/extensions/trello/` (Frappe Gantt MIT); workspace/board discovery; `TrelloMemberMap` links Trello handles to helpdesk users for fairness routing
Browser2 ticket UI	--	--	Yes	2-col layout, urgency banner (SLA/FQR/strike/overdue), summary bar, context sidebar with past tickets by user+resource, suggested collaborators, similar open tickets, skill match
Note categori zation + email	--	--	Yes	`FollowUpMeta`: note kind (tech/RCA/internal/decision/customer_reply) + email class (FQR/strike 1-3/LQR); picking a classification pastes `ResponseTemplate.body` with `{name}`/`{your_name}`/`{issue}`/`{ticket_number}`/`{survey_url}`/`{signature}` already substituted (pre_save scrub safety net catches bypassed tokens); advances `TicketFollowUpTracker`
classification AI triage + pre-diag nosis	--	--	Yes	`ai-engine` container (llama.cpp + qwen3-4b-instruct, internal network only) classifies every new ticket (`TriageResult` badge + staff Apply) and analyzes the guided pre-diagnosis form (`hd_prediag`); provider registry in admin (`AIProvider`/`AIEngineConfig`, keys stay in env); staff dashboard at `/extensions/ai/`; user data never leaves infra, `CLAUDE_API_KEY` optional -- see helpdesk/AI_ARCHITECTURE.md
Deletion approval queue	--	--	Yes	`TicketDeletionRequest`: engineer must justify (min 10 chars); superuser-only queue at `/extensions/deletion-requests/` approves or rejects
Queue pool	--	--	Yes	`QueuePool(queue, members)`: Owner dropdown filtered to pool members; permission checks allow pool members to work tickets
Collabor ator permissi ons	--	--	Yes	Non-owner/non-collab/non-pool staff blocked from adding comments; collaborators forced `public=False` (internal notes only); owners send customer emails; permission failures on collaborator request/approve/deny/remove show dismissible banners via `messages.error` + redirect to ticket (no more dead-end HTTP 403 pages)



Public feedback categories +	--	--	Yes	`TicketCategory.show_on_submit` renders opt-in "Feature feedback?" dropdown on the public submit form; `TicketCategory.default_assignee` auto-routes the ticket to the product owner. `TicketCategory.default_queue` (FK) moves the ticket to a dedicated Feedback & Suggestions queue (slug `feedback`, owner=`ARCH_SUPERUSERS[0]`) regardless of the queue the submitter picked. Seeded by migration 0029; ColdFront + Helpdesk feedback categories both routed to the feedback queue.
Dedicated sticky save-at-top	--	--	Yes	Sticky save bar injected via `{% block form_top %}` override -- no scroll on long change forms
Ticket Analytics -- dynamic Resource combo	--	--	Yes	Three tabs on `/dashboard/`: Department/Project , PI , Resource . Resource tab carries a combo (GPU/CPU/Storage/Backup) populated dynamically from ColdFront `GET /api/v1/resources/catalog/` -- GPU items from `SlurmPartition.hardware` distinct, CPU items from partition names, Storage/Backup from `AllocationBilling`. Per-item columns + the matching usage metric (`gpu_hours`/`cpu_hours`/`storage_tb`/`backup_tb`) merged from `GET /api/v1/stats/slurm/?metric=...`. Graceful degradation to a GPU-only view when ColdFront is offline.
Resource + PI/Project auto-link	--	--	Yes	Two passes on `new_ticket_done`: (a) regex scan of title+description against the dynamic ColdFront catalog creates `TicketResourceLink(kind='resource')` rows (GPU unprefixed, others namespaced like `cpu:high`); (b) submitter-email lookup against ColdFront `/api/v1/allocations/by_submitter/` creates `kind='pi'`, `kind='project'`, `kind='allocation'` (department) links so the PI / Department tabs render automatically. Backfill command: `python manage.py backfill_ticket_links [--window-days N]`.
Engineer tickets Schedule Grid + Time	--	--	Yes	`/extensions/schedule/`: per-engineer x day x 30-min-slot grid. Time Allocations (model `StaffBlockedPeriod`) group unavailability (holiday/vacation/sick/OOF) and committed work (project/tickets/learning/on-call) with color-coded cells. Admin drag-select + change-request flow for non-admin staff; denied requests surface red dashed cells with reason on hover. "Next N days" summary panel with configurable window (3/7/14).
Request Assignment Request follow-u	--	--	Yes	Staff submit weekly allocation requests at `/extensions/schedule/request/`. /extensions/schedule/my-requests/ lists the staff's own history with status badges and links to generated blocks. /extensions/schedule/admin/requests/ (admin) groups all requests by `StaffMember` with follow-up flags for pending > 7 days or approved-without-block. Grid cells generated from an approved request carry a blue inset border + tooltip linking back to the origin request.
Password change	--	--	Yes	`HELPDESK_SHOW_CHANGE_PASSWORD=True` keeps the dropdown item; clicking it shows a banner page (no form) with a button to `\${COLDFRONT_URL}/user/user-profile/`. Password management stays in ColdFront/identity providers.
Helpdesk rollout (TOTP + Entra)	--	--	Yes	No `userPassword` written to LDAP for any non-system user. JHED (incl. cn=admin) authenticate via Entra Federation -> MFA from JHU; ext-*/ssci-* enrol TOTP in the Keycloak Account Console (Schmidt realm browser flow rewritten to Username Form -> OTP Form only, no password page). `CONFIGURE_TOTP` required action `enabled=true defaultAction=true` on both realms -> first login forces enrolment. ColdFront banner block in `base_black.html` + `base_light_sidebar.html` checks `UserProfile.totp_enrolled` (sync'd every 30 min by `sync_totp_status` qcluster task polling Keycloak admin REST). `/admin/` break-glass stays local-password + `django-axes` brute-force lockout + `AdminLoginRateLimitMiddleware` per-IP throttle + `admin_login_alert_task` email on every superuser login. Quarterly rotation automated by `scripts/rotate-admin-password.sh`. Merged to dev 2026-06-09 (commit `aef43b6`).

Scheduled Tasks (django-q)

Task	prod	dev	Schedule
`sync_ldap_all`	Yes	Yes	Daily 2 AM
`sync_slurm_all`	Yes	Yes	Daily 3 AM
`collect_sacct_task`	--	Yes	Every 15 min
`check_and_run_core_hours_reset`	--	Yes	Daily 1 AM
`cleanup_task_queue`	--	Yes	Hourly :30
~~`sync_weekly_cap_usage`~~	--	~~Removed~~	Removed (Phase 2) -- replaced by live endpoint



All tasks in dev use **hook=TASK_HOOK** for Django Admin LogEntry audit trail.

UI & Admin Improvements

Feature	main	prod	dev	Notes
Sidebar navigation (light theme)	--	--	Yes	Admin Tools, Staff Tools, Clusters, Billing Management
Staff Help page	--	--	Yes	`/staff/help/` -- 21 DB-driven accordion sections editable via Django Admin
Allocation detail enhancements	--	--	Yes	Weekly cap gauge, next reset date, billing context
Job metrics dashboard	--	--	Yes	`/jobs/dashboard/` with per-cluster/partition/user filtering
Cluster list + detail views	--	--	Yes	Full cluster management UI
Inline attribute editing	--	--	Yes	Edit Core Usage, Next Reset, Weekly Cap directly on allocation page
Allocation sync trigger	--	--	Yes	"Sync Now" button on allocation detail
Project creation with storage/backup	--	--	Yes	Storage + backup fields on create/update forms

Infrastructure & DevOps

Feature	main	prod	dev	Notes
`start.sh` orchestr	Basic	Basic	Enhanced	dev: init, start, restart, rebuild, populate, slurm, ldap-reinit, prune-*, test
start.sh test comman	--	--	Yes	Run Django unit tests inside container
Docker Slurm cluster	--	--	Yes	Full containerized Slurm: slurmctld, slurmd, slurmdbd, slurmrestd
(dev) tests (arch_sy	--	--	Yes	test_billing, test_models, test_slurm_sync, test_signals
SPANK plugins	--	--	Yes	Cost telemetry compiled into base image
Contain er jobs (Pyxis +	--	--	Yes	NVIDIA Pyxis SPANK in base image (`spank_pyxis.so` v0.20.0); NVIDIA enroot runtime on compute + login nodes (`arch/slurmd` v3.5.0). Enables `sbatch --container-image=...`. GPU support (`libnvidia-container-tools`) installed via Ansible on bare-metal GPU nodes only.
enforce.lua (cost estimates)	--	--	Yes	Client-side cost prompt at job submit (login node only)
enforce.lua (enforce)	--	--	Yes	Controller-side weekly cap + billing validation
enforce.lua (billing rates)	--	--	Yes	Per-partition CPU/GPU rate definitions for Lua scripts
Django- Q audit	--	--	Yes	Every async_task uses TASK_HOOK for LogEntry tracking
Free-bas ed email (dev)	--	--	Yes	Emails logged to `coldfront/report/emails/` instead of SMTP



Per-Queue routing	--	--	Yes	`QueueRoutingOverride` sidecar pins the starting tier per queue (e.g. `storage` -> `se`, `software` -> `sse`). Signal precedence: `TicketCategory.default_policy` > `Queue.routing_override` > `first_tier`. Admin at `/admin/helpdesk_extensions/queueroutingoverride/`.
Ticket affiliation	--	--	Yes	`TicketAffiliation` OneToOne sidecar populated on ticket creation by 4-tier detect: gateway queue (`ssci` -> `schmidt`, `skipjack` -> `jhu`) > authenticated OIDC user profile > email domain (`jhu.edu`, `schmidtsciences.org` etc.) > `unknown`. Admin filter on Ticket list; override maps via `HELPDESK_GATEWAY_QUEUE_AFFILIATION` / `HELPDESK_AFFILIATION_DOMAINS`
FQR/SLA (SLA) timers	--	--	Yes	SLA: `FQRMetrics.fqr_due_at` (3 biz hours) + `TicketFollowUpTracker.followup_due_at` (3 biz days, rolling on each classified follow-up). `business_time.py` helper skips Sat/Sun. Banner chips on ticket detail: SLA always visible, FQR halfway/overdue/sent, follow-up idle. `SLAPolicy.clean()` enforces `first_response_minutes` <= `resolution_minutes`.
Related Tickets live-sug	--	--	Yes	`/extensions/tickets/related/` endpoint with 3-tier auth (staff=global, user=own email, anon=explicit `?email=`). Debounced JS widget on public + staff Create Ticket forms renders up to 5 cards with title, queue, status. Privacy-safe: anonymous users without email get empty results.
Queue admin UX (fieldset)	--	--	Yes	Custom `ArchQueueAdmin` groups 26 upstream fields into 5 fieldsets (Main / Escalation / Notifications-collapsed / Email intake-collapsed / Advanced-collapsed). `QueueRoutingOverrideInline` + `QueuePoolInline` surface routing + pool config on the queue detail page. "Starts at" column in list view.
SME + classification + Buddy	--	--	Yes	SME primitives (`StaffSkill.SME_THRESHOLD=3`, `is_sme`, `StaffMember.is_sme_in(cat)`, `sme_categories()`) orthogonal to cargo seniority. In `skill_match` learning mode, `_strategy_skill_match` auto-picks an SME mentor when the junior lands a ticket; `TicketCollaborator(role='mentor', status='approved')` created automatically. ? chips on ticket banner
System Promotion Points	--	--	Yes	Skills matrix (mentor status, category) via activity feed. FQR=3 / LQR=2 / tech_note>=80chars=2 / public_followup>=100chars=1 / own_resolved=3 / mentor_on_resolved=1, capped at 10 pts per ticket. `/extensions/my-skill-path/` personal page with progress-to-threshold bar + activity feed. Admin leaderboard per category. Idempotent via `SMEPointEvent` unique_together. Promotion remains manual (human-in-the-loop).

What prod has over main (69 commits)

Area	What was added
LDAP sync	Enhanced `ldap_sync.py` -- home directory creation, improved error handling
Slurm sync	Basic `slurm_sync.py` -- account creation, user assignment, usage snapshots
Signals	Auto-user creation on OIDC login, basic allocation signals
Staff views	Initial scoped staff views for project/allocation management
Tasks	Basic django-q scheduled tasks (LDAP sync, Slurm sync)
docker-compose-dev.yml	Standalone dev compose for local development -- redefines all services (not an overlay), uses named volumes for /home and /scratch to bypass macOS virtiofs UID remap, mirrors helpdesk: so compose-hash stays consistent across commands

What dev has over prod (259 commits)

Area	What was added
Billing	Full billing infrastructure: BillingRate, AllocationBilling, HardwareCreditPool, InvoiceSettings, PaymentStatus, SlurmAccountUsagePeriod
Weekly Spend Cap	Dollar-based caps via account-level GrpTRESMins; live AJAX gauge via REST endpoint (Phase 2); diagnostic `shadow_compare_weekly_cap_usage`
Storage/Backup	Project-level fields, form validation, AllocationBilling integration
Core Hours Reset	6-period configurable reset with per-allocation overrides



Slurm Cluster	Full Django models (SlurmCluster/Node/Partition), dynamic slurm.conf, Docker cluster management
QOS	SlurmQOS model (DB-driven), per-allocation override, export_qos_config, AllowQos per partition
Hierarchy	School model, department propagation signals, orphan cleanup
TRES	Configurable tres_types, TRESBillingWeights per partition
CLI	cli_filter.lua, job_submit.lua, rates.lua, SPANK plugins, sbalance
User	CSV import, staff resource assignment, SubProject scheme, system account exclusions
Management	Sidebar nav, staff help page, job dashboard, inline editing, billing report, PI billing summary
Testing	Unit tests (billing, models, slurm_sync, signals) + start.sh test command
DevOps	Enhanced start.sh, Docker Slurm cluster, django-q audit trail, file-based email

Billing-Free Resources

These resources are never billed enforced in both UI and Slurm CLI:

Resource	Rule
`scavenger` partition	Always free -- preemptable jobs
`ssci-*` accounts	Schmidt Sciences PIs -- always free

Slurm Account Naming Convention

Project Type	Slurm Account	Example
Default PI project	`<pi_username>`	`rdesouz4`
Custom project	`<pi_username>_<abbreviation>`	`rdesouz4_proj1`
SubProject	`<pi_username>_<abbreviation>_<seq>`	`rdesouz4_proj1_A`

Full hierarchy: **root -> school -> department -> pi -> project -> users**

UID/GID Ranges

Domain	Range
JHU users	10000+
Schmidt users	30000+
Slurm system	450
Munge system	449



Staff Help & Reference Guide

Operational reference for ARCH ColdFront staff. Covers billing, Slurm integration, weekly spend caps, core hours reset, and more.

{{ section.title }}

{{ section.content|safe }}



ColdFront

Django-based allocation portal for ARCH at Johns Hopkins University. Provides user authentication (multi-realm Keycloak OIDC), allocation management, Slurm integration, billing/invoicing, job metrics, and Docker-based cluster orchestration.

Build

```
docker build -t arch/coldfront coldfront/
```

Image: **python:3.10-slim** with PostgreSQL, LDAP, Slurm CLI, Munge, and Docker CLI.

Startup (entrypoint.sh)

1. Deploy custom code from **/opt/coldfront/custom/** into site-packages
2. Run database migrations (default + slurm_jobs)
3. Create superuser (if **DJANGO_SUPERUSER_USERNAME/PASSWORD** set)
4. Seed LDAP users into Django (if **SEED_ON_RESTART=true**)
5. Collect static files
6. Start SSSD (LDAP user resolution)
7. Install cron (nightly backups, periodic Slurm sync)
8. Start Munge (if munge.key present)
9. Start Gunicorn on port 8000

Directory Structure

```

coldfront/
|-- Dockerfile
|-- entrypoint.sh
|-- custom/          -- All JHU/Schmidt-specific code (see custom/README.md)
|   |-- config/      -- Django settings (base, auth, database, plugins)
|   |-- core/user/    -- Extended UserProfile, registration, PI requests
|   |-- external/     -- PI promotion handling
|   |-- etc/          -- Runtime config (coldfront.env, local_settings.py)
|   |-- plugins/
|   |-- arch_oidc/    -- Multi-realm Keycloak OIDC backends
|   |-- arch_sync/    -- Sync, billing, cluster mgmt, job metrics
|   |-- static/       -- CSS/JS assets
|   |-- templates/    -- HTML overrides
L-- init/            -- Startup scripts (seed_initial_data.py, LDAP sync triggers)

```

Key Config



File	Purpose
`custom/etc/coldfront.env`	Django env vars (loaded by docker-compose)
`custom/etc/local_settings.py`	Runtime overrides (mounted, no rebuild needed)
`custom/config/base.py`	Installed apps, django-q schedule, middleware
`custom/config/auth.py`	LDAP + OIDC authentication backends

Environment Variables

Variable	Purpose
`DJANGO_SUPERUSER_USERNAME`	Auto-create superuser
`DJANGO_SUPERUSER_PASSWORD`	Superuser password
`SEED_ON_RESTART`	`true` = sync LDAP users on every start
`PLUGIN_SLURM`	`1` = enable Slurm plugin
`SLURM_NOOP`	`1` = no-op Slurm mode for testing
`LDAP_ADMIN_PASS`	LDAP bind password (for SSSD)

Port

- 8000 (Gunicorn HTTP)

Development

Custom files are deployed from mounted volume into site-packages on every container start. Code changes take effect on restart without rebuilding:

```
docker restart coldfront
```

Management commands:

```
docker exec coldfront python manage.py migrate
docker exec coldfront python manage.py collect_sacct --current-week
docker exec coldfront python manage.py sync_slurm
docker exec coldfront python manage.py shell
```

Admin Features

- CSV Import Admin > Import Departments / Import Users (sidebar links). Users created with **/bin/false** shell in LDAP
- Core Hours Reset Admin > Billing > Reset Core Hours. Configurable period (weekly/monthly/quarterly), manual trigger, daily django-q check
- Deny Allocation Modal with pre-filled reasons (Billing Info missing, etc.). Saved to **allocation.description**, shown as red alert banner to PI
- SubProject Scheme Set numeric (0,1,2) or alphabetic (A,B,C) per PI/resource. Locked after first SubProject is assigned
- Inline Attribute Editing Staff can edit Core Usage, Extra Core, Next Reset directly on Allocation Detail
- AllocationBilling dept -> PI profile When billing department is set, auto-added to PI's UserProfile.departments



- Service Health & Logs [/user/service-health/](#) and [/user/service-logs/](#) open to all authenticated users (works with admin impersonation)

LDAP Sync Monitoring

`coldfront/custom/plugins/arch_sync/utils/ldap_sync_monitor.py` is a read-only drift detector that compares ColdFront DB state against live LDAP. Never writes [safe to run on any schedule.](#)

On-demand (full report):

```
docker exec coldfront coldfront shell -c "  
from coldfront.plugins.arch_sync.utils.ldap_sync_monitor import check  
check()  
"
```

Reports counts for: users missing from LDAP, orphan LDAP entries, wrong **loginShell**, PI groups missing, and membership drift per PI.

One-line summary (for cron / scraping):

```
docker exec coldfront coldfront shell -c "  
from coldfront.plugins.arch_sync.utils.ldap_sync_monitor import summary  
summary()  
"  
# [ldap_sync_monitor] OK | active_users=95 | pis=31 | ldap=99 | missing=0 | orphan=0 | shell=0 | group_miss...
```

Prometheus exposure via `node_exporter` textfile collector

The **summary()** output is already in key=value form. Convert to Prometheus textfile format with a wrapper script + cron, so Grafana/Prometheus (already in the stack) can scrape it:



```
cat > /usr/local/bin/ldap_monitor_prom.sh <<'SCRIPT'
#!/bin/bash
OUT=$(docker exec coldfront coldfront shell -c "
from coldfront.plugins.arch_sync.utils.ldap_sync_monitor import summary
r = summary()
" 2>&1)

# Extract metrics and write in Prometheus textfile format
TMPFILE=/var/lib/node_exporter/textfile/ldap_sync.prom.tmp
PROMFILE=/var/lib/node_exporter/textfile/ldap_sync.prom

active=$(echo "$OUT" | grep -oP 'active_users=\K\d+')
pis=$(echo "$OUT" | grep -oP 'pis=\K\d+')
missing=$(echo "$OUT" | grep -oP 'missing=\K\d+')
shell=$(echo "$OUT" | grep -oP 'shell=\K\d+')
drift=$(echo "$OUT" | grep -oP 'membership_drift=\K\d+')

cat > $TMPFILE <<EOF
ldap_sync_active_users $active
ldap_sync_pis $pis
ldap_sync_missing_in_ldap $missing
ldap_sync_wrong_shell $shell
ldap_sync_membership_drift $drift
EOF
mv $TMPFILE $PROMFILE
SCRIPT

chmod +x /usr/local/bin/ldap_monitor_prom.sh
```

Wire it into cron (every minute):

```
* * * * * /usr/local/bin/ldap_monitor_prom.sh
```

Notes:

- Script writes to **.tmp** first, then **mv**s atomically prevents node_exporter from reading a half-written file.
 - The textfile directory (**/var/lib/node_exporter/textfile/**) must match the **collector.textfile.directory=** flag on the exporter.
 - Metrics are gauges, no **_total** suffix needed.
 - Sample Prometheus alert:
- ```
- alert: LDAPSyncDrift
 expr: ldap_sync_missing_in_ldap + ldap_sync_wrong_shell + ldap_sync_membership_drift > 0
 for: 10m
 annotations:
 summary: "ColdFront <-> LDAP drift detected"
```



# Custom Configurations

This folder contains all customizations for ColdFront specific to the ARCH project at Johns Hopkins University.

## Structure

```
custom/
|-- config/ -- Overrides for coldfront.config module
| |-- auth.py -- Authentication settings and pipelines
| |-- base.py -- Base settings (DATABASE, INSTALLED_APPS, MIDDLEWARE, etc.)
| |-- core.py -- Core functionality settings (features, options)
| |-- database.py -- Database configuration
| |-- email.py -- Email backend settings
| |-- env.py -- Environment variable parsing utilities
| |-- logging.py -- Logging configuration
| |-- plugins/ -- Plugin-specific settings (LDAP, Slurm, etc.)
| |-- settings.py -- Additional Django settings
| |-- urls.py -- URL routing with ARCH custom routes
| |-- wsgi.py -- WSGI application entry point
|
|-- core/ -- Overrides for coldfront.core module
| |-- user/ -- Custom user app
| | |-- models.py -- Extended UserProfile with has_pi_request field
| | |-- views.py -- Custom RegisterView, promote_view
| | |-- forms.py -- Custom registration form
| | |-- admin.py -- Custom admin configuration
| | |-- signals.py -- User creation signals
| | |-- urls.py -- User app routes
| | |-- migrations/ -- Database migrations
| |-- templates/user/
| | |-- login.html -- Custom login with JHU/Schmidt realm buttons
|
|-- external/ -- Overrides for coldfront.external module
| |-- promote.py -- PI promotion request handling
| |-- __init__.py
|
|-- etc/ -- Runtime configuration and utilities
| |-- coldfront.env -- Environment variables for Django
| |-- local_settings.py -- Local Django settings override
| |-- ldap_sync_mysql_run_it_once.py -- One-time LDAP -> ColdFront sync
| |-- README.md -- Configuration documentation
|
L-- plugins/ -- Custom authentication and sync plugins
| |-- arch_oidc/ -- Multi-realm Keycloak OIDC authentication
| | |-- backends.py -- JHU and Schmidt realm OIDC backends
| | |-- urls.py -- OIDC login/logout routes
| | |-- views.py -- Custom OIDC views
|
L-- arch_sync/ -- User synchronization, Slurm, billing, job metrics
| |-- apps.py -- App configuration
| |-- signals.py -- Post-create user signals, allocation sync dispatch with admin audit
| |-- tasks.py -- Async sync tasks with completion hook, admin LogEntry audit trail
| |-- models.py -- SlurmJob, BillingRate, AllocationBilling, Department, etc.
| |-- admin.py -- Admin config with weekly cap reset action
| |-- job_views.py -- Job Usage Metrics dashboard, detail, CSV views
| |-- billing_views.py -- Billing report, PI dashboard, invoice views
| |-- billing_urls.py -- URL routing for billing + job metrics
| |-- cluster_views.py -- Cluster management views (Build Images, Deploy, Restart, Stop)
| |-- cluster_manager.py -- Docker-based Slurm cluster orchestration (images_build, build_images, deplo...
| |-- utils/
| | |-- slurm_sync.py -- Slurm <-> ColdFront sync (accounts, QOS, caps, usage)
| | |-- ldap_sync.py -- LDAP user/group provisioning (adds to realm users group)
| |-- management/commands/
| | |-- collect_sacct.py -- Import per-job data from sacct or file
| | |-- sync_slurm.py -- Slurm sync management command
| | |-- populate_fake_billing_data.py -- Dev data seeding
| | |-- ...
| |-- templates/arch_sync/
| | |-- jobs/dashboard.html -- Job Metrics dashboard template
| | |-- jobs/detail.html -- Raw job list template
| | |-- billing/ -- Billing report, invoice templates
```

## How These Customizations Work

1. Config Module Override: The **config/** folder is copied into **site-packages/coldfront/config** during Docker build, overriding upstream settings.
2. Core Module Override: The **core/** folder is copied into **site-packages/coldfront/core**, allowing custom user app to



extend upstream models.

3. External Module Override: The **external/** folder is copied into **site-packages/coldfront/external** for PI promotion handling.

4. Plugin Registration: **arch\_oidc** and **arch\_sync** are custom plugins registered in **config/base.py** under **EXTRA\_APPS**.

5. Local Settings: **etc/local\_settings.py** is mounted at runtime to allow configuration changes without rebuild.

6. Environment Variables: **etc/coldfront.env** is sourced by **docker-compose.yml** to configure Django environment.

7. Startup Seed (**coldfront/init/seed\_initial\_data.py**): Runs at every container start to sync LDAP users into Django, create per-PI projects/allocations, and add allocation members from LDAP groups. Controlled by **SEED\_ON\_RESTART** in **.env**:

| Value            | Behavior                                                                                          |
|------------------|---------------------------------------------------------------------------------------------------|
| `true` (default) | Sync LDAP -> ColdFront on every container start (idempotent)                                      |
| `false`          | Skip LDAP sync on routine restarts -- use when coldfront restarts are frequent and LDAP is stable |

Set in **.env** (written by **./start.sh init**). To force a one-time re-sync after LDAP changes, use **./start.sh ldap-reinit** which also removes the OpenLDAP idempotency marker and restarts dependent services.

## Key Customizations

### Authentication (**arch\_oidc**)

- Supports multiple Keycloak realms: JHU and Schmidt
- Separate OIDC backends for each realm
- Custom login template with dual-button UI (**custom/core/user/templates/user/login.html**)

### User Management (**custom/core/user**)

- Extended UserProfile with **has\_pi\_request** field
- Custom registration form with Slurm integration
- PI promotion request handling via **custom/external/promote.py**

### User Sync, Billing & Job Metrics (**arch\_sync**)

- Syncs users from LDAP (OpenLDAP for ARCH pilot); new users added to realm **users** group
- Integrates with Slurm for compute allocation mapping
- Automatic user creation on first OIDC login
- Billing & Invoicing System: Two-track model with hardware credits and O&M charges
  - Job Usage Metrics: XDMoD-style dashboard with group-by analytics (User, Account, Partition, QOS, State, Department, School)
- Per-job tracking: **SlurmJob** model populated by **collect\_sacct** from **sacct** output or CSV file
- Weekly spend cap: Account-level **GrpTRESMins=billing** enforcement (dollar cap when billing != cpu)
- Slurm account naming: default = , **project** = \_
- Default project protection: Abbreviation and IO Number fields are disabled (read-only) on the edit form for default projects



## Billing & Invoicing System

The billing system supports two billing models per allocation:

1. Pay Per Use: Users pay full rate (compute + O&M bundled) for all hours used
2. Hardware Credit Holder: Users who purchased hardware receive credits (prepaid compute hours) that offset usage. O&M charges (power, cooling) are billed separately in cash.
3. No Charge: Allocations with no billing

### Key Features:

- Staff Billing Report (**/billing/report/**): Filterable DataTable showing usage by allocation, PI, department, school, or affiliation. Columns: Account, PI, IO#, Jobs, Partition, CPU Hours, GPU Hours, Credits Used, O&M Charges, Total. Download buttons: Summary CSV (PI/IO/Total), IO CSV (by IO number), Detail CSV (per allocation), and formatted PDF report
- Rate Management (**/billing/rates/**): Superuser can set billing rates for CPU/GPU/storage/backup with effective date ranges
- Hardware Credit Pools (**/billing/credits/**): Track hardware purchases converted to compute credit pools with expiry dates
- PI Billing Dashboard (**/billing/my-usage/**): Shows credit balance, O&M charges owed, and credit runway in separate sections (prevents confusion between credits and cash)
- Printable Invoices (**/billing/my-invoice/**): Two-section invoice showing compute credits consumed (non-cash) and O&M charges due (cash). Includes IO/cost center number for payment processing. Supports email sending.
- Allocation Billing Config: Staff can set billing type, IO number, department, and storage/backup TB per allocation via "Billing Info" card on Allocation Detail page
- Department Tracking: Users can belong to multiple departments; managed from Department Admin page (dual-pane filter widget) or UserProfile Admin. When a PI is added to a department, their projects without a department are auto-assigned and Slurm sync is triggered to reparent accounts.
- CSV Import (Admin Tools -> Import Departments / Import Users): Superuser-only views for bulk importing departments and users from CSV files. Department CSV: **name,code,school,school\_code,affiliation**. User CSV: **full\_name,email,uid**. SSCI users must have **ssci-** prefix; JHU users must use JHED ID with **@jhu.edu/@jh.edu/@jhmi.edu** email (or **ext-** prefix for non-JHU affiliates). Existing records are skipped. Users are created with **/bin/false** shell in LDAP and activated when a PI adds them to an allocation.
- Usage Collection: **manage.py collect\_billing** command collects CPU/GPU hours and job counts from Slurm per billing period

### Database Models:

- **BillingRate**: Rate catalog with cluster/partition/type (cpu\_full, cpu\_om, gpu\_full, gpu\_om, storage, backup) and effective dates
- **AllocationBilling**: Per-allocation billing config (io\_number, billing\_type, department, storage/backup TB)
- **HardwareCreditPool**: Hardware purchase records with total credits and credits consumed tracking
- **SlurmAccountUsagePeriod**: Per-period usage snapshot (source of truth for billing: cpu\_hours, gpu\_hours, job\_count)
- **School**: Organizational unit (name, code, affiliation, shares) manages Slurm fairshare at school level
- **Department**: FK to School; organizational filtering and Slurm dept-level fairshare
- **UserProfile.departments** (M2M): Users can belong to multiple departments

Setup:

Dev (all-in-one):

```
./start.sh populate dev --partitions premium,scavenger,gpu
```

This runs in order: **import\_departments** -> **populate\_fake\_billing\_data** (rates, **AllocationBilling**, credit pools) ->



**generate\_slurm\_usage\_data** (SlurmAccountUsagePeriod) -> **collect\_sacct clear** (SlurmJob).

Production / manual steps:

1. Import departments: via Admin Tools -> Import Departments (CSV upload), or **manage.py import\_departments departments.csv**
2. Import users: via Admin Tools -> Import Users (CSV upload). Users are created with unusable password + LDAP **/bin/false** shell. Activated when a PI adds them to an allocation (triggers ldap\_sync + slurm\_sync signals)
3. Assign PIs to departments: **manage.py assign\_pi departments pi\_departments.csv**
3. Create billing rates via Django Admin: **/admin/arch\_sync/billingrate/add/**
4. When an allocation is approved (status -> Active), **AllocationBilling** is created automatically. If the project has an IO number, **is\_billable=True** is set automatically. Otherwise, staff updates billing info via Allocation Detail -> "Billing Info" card.
5. Collect job data: **manage.py collect\_sacct current-month**

See **coldfront/templates/examples\_departments.csv** and **coldfront/templates/examples\_pi\_departments.csv** for CSV format examples.

## Cluster Management (arch\_sync)

Docker-based Slurm clusters are managed from the Cluster Detail page (**/cluster/**). Actions follow a required sequence:

| Step | Button                  | Condition                           | Action                                                                                                                                                                    |
|------|-------------------------|-------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1    | <b>**Build Images**</b> | Always available                    | Builds all 6 Slurm Docker images (`arch/cn-rockylinux`, `arch/slurm-munge`, `arch/slurmdbd`, `arch/slurmctld`, `arch/slurmd`, `arch/slurmrestd`) -- takes several minutes |
| 2    | <b>**Deploy**</b>       | Enabled only after images are built | Creates the compose overlay (`docker-compose-slurm-<name>.yaml`) and starts all services                                                                                  |
| 3    | <b>**Restart**</b>      | Only after deploy                   | Restarts all running Slurm containers                                                                                                                                     |
| 3    | <b>**Stop**</b>         | Only after deploy                   | Stops all Slurm containers                                                                                                                                                |
| --   | <b>**Redeploy**</b>     | After deploy, replaces Deploy       | Recreates containers with current config (rolling redeploy)                                                                                                               |

button

The Deploy button is disabled (grayed out with tooltip) until **images\_built()** confirms all required images are present on the Docker host. This prevents the common failure mode where **deploy()** fails silently because **arch/slurmd** or **arch/slurmrestd** images are missing.

**cluster\_manager.images\_built()** checks the local Docker daemon for the presence of all 6 required Slurm images. Returns **False** if Docker is unavailable.

**ClusterActionView** handles all actions via POST to **/cluster/action/**. Long-running operations (**build\_images**, **deploy**) run in a background thread to avoid Gunicorn timeout.

## Backend Services

- LDAP authentication and user directory
- MariaDB relational database
- Keycloak identity provider with Entra ID integration



## Deployment Notes

When deploying, ensure:

1. **coldfront/custom/** folder is included in git history (not in .gitignore)
2. Environment variables from **custom/etc/coldfront.env** are loaded
3. Volume mounts in docker-compose point to correct paths
4. Database migrations are current: **python manage.py migrate**
5. Static files collected: **python manage.py collectstatic**



## Billing Setup Guide

This guide walks you through populating example data for the billing system.

### Step 1: Import Departments

Create departments from CSV:

```
docker exec coldfront python manage.py import_departments /path/to/examples_departments.csv
```

Or use the example file provided:

```
docker exec coldfront python manage.py import_departments examples_departments.csv --dry-run
docker exec coldfront python manage.py import_departments examples_departments.csv
```

CSV Format: **name,code,school,affiliation**

- Example: "Computer Science","CS","Whiting School of Engineering","jhu"

See **examples\_departments.csv** for 12 example departments across JHU schools and Schmidt Science.

### Step 2: Assign PIs to Departments

Bulk assign PIs to departments from CSV:

```
docker exec coldfront python manage.py assign_pi_departments /path/to/examples_pi_departments.csv --dry-run
docker exec coldfront python manage.py assign_pi_departments /path/to/examples_pi_departments.csv
```

CSV Format: **username,department\_name**

- Example: rdesouz4,"Computer Science"

See **examples\_pi\_departments.csv** for a sample file.

### Step 3: Create Billing Rates

Automatic (first deploy): Example billing rates are seeded automatically on first startup if none exist. The default rates match **slurm/conf/rates.lua**:

| Rate Type | Partition | Rate            | Description                 |
|-----------|-----------|-----------------|-----------------------------|
| CPU Full  | (all)     | \$0.35/SU       | Pay-per-use (compute + O&M) |
| CPU O&M   | (all)     | \$0.10/SU       | Credit holders (O&M only)   |
| GPU Full  | (all)     | \$1.20/GPU-hour | Pay-per-use                 |
| GPU O&M   | (all)     | \$0.40/GPU-hour | Credit holders              |
| CPU Full  | premium   | \$0.35/SU       | Premium partition           |



|          |           |                  |                         |
|----------|-----------|------------------|-------------------------|
| CPU O&M  | premium   | \$0.10/SU        | Premium O&M             |
| CPU Full | scavenger | \$0.10/SU        | Scavenger (preemptible) |
| CPU O&M  | scavenger | \$0.05/SU        | Scavenger O&M           |
| Storage  | (all)     | \$0.025/TB/month | Storage                 |
| Backup   | (all)     | \$0.015/TB/month | Backup                  |

Manual: Go to Billing > Manage Rates > Create Rate in the UI, or re-seed by deleting existing rates and restarting:

```
docker exec coldfront python -c "import django; django.setup(); from coldfront.plugins.arch_sync.models imp...
docker restart coldfront
```

## Step 4: Configure Allocation Billing

For each allocation you want to bill:

1. Go to Allocation Detail
2. Scroll to Billing Info card
3. Set:
  - IO Number: Cost center (e.g., **1010500000**)
  - Billing Type:
  - **Pay Per Use** charged at full rate
  - **Credit Holder** credits offset compute, O&M charged separately
  - **No Charge** not billed
  - Storage TB: Provisioned storage amount (manual)
  - Backup TB: Provisioned backup amount (manual)
4. Click Update Billing

## Step 5: Compute Billing (Manual)

Once allocations have Slurm usage, compute the billing ledger. This is always manual no scheduled task writes **SlurmAccountUsagePeriod** or **AllocationBilling** rows.

Preferred UI button: go to **/billing/report/** and click Compute Billing. That runs **recalculate billing** (billing rules) followed by **sync job usage** (usage periods) in sequence.

CLI equivalent:



```
Recompute is_billable classification on all AllocationBilling rows
docker exec coldfront coldfront recalculate_billing --all

Aggregate SlurmJob -> SlurmAccountUsagePeriod for the last 30 days
docker exec coldfront coldfront sync_job_usage

Specific month / custom window
docker exec coldfront coldfront sync_job_usage --month 2026-04
docker exec coldfront coldfront sync_job_usage --start 2026-03-01 --end 2026-03-31

Clear + re-sync a month (removes stale rows)
docker exec coldfront coldfront sync_job_usage --month 2026-04 --clear
```

The 15-min **sync\_slurm\_all** task keeps the Core Usage gauge and per-user snapshots fresh, but deliberately does not touch the billing ledger. See **project\_billing\_manual\_only** memory for rationale.

## Offline Recompute (Audit / Finance Handoff)

For audit, historical recompute, or air-gapped review, **scripts/billing\_offline.py** produces an identical **PI,IO,Total** summary CSV from three plain files no ColdFront DB, no network:

```
On the cluster login node -- sacct gives partition + AllocTRES, so the
script can count CPU + GPU and apply per-partition tiers exactly like
the in-portal report.
sacct --allusers --parsable2 --noheader --allocations \
 --format jobid,partition,qos,account,user,start,end,elapsed,state,ncpus,allocres \
 --state CANCELLED,COMPLETED,FAILED,NODE_FAIL,PREEMPTED,TIMEOUT \
 --starttime 2026-05-01 --endtime 2026-06-01 > sacct.txt

Anywhere with Python 3.8+
python3 scripts/billing_offline.py \
 --sacct sacct.txt \
 --accounts accounts.csv \ # finance-supplied [pi,]account,io[,share,storage_tb,backup_tb]
 --rates slurm/conf/rates.lua \
 --output billing_2026-05.csv
```

Multi-IO splits use **share\_percent** (1-100, sum per account = 100), mirroring **AllocationCostCenter.share\_percent**. The optional **pi** column makes the script work with project-based account naming (e.g. Discovery's **pmcmood-projects**) where PI cannot be derived from the account name. Full schema, assumptions, and the Discovery walkthrough are documented in **scripts/README\_billing.md** and **scripts/README\_billing\_DISCOVERY.md**.

## Step 6: View Reports

### Staff Views

- Billing Report (filterable): **/billing/report/**
- Summary: total jobs, CPU hours, GPU hours, average
- Filter by: PI, school, department, billing type, date range
- Export: CSV button at bottom
- Invoice List: **/allocation/invoice/**



- All allocations with **is\_billable=True**
- Click allocation to see details
- Rate Management: **/billing/rates/**
- View all rates
- Add/edit rates (superuser only)
- Hardware Credits: **/billing/credits/**
- View all credit pools
- Add hardware purchase -> convert to credits

## PI Views

- My Usage & Billing: **/billing/my-usage/**
- Left: Credit balance, dollar equivalent, runway estimate
- Right: O&M charges owed (30-day rolling)
- Recent usage table
- My Invoice: **/billing/my-invoice/**
- Printable invoice
- Two sections:
  1. Compute Credits Used (hardware-backed, not a cash charge)
  2. O&M Charges (cash due)
    - Browser print (Ctrl+P) to save as PDF

## Example Workflow

1. Import departments:

```
docker exec coldfront python manage.py import_departments examples_departments.csv
docker exec coldfront python manage.py assign_pi_departments examples_pi_departments.csv
```

2. Create rates:

```
docker exec coldfront python manage.py shell < populate_example_rates.py
```

3. Configure allocations:

- Go to rdesouz4's allocation in ColdFront
- Set IO: **1010500000**
- Set Billing Type: **Pay Per Use**
- Click Update Billing

4. Compute billing (manual):



```
Preferred: click "Compute Billing" on /billing/report/
CLI equivalent:
docker exec coldfront coldfront recalculate_billing --all
docker exec coldfront coldfront sync_job_usage
```

## 5. View reports:

- Staff: **/billing/report/** -> filter by department "Computer Science"
- PI (as rdesouz4): **/billing/my-usage/** -> see charges
- PI: **/billing/my-invoice/123/** -> print invoice

## Key Concepts

### Two-Track Billing Model

- Credits (from hardware purchase)
- Cover compute hours ONLY
- Do NOT pay O&M costs
- UI shows clearly as separate from cash
- O&M Charges (always cash)
- Power, cooling, maintenance
- Applied separately per hour used
- Cannot be paid with credits

### Billing Types

| Type          | CPU Charge           | O&M Charge        | Uses Credits?      |
|---------------|----------------------|-------------------|--------------------|
| Pay Per Use   | Full rate (cpu_full) | Included          | No                 |
| Credit Holder | \$0                  | O&M rate (cpu_om) | Yes (compute only) |
| No Charge     | \$0                  | \$0               | No                 |

## Troubleshooting

### Q: Billing Report shows no data

- Check that allocations have **is\_billable=True** in the Billing Info card
- Click Compute Billing on **/billing/report/** (or run **coldfront recalculate\_billing all && coldfront sync\_job\_usage**). Billing is manual-only nothing populates it automatically, not even **./start.sh deploy**

### Q: Invoice List is empty

- Verify allocations have Billing Info card configured
- Check that **is\_billable=True** is set

### Q: PI can't see My Usage & Billing

- Verify they have at least one allocation with billing configured
- Check they're logged in as that PI user



## Q: Rates not applying

- Verify **effective\_from** date is on or before usage period end date
- Check **effective\_to** is null (active) or after usage period

## Files Included

- **examples\_departments.csv** 12 departments (JHU Whiting, Krieger, Medicine + Schmidt Science)
- **examples\_pi\_departments.csv** 10 PIs assigned to departments (including rdesouz4)
- **populate\_example\_rates.py** Create 6 example rates via Django shell
- **BILLING\_SETUP\_GUIDE.md** This file

## Departments Included

### Whiting School of Engineering (JHU)

- Computer Science (rdesouz4)
- Electrical & Computer Engineering (agarcia7)
- Applied Physics (mchen2)
- Mechanical Engineering (ekim6)
- Biomedical Engineering

### Krieger School of Arts & Sciences (JHU)

- Biology (jsmith5)
- Chemistry (bpatel9)
- Physics (lrodriguez4)

### School of Medicine (JHU)

- Neuroscience (kwilson3)
- Molecular Biology & Genetics

### Schmidt Science

- Structural Biology (dgonzalez1)
- Computational Biology (hjones8)

## Quick Start



```
1. Import 12 departments
docker exec coldfront python manage.py import_departments examples_departments.csv

2. Assign 10 PIs to departments
docker exec coldfront python manage.py assign_pi_departments examples_pi_departments.csv

3. Create 6 billing rates
docker exec coldfront python manage.py shell < populate_example_rates.py

4. Configure allocations for billing
For each PI with an allocation:
- Go to Allocation Detail
- Scroll to Billing Info card
- Set IO Number (e.g., 1010500000)
- Set Billing Type: Pay Per Use or Credit Holder
- Click Update Billing

5. Compute billing (manual -- click Compute Billing on /billing/report/, or CLI)
docker exec coldfront coldfront recalculate_billing --all
docker exec coldfront coldfront sync_job_usage

6. View reports
Staff: /billing/report/ (filter by school/department/PI)
PI: /billing/my-usage/ (see your charges)
Invoice: /allocation/invoice/ (all billable allocations)
```

Test Users (10 PIs across multiple departments):

- rdesouz4 (Computer Science, Whiting)
- jsmith5 (Biology, Krieger)
- mchen2 (Applied Physics, Whiting)
- agarcia7 (ECE, Whiting)
- bpatel9 (Chemistry, Krieger)
- kwilson3 (Neuroscience, Medicine)
- dgonzalez1 (Structural Biology, Schmidt)
- hjones8 (Computational Biology, Schmidt)
- lrodriguez4 (Physics, Krieger)
- ekim6 (Mechanical Engineering, Whiting)



# Billing Data Import

---

This directory contains sample CSV files and instructions for populating the ColdFront billing system with departments, rates, billing configurations, and hardware credit pools.

## CSV Files Overview

### 1. departments.csv

Purpose: Define the organizational structure (schools, departments, affiliations)

Format:

```
name,code,school,affiliation
Computer Science,CS,Whiting School of Engineering,jhu
Data Science,DS,Schmidt Sciences Institute,schmidt
```

Fields:

- **name** (required): Department name (e.g., "Computer Science")
- **code** (optional): Department code (e.g., "CS")
- **school** (optional): School/Institute name (e.g., "Whiting School of Engineering")
- **affiliation** (optional): Either "jhu" or "schmidt" for filtering/organization

Import Command:

```
python manage.py import_departments departments.csv
```

### 2. pi\_departments.csv

Purpose: Assign PIs to departments (supports multiple departments per PI)

Format:

```
username,department_names
testuser1,Computer Science;Applied Physics
testuser2,Data Science
```

Fields:

- **username** (required): JHU username
- **department\_names** (required): Semicolon-separated list of department names

Import Command:

```
python manage.py assign_pi_departments pi_departments.csv --dry-run
python manage.py assign_pi_departments pi_departments.csv # Run without --dry-run after verification
```



### 3. billing\_rates.csv

Purpose: Define billing rates for compute, storage, and backup resources

Format:

```
cluster,partition_name,rate_type,rate_per_unit,effective_from,effective_to,notes
ARCH,gpu,cpu_full,0.45,2026-01-01,,Full rate for CPU hours (compute + O&M bundled)
ARCH,gpu,gpu_om,0.40,2026-01-01,,O&M only charge for GPU hours
```

Fields:

- **cluster** (optional): SLURM cluster name
- **partition\_name** (optional): Partition name (empty = all partitions)
- **rate\_type** (required): One of:
  - **cpu\_full**: CPU hours at full rate (compute + O&M)
  - **cpu\_om**: CPU hours at O&M only (for credit holders)
  - **gpu\_full**: GPU hours at full rate
  - **gpu\_om**: GPU hours at O&M only
- **storage**: Storage per TB/month
- **backup**: Backup per TB/month
- **rate\_per\_unit** (required): Cost per unit (decimal, up to 6 decimals)
- **effective\_from** (required): Start date (YYYY-MM-DD)
- **effective\_to** (optional): End date or blank for currently active
- **notes** (optional): Description or comments

Import via Django Admin:

Go to `/admin/arch_sync/billingrate/` and add rates manually, or create a management command for bulk import.

### 4. allocations\_billing.csv

Purpose: Configure billing per allocation (IOS code, billing type, storage/backup provisioning)

Format:

```
allocation_id,ios_number,department_name,billing_contact,billing_type,storage_tb,backup_tb,is_billable
1,I0-2024-001,Computer Science,john.doe@jhu.edu,pay_per_use,10.0,5.0,True
2,I0-2024-002,Data Science,jane.smith@ssci.jhu.edu,credit_holder,50.0,25.0,True
```

Fields:

- **allocation\_id** (required): ColdFront allocation ID
- **ios\_number** (required): IOS/cost center for billing
- **department\_name** (required): Department name (must exist)
- **billing\_contact** (required): Email or name of billing contact
- **billing\_type** (required): One of:
  - **pay\_per\_use**: Full rate billed to user
  - **credit\_holder**: Hardware credits offset compute; O&M charged separately
  - **no\_charge**: No billing for this allocation



- **storage\_tb** (required): Provisioned storage in TB (0 if none)
- **backup\_tb** (required): Provisioned backup in TB (0 if none)
- **is\_billable** (required): True/False - whether to include in billing reports

Import via Management Command:

```
python manage.py import_allocation_billing allocations_billing.csv --dry-run
python manage.py import_allocation_billing allocations_billing.csv
```

## 5. hardware\_credits.csv

Purpose: Record hardware purchases and convert to credit pools

Format:

```
allocation_id,description,hardware_value,total_credits,credit_type,granted_date,expiry_date,notes
2,"4x H100 GPU purchase FY2025",85000.00,87600,gpu,2025-01-15,2027-01-15,Purchased through grant funding
```

Fields:

- **allocation\_id** (required): ColdFront allocation ID (allocation must exist and have billing config)
- **description** (required): Hardware purchase description
- **hardware\_value** (required): Cost of hardware in dollars
- **total\_credits** (required): Total compute hours granted
- **credit\_type** (required): "cpu" or "gpu"
- **granted\_date** (required): Date credits become available (YYYY-MM-DD)
- **expiry\_date** (optional): Date credits expire or blank for no expiration
- **notes** (optional): Additional notes

Import via Management Command:

```
python manage.py import_hardware_credits hardware_credits.csv --dry-run
python manage.py import_hardware_credits hardware_credits.csv
```

## Quick Start - Populate Fake Data for Testing

To quickly test the billing system without manual CSV setup:

```
1. Create fake billing data (departments, rates, allocations, usage)
python manage.py populate_fake_billing_data

2. Create allocations and usage specifically for a user
python manage.py populate_user_allocations rdesouz4
python manage.py populate_user_allocations testuser1
```

Then visit:

- **/billing/report/** - See all billing report with fake data
- Click "Print / Save as PDF" to download as PDF



- Click filter by PI and search for the user

## Import Workflow

Recommended order:

1. Import departments (creates department records)
2. Assign PIs to departments (links existing users to departments)
3. Import billing rates (sets up rate catalog)
4. Import allocation billing configs (configures each allocation for billing)
5. Import hardware credits (records hardware purchases and credit pools)

Example full workflow:

```
cd /path/to/coldfront

1. Import departments
python manage.py import_departments departments.csv

2. Assign PIs to departments (with dry-run first)
python manage.py assign_pi_departments pi_departments.csv --dry-run
python manage.py assign_pi_departments pi_departments.csv

3. Import billing rates (via admin or custom command)
(Use Django admin interface at /admin/arch_sync/billingrate/ for now)

4. Import allocation billing configs
python manage.py import_allocation_billing allocations_billing.csv --dry-run
python manage.py import_allocation_billing allocations_billing.csv

5. Import hardware credits
python manage.py import_hardware_credits hardware_credits.csv --dry-run
python manage.py import_hardware_credits hardware_credits.csv
```

## Verification

After importing, verify data in Django admin:

- [/admin/arch\\_sync/department/](#) - View departments
- [/admin/arch\\_sync/billingrate/](#) - View billing rates
- [/admin/arch\\_sync/allocationbilling/](#) - View allocation billing configs
- [/admin/arch\\_sync/hardwarecreditpool/](#) - View hardware credit pools

## Staff Views

Once data is imported:

- Billing Report: [/billing/report/](#) - View all charges and credits
- Rate Management: [/billing/rates/](#) - View/edit rates
- Hardware Credits: [/billing/credits/](#) - View/manage credit pools
- Allocation Billing: [/billing/allocation/set/](#) - Configure per-allocation billing



## PI Views

Once data is imported, PIs can view:

- My Usage & Billing: **/billing/my-usage/** - Dashboard with credit balance and O&M charges
- My Invoice: **/billing/my-invoice/** - Printable invoice with costs and credit usage

## Management Commands

### populate\_fake\_billing\_data.py

Creates fake billing data for testing:

```
python manage.py populate_fake_billing_data # Create fake data
python manage.py populate_fake_billing_data --clear # Clear and recreate
python manage.py populate_fake_billing_data --no-usage # Skip usage periods (faster)
```

Creates:

- 7 sample departments (JHU + Schmidt Sciences)
- 8 billing rates (CPU, GPU, storage, backup)
- 5 allocation billing configs
- 15 usage periods with realistic hours
- 1 hardware credit pool

### populate\_user\_allocations.py

Creates allocations and usage for a specific user:

```
python manage.py populate_user_allocations rdesouz4 # Create for rdesouz4
python manage.py populate_user_allocations testuser1 # Create for testuser1
python manage.py populate_user_allocations username --clear # Clear and recreate
```

Creates:

- 1-3 allocations for the user on available resources
- Realistic billing configurations (pay-per-use or credit-holder)
- 3 monthly usage periods with jobs, CPU hours, GPU hours
- Hardware credit pool for credit-holder allocations

## Notes

- All CSV imports are idempotent (safe to re-run with same data)
- Use **dry-run** flag to preview changes before committing
- Affiliation is case-sensitive: use "jhu" or "schmidt" (lowercase)
- Credit types: "cpu" for CPU hours, "gpu" for GPU hours
- Rates use up to 6 decimal places for precision (e.g., 0.123456)



- Storage/backup provisioning is manual - update via allocation billing config form
- Use **populate\_fake\_billing\_data** for system-wide testing
- Use **populate\_user\_allocations USERNAME** for user-specific testing



# Billing & Storage Changes

---

## Changes Made

### 1. Enhanced Billing Information Form

File: `coldfront/custom/plugins/arch_sync/templates/arch_sync/billing/set_billing_form.html`

PIs and admins now have choice interfaces for:

- IO Number: Use project's IO or set a custom one per allocation
- Storage (TB): Use project's storage or set custom per allocation
- Backup (TB): Directly manage backup capacity

Features:

- Shows project defaults as reference
- Radio buttons for easy selection
- Custom fields slide in/out based on choice
- Responsive UI with dynamic visibility

### 2. Updated Billing Views

File: `coldfront/custom/plugins/arch_sync/billing_views.py`

**SetAllocationBillingView** changes:

- GET: Now passes **default\_storage\_tb** to template from project
- POST: Intelligently handles choice logic
- If **io\_choice='project'**: uses project's IO number
- If **storage\_choice='project'**: uses project's storage
- Otherwise: uses custom values entered by user

Key Logic:

```
IO Number
if io_choice == 'project':
 io_number = project_io_number
else:
 io_number = request.POST.get('io_number', '').strip()

Storage
if storage_choice == 'project':
 storage_tb = project_storage_tb
else:
 storage_tb = float(request.POST.get('storage_tb', '0'))
```

### 3. New Allocation Attribute Type: Backups (TB)

File: `coldfront/custom/plugins/arch_sync/management/commands/create_backup_attribute.py`



Creates a new allocation attribute type allowing:

- PIs/staff to manage backup quotas through allocation attributes
- Consistent with existing Storage Quota management
- Type: Float, changeable, no usage tracking

## 4. Admin Interface

Already Configured In: **coldfront/custom/plugins/arch\_sync/admin.py**

**AllocationBillingAdmin** fieldsets:

```
Allocation
| - allocation
L - department
Billing Configuration
| - io_number
| - billing_type
L - billing_contact
Resources
| - storage_tb
L - backup_tb
Status
L - is_billable
Administration
| - updated_by
L - updated_at
```

## Deployment Steps

### Step 1: Create the Backups Allocation Attribute Type

```
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py create_backup_attribute
```

Expected output:

```
v Created allocation attribute type: Backups (TB)
```

### Step 2: Verify Changes

The following are now available:

1. For PIs (on their own allocations):

- Navigate to allocation detail page
- Click "Set Billing Info" button
- See choice interfaces for:
  - IO Number (project default vs custom)
  - Storage (project default vs custom)
  - Backup capacity (direct entry)



## 2. For Staff (in admin panel):

- Allocation > Set Billing Info
- View all billing fields with same choice interface
- Edit allocation attributes to set "Backups (TB)" quota

## 3. In Admin Dashboard:

- Go to Allocation Billing
- View storage\_tb and backup\_tb in Resources section
- Filter and search by billing configuration

## Data Model

### AllocationBilling Model (Updated)

```
allocation (FK) -> Associated allocation
io_number (CharField) -> Cost center for billing
billing_type (CharField) -> pay_per_use, credit_holder, no_charge
storage_tb (FloatField) -> Storage capacity in TB
backup_tb (FloatField) -> Backup capacity in TB
billing_contact (Text) -> Contact info
is_billable (Boolean) -> Include in billing reports
department (FK) -> Department for attribution
updated_by (FK) -> Last modified by user
updated_at (DateTime) -> Last modification timestamp
```

### New Allocation Attribute Type

```
Name: Backups (TB)
Type: Float
Changeable: Yes
Has Usage: No
Purpose: Track backup quota per allocation
```

## Usage Examples

### Example 1: PI Using Project IO

1. Open allocation detail -> "Set Billing Info"
2. IO Number section shows: "Project IO Number: IO-2026-0001"
3. Select "Use project's IO Number"
4. Save -> AllocationBilling.io\_number = "IO-2026-0001"

### Example 2: PI Overriding Storage

1. Open allocation detail -> "Set Billing Info"
2. Storage section shows: "Project Storage: 100 TB"
3. Select "Use a different storage allocation"



4. Enter custom value: 50
5. Save -> AllocationBilling.storage\_tb = 50.0

### Example 3: Staff Managing Backups via Attributes

1. Go to Admin -> Allocation Attributes
2. Select attribute type: "Backups (TB)"
3. Enter backup quota: 10 (TB)
4. Also set in AllocationBilling for billing purposes

### Verification Checklist

- [ ] Backups (TB) attribute type created
- [ ] Billing form displays IO Number choice interface
- [ ] Billing form displays Storage choice interface
- [ ] Radio buttons toggle custom field visibility
- [ ] Form submits with correct values
- [ ] Admin can see storage\_tb and backup\_tb in AllocationBilling
- [ ] PI can access and modify billing info (via StaffOrAllocationPIMixin)
- [ ] Database stores custom values correctly

### Testing Commands

```
Create backups attribute
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py create_backup_attribute

Check if attribute was created
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py shell -c "
from coldfront.core.allocation.models import AllocationAttributeType
attr = AllocationAttributeType.objects.get(name='Backups (TB)')
print(f'v Found: {attr.name} (type: {attr.attribute_type})')
"

Test PI access (as PI)
Navigate to: /allocation/{id}/set-billing-info/
Should see choice interfaces for IO Number and Storage
```

### Files Modified

1. **coldfront/custom/plugins/arch\_sync/templates/arch\_sync/billing/set\_billing\_form.html**
  - Added choice interface for IO Number
  - Added choice interface for Storage
  - Added JavaScript for toggling visibility
2. **coldfront/custom/plugins/arch\_sync/billing\_views.py**



- Updated GET to pass default\_storage\_tb
- Enhanced POST with choice logic for io\_choice and storage\_choice

### 3. coldfront/custom/plugins/arch\_sync/management/commands/create\_backup\_attribute.py (NEW)

- Creates "Backups (TB)" allocation attribute type

## Notes

- AllocationBilling model already has storage\_tb and backup\_tb fields
- Admin interface already displays these fields in Resources fieldset
- StaffOrAllocationPIMixin ensures PIs can only modify their own allocations
- Choice interface gracefully falls back to simple input if no project defaults exist
- Both IO Number and Storage can be managed per-allocation independently



## Invoice Settings

---

The billing system now includes customizable invoice headers and branding options. Superusers can configure how invoices appear to PIs.

### How to Customize Invoices

1. Log in as superuser (admin)
2. Go to Django Admin -> **/admin/**
3. Navigate to ARCH Sync -> Invoice Settings
4. Click on the existing settings record (or create one)
5. Customize the following fields:

### Available Settings

#### Organization Information

- Organization Name: Name displayed on invoice header (default: "ARCH | Research Computing Service Center")
- Organization Address: Full address including street, city, state, zip
- Organization Phone: Contact phone number
- Organization Email: Contact email address
- Logo URL: URL to organization logo image (optional)

#### Invoice Appearance

- Invoice Title: Title displayed at top of invoice (default: "BILLING INVOICE")
- Invoice Footer Text: Text shown at bottom of invoice

#### Invoice Text & Explanations

- Credit Explanation: Explanation text for hardware credits section
- Default: "Credits represent hardware-backed compute hours purchased by your grant. They cover CPU/GPU runtime only and do NOT offset operational costs (power, cooling, maintenance)."
- O&M Explanation: Explanation text for O&M charges section
- Default: "O&M (Operations & Maintenance) charges cover the operational costs of the facility including power, cooling, and infrastructure maintenance. These charges are real cash costs and cannot be paid with credits."

### How It Works

For PIs (Users):

- When a PI views their invoice at **/billing/my-invoice/**
- The invoice displays the customized organization name, address, and logo
- Invoice text explanations help PIs understand credits vs cash charges



For Staff:

- Same customizations appear on staff billing reports

## Example Configuration

```
Organization Name: Johns Hopkins University - ARCH Computing
Organization Address: 3400 North Charles Street, Baltimore, MD 21218
Organization Phone: (410) 555-0100
Organization Email: arch-admin@jhu.edu
Logo URL: https://www.jhu.edu/static/images/jhu-logo.png

Invoice Title: MONTHLY BILLING INVOICE

Invoice Footer Text:
"Thank you for using the ARCH Computing Service.
Questions? Email: arch-admin@jhu.edu"

Credit Explanation:
"Your hardware purchase credits represent compute hours on our systems.
These credits cover CPU/GPU runtime hours only.
Operational & Maintenance costs (power, cooling, facility) are billed separately."

O&M Explanation:
"Operations & Maintenance (O&M) charges cover the operational costs of maintaining
the computing facility. These are real cash costs and must be paid separately.
They cannot be covered by hardware credits."
```

## Important Notes

- Singleton Model: Only ONE invoice settings record can exist
- Deletion Protection: Settings cannot be deleted
- Updated By: Tracks which user last modified the settings
- Timestamps: Shows when settings were created and last updated
- Logo URL: Must be a valid URL; images are embedded in the invoice HTML

## API Access

To get the current settings programmatically:

```
from coldfront.plugins.arch_sync.models import InvoiceSettings

settings = InvoiceSettings.get_settings()
print(settings.organization_name)
print(settings.organization_email)
```

## Testing

After customizing invoice settings:

1. Log in as a PI
2. Go to My Invoice -> **/billing/my-invoice/**



3. Verify organization name, address, and logo appear correctly
4. Click Print / Save as PDF to download the invoice
5. Verify all customized text appears in the PDF



# School & Department Tracking

---

## Overview

Added school and department information to job tracking (**SlurmJob** model) for comprehensive billing reports and analytics by organizational structure.

## Changes Made

### 1. Database Schema

File: `coldfront/custom/plugins/arch_sync/models.py`

Added two new fields to **SlurmJob** model:

```
school = models.CharField(max_length=256, blank=True, db_index=True)
department = models.CharField(max_length=256, blank=True, db_index=True)
```

Denormalization Strategy:

- School and department are stored directly in SlurmJob (denormalized)
- Populated from **Allocation.Project.Department.school** and **Allocation.Project.Department.name**
- This enables faster queries and aggregations without complex joins

### 2. Migration

File: `coldfront/custom/plugins/arch_sync/migrations/0013_slurmjob_school_department.py`

Creates the two new indexed columns in the **arch\_sync\_slurmjob** table.

### 3. Job Collection

File: `coldfront/custom/plugins/arch_sync/management/commands/collect_sacct.py`

Enhanced to populate school and department when collecting jobs:

- Updated **acct\_to\_alloc** lookup to include school and department info
- For each job, resolves: **account** -> **allocation** -> **project** -> **department** -> {name, school}
- Handles missing allocations gracefully (empty string defaults)

Updated Logic:

```
acct_to_alloc[account] = {
 'allocation_id': alloc.id,
 'school': school_name or '',
 'department': department_name or '',
}
```



## 4. CSV Export

File: `coldfront/custom/plugins/arch_sync/management/commands/export_jobs_usage_csv.py` (NEW)

New management command to export job data with school and department:

```
python manage.py export_jobs_usage_csv [--start YYYY-MM-DD] [--end YYYY-MM-DD] [--output FILE]
```

Features:

- Exports to `coldfront/templates/jobs_usage.csv` by default
- Includes 33 columns with proper CSV headers
- Includes school and department columns
- Filters by date range (default: last 30 days)

CSV Columns:

```
job_id | job_id_raw | cluster | partition | qos | account | group_name | gid_number |
user | uid_number | submit_time | eligible_time | start_time | end_time |
wall_duration_secs | wait_duration_secs | exit_code | state | n_nodes | n_cpus |
req_cpus | req_mem | req_tres | alloc_tres | time_limit | node_list | job_name |
cpu_time_hours | gpu_time_hours | node_time_hours | school | department | allocation_id
```

## 5. Job Metrics Dashboard

File: `coldfront/custom/plugins/arch_sync/job_views.py`

Optimized grouping by school and department:

- Changed from Python-side lookup to direct ORM field aggregation
- Much faster queries with `values('school')` instead of account-based lookup
- Enables filtering by school/department at the database level

Updated GROUP\_BY\_OPTIONS:

```
('department', 'department', 'Department', False), # direct field
('school', 'school', 'School', False), # direct field
```

## Deployment Steps

### Step 1: Migrate Database

```
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py migrate arch_sync
```

Expected output:

```
Applying arch_sync.0013_slurmjob_school_department... OK
```



## Step 2: Re-collect Job Data (Optional but Recommended)

To populate school and department for existing jobs:

```
Re-collect last 30 days of jobs
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py collect_sacct --current-month

Or specify a custom date range
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py collect_sacct --start 2026-01-01 --end 2026-03-28
```

## Step 3: Export Jobs Usage CSV

```
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py export_jobs_usage_csv

Or with custom date range
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py export_jobs_usage_csv --start 2026-03-01 --end 2026-03-28
```

The CSV will be saved to: **coldfront/templates/jobs\_usage.csv**

## Usage Examples

### Example 1: Group Jobs by School

Navigate to **/billing/jobs/** and select "School" from the "Group by" dropdown.

Result: See total CPU hours, job count, and other metrics aggregated by school.

### Example 2: Group Jobs by Department

Select "Department" from the "Group by" dropdown.

Result: See metrics aggregated by department across the organization.

### Example 3: Export School/Department Data

```
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py export_jobs_usage_csv --start 2026-03-01 --output /tmp/jobs_march.csv
```

Then import **/tmp/jobs\_march.csv** into your analysis tool (Excel, Python pandas, etc.).

## Data Structure Example

### Before (Old sacct output)



```
job_id|cluster|account|user|cpu_hours|gpu_hours|...[no school/department]
1001|arch-dev|ricardo|ricardo|2.5|0|...
```

### After (New with school/department)

```
job_id|cluster|account|user|cpu_hours|gpu_hours|...|school|department|...
1001|arch-dev|ricardo|ricardo|2.5|0|...|Whiting School of Engineering|Applied Physics|...
```

## Query Performance

### Before (Python-side lookup)

```
1. Query all accounts: SELECT account FROM arch_sync_slurmjob
2. Lookup account -> department/school: 1000+ Python dict lookups
3. Aggregate in Python: manually sum values
~ O(n) time complexity with manual aggregation
```

### After (Direct ORM query)

```
SELECT department, SUM(cpu_time_hours)
FROM arch_sync_slurmjob
WHERE start_time BETWEEN X AND Y
GROUP BY department
~ O(1) database-level aggregation
```

Performance Improvement: 10-100x faster for large datasets.

## Verification Checklist

- [ ] Migration applied successfully
- [ ] New columns (school, department) exist in arch\_sync\_slurmjob table
- [ ] **collect\_sacct** populates school/department for new jobs
- [ ] Job Metrics Dashboard groups by School/Department
- [ ] **export\_jobs\_usage\_csv** command works
- [ ] CSV includes school and department columns
- [ ] Existing jobs can be re-collected with school/department info
- [ ] Dashboard aggregation is fast (< 1 second)

## Testing Commands



```
Check if migration is applied
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py migrate arch_sync --plan

Verify new columns exist
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py dbshell -c "\d arch_sync_slurmjob" | grep -E "school|department"

Test collect_sacct with new fields
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py collect_sacct --dry-run --current-week

Verify data in dashboard
Navigate to: http://localhost/billing/jobs/?group_by=school
Should show jobs grouped by school with proper aggregations

Check CSV export
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py export_jobs_usage_csv && \
 head -1 /opt/arch/coldfront/templates/jobs_usage.csv
```

## Files Modified/Created

| File                                                                               | Type     | Change                                              |
|------------------------------------------------------------------------------------|----------|-----------------------------------------------------|
| `coldfront/custom/plugins/arch_sync/models.py`                                     | Modified | Added school, department fields to SlurmJob         |
| `coldfront/custom/plugins/arch_sync/migrations/0013_slurmjob_school_department.py` | Created  | Migration for new columns                           |
| `coldfront/custom/plugins/arch_sync/management/commands/collect_sacct.py`          | Modified | Populate school/department when collecting jobs     |
| `coldfront/custom/plugins/arch_sync/management/commands/export_jobs_usage_csv.py`  | Created  | New CSV export command with school/department       |
| `coldfront/custom/plugins/arch_sync/job_views.py`                                  | Modified | Optimize department/school grouping with direct ORM |

## Billing Report Enhancement

The school and department information can now be used for:

1. Billing by Department: Group charges by department for cost allocation
2. School Reports: Generate reports showing each school's total resource usage
3. Cross-organizational Analytics: Compare usage patterns between schools/departments
4. Hierarchical Reporting: JHU Affiliation -> Schools -> Departments -> PIs -> Projects -> Allocations -> Jobs

## Notes

- School and department are denormalized in SlurmJob for query performance
- They're populated at job collection time from the allocation's project department
- If a job's allocation doesn't have a department, school/department will be empty
- The fields are indexed for fast grouping/filtering operations
- Existing jobs will have empty school/department unless re-collected

## Future Enhancements



- Add affiliation (JHU vs Schmidt) as another denormalized field
- Create a hierarchical tree view of resource usage (JHU -> School -> Department -> PI -> Project)
- Add department-level billing summaries and cost allocation
- Implement department quota/cap tracking



# Re-collect Historical Jobs

---

## Overview

Re-collect past job data from Slurm sacct to populate the new school and department fields in SlurmJob records.

## Prerequisites

- ? Migration 0013\_slurmjob\_school\_department applied
- ? All allocations have projects with departments assigned
- ? Slurm cluster is accessible (slurmrestd running)

## Step-by-Step Process

### Step 1: Identify Date Range

Determine which historical period you want to collect:

```
Example ranges:
Last 3 months: 2025-12-28 to 2026-03-28
Last 6 months: 2025-09-28 to 2026-03-28
All of 2026: 2026-01-01 to 2026-03-28
Specific month: 2026-03-01 to 2026-03-31

START_DATE="2026-01-01"
END_DATE="2026-03-28"
```

### Step 2: Re-collect Job Data with Dry-Run (Optional but Recommended)

First, test the collection without writing to the database:

```
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py collect_sacct --start $START_DATE --end $END_DATE --dry-run
```

Expected output:

```
collect_sacct: 2026-01-01 -> 2026-03-28 [DRY-RUN]
Loaded 2847 lines from sacct
collect_sacct done: 2847 parsed, 0 updated -- 2847 raw lines parsed
```

This shows how many jobs would be collected without actually writing them.

### Step 3: Actually Collect the Data

Once you're confident, run without **dry-run**:



```
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py collect_sacct --start $START_DATE --end $END_DATE
```

Expected output:

```
collect_sacct: 2026-01-01 -> 2026-03-28
Loaded 2847 lines from sacct
collect_sacct done: 2847 new, 0 updated -- 2847 raw lines written
```

## Step 4: Verify School & Department Were Populated

```
Check a sample of jobs
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py shell -c "
from coldfront.plugins.arch_sync.models import SlurmJob
jobs = SlurmJob.objects.filter(school__isnull=False).exclude(school='')[:5]
for job in jobs:
 print(f'Job {job.job_id}: {job.account} @ {job.school} / {job.department}')
"
```

Expected output:

```
Job 1001: ricardo @ Whiting School of Engineering / Applied Physics
Job 1002: jasonw @ Whiting School of Engineering / Computer Science
Job 1003: ssci-gomes @ Schmidt Sciences / Biology
...
```

## Step 5: Export Jobs Usage CSV

Now export the enriched data to CSV with school and department:

```
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py export_jobs_usage_csv --start $START_DATE --end $END_DATE
```

Expected output:

```
Exporting 2847 jobs from 2026-01-01 to 2026-03-28
v Successfully exported 2847 jobs to /opt/arch/coldfront/templates/jobs_usage.csv
```

## Step 6: Verify CSV Content

Check that the CSV includes school and department columns:

```
Show header and first few rows
docker compose -f docker-compose-dev.yml exec coldfront \
bash -c "head -2 /opt/arch/coldfront/templates/jobs_usage.csv | cut -d',' -f1-15,30-32"
```



Expected output:

```
job_id,job_id_raw,cluster,partition,qos,account,group_name,gid_number,user,uid_number,submit_time,eligible_...
1001,1001,arch-dev,cpu,normal,ricardo,ricardo,1001,ricardo,1001,2026-01-15T10:30:00,2026-01-15T10:31:00,202...
```

## Collecting in Batches (Large Datasets)

If you need to collect many months of data, do it in batches to avoid large memory usage:

```
Collect month by month
for month in {01..03}; do
 START="2026-${month}-01"
 END="2026-${month}-28"
 echo "Collecting $START to $END..."
 docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py collect_sacct --start $START --end $END
done
```

## Troubleshooting

### Problem: "No jobs found" / Empty output

Solution: Check that:

1. Slurm sacct has data for that date range: **sacct start=2026-01-01 end=2026-01-31**
2. Jobs exist in the database but without school/department (old jobs)
3. The slurmrestsd service is running: **docker compose -f docker-compose-dev.yml ps slurmrestsd**

### Problem: "school/department are empty"

Solution:

1. Check allocation has a project: **docker compose -f docker-compose-dev.yml exec coldfront python manage.py shell -c "from coldfront.core.allocation.models import Allocation; print(Allocation.objects.count())"**
2. Check project has a department: **docker compose -f docker-compose-dev.yml exec coldfront python manage.py shell -c "from coldfront.core.project.models import Project; print(Project.objects.filter(departmentisnull=False).count())"**
3. Re-run collect\_sacct - it will populate school/department for all jobs

### Problem: "account not found in lookup"

Solution: Jobs with accounts that don't map to allocations will have empty school/department, which is fine. Check logs:

```
docker compose -f docker-compose-dev.yml exec coldfront \
 python manage.py collect_sacct --start 2026-03-01 --end 2026-03-28 2>&1 | grep -i "account\|error"
```

## Typical Usage Examples



## Example 1: Collect Last Month

```
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py collect_sacct --current-month && \
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py export_jobs_usage_csv --end $(date +%Y-%m-%d)
```

## Example 2: Collect Q1 2026

```
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py collect_sacct --start 2026-01-01 --end 2026-03-31 && \
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py export_jobs_usage_csv --start 2026-01-01 --end 2026-03-31
```

## Example 3: One-Time Historical Collection (Entire Year)

```
Dry-run first
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py collect_sacct --start 2025-01-01 --end 2025-12-31 --dry-run

Actually collect
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py collect_sacct --start 2025-01-01 --end 2025-12-31

Export
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py export_jobs_usage_csv --start 2025-01-01 --end 2025-12-31
```

## Performance Tips

- For large date ranges: Collect in monthly batches to reduce memory usage
- For first-time collection: Use [dry-run first to verify data](#)
- After collection: Run export immediately while data is fresh in cache
- Parallel collection: You can run multiple collect\_sacct commands on different date ranges simultaneously in different shells

## Verification After Collection

### Dashboard Check

Navigate to **/billing/jobs/** and verify:

- ? "Group by: School" shows jobs grouped by school name
- ? "Group by: Department" shows jobs grouped by department name
- ? CPU hours are aggregated correctly by school/department
- ? Load time is fast (< 1 second)

### Database Check



```
Count jobs with school/department populated
docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py shell -c "
from coldfront.plugins.arch_sync.models import SlurmJob
total = SlurmJob.objects.count()
with_dept = SlurmJob.objects.exclude(school='').count()
print(f'Total jobs: {total}')
print(f'Jobs with school: {with_dept}')
print(f'Coverage: {with_dept/total*100:.1f}%')
"
```

## CSV Check

```
Verify CSV has data and includes school/department
docker compose -f docker-compose-dev.yml exec coldfront \
 bash -c "
echo 'Total rows:' && wc -l /opt/arch/coldfront/templates/jobs_usage.csv
echo 'Rows with school data:' && grep -v '^[,]*,[,]*,[,]*,[,]*,[,]*,[,]*,[,]*,[,]*,[,]*,[,]*,...'
"
```

## Next Steps

After re-collecting historical jobs with school/department:

1. Use in Billing Reports: View billing aggregated by school/department
2. Create Department Hierarchies: Build cost allocation reports
3. Export for Analysis: Use jobs\_usage.csv in Excel or Python for analysis
4. Set Departmental Quotas: Use school/department data to set resource limits

## Schedule Regular Collection

To keep jobs\_usage.csv up-to-date, run collection daily:

```
Add to crontab (runs daily at 2 AM)
0 2 * * * docker compose -f docker-compose-dev.yml exec coldfront \
python manage.py collect_sacct --current-week && \
python manage.py export_jobs_usage_csv --end $(date +%Y-%m-%d)
```

Or add to Django-Q task scheduler for automatic execution.



## Initialization Scripts

---

### test\_ldap\_provision.py

Smoke-test for **provision\_web\_user**. Verifies that a newly registered user would be (or is) correctly propagated to LDAP.

#### Usage inside the container

```
Dry-run -- reads LDAP, prints what it WOULD do, no writes:
docker exec coldfront python /opt/coldfront/init/test_ldap_provision.py

Actually create the test user (JHU tree):
docker exec coldfront python /opt/coldfront/init/test_ldap_provision.py --create

Schmidt tree with a custom username:
docker exec coldfront python /opt/coldfront/init/test_ldap_provision.py \
 --username ssci-test --affiliation schmidt --create
```

#### Dry-run output

```
[DRY-RUN] provision_web_user for 'test-ldap-user' in tree='jhu'

[DRY-RUN] Would create LDAP entries:
User DN : uid=test-ldap-user,ou=People,ou=jhu,dc=arch,dc=cluster
uidNumber: 10001
gidNumber: 10000 (users group fallback)
givenName: Test
sn : User
mail : test@example.com
Group DN : cn=test-ldap-user,ou=Groups,ou=jhu,dc=arch,dc=cluster gidNumber=10001
memberUid: add test-ldap-user -> cn=users,ou=Groups,ou=jhu,dc=arch,dc=cluster
```

With **create** it writes to LDAP and then verifies the user can be found with an **ldapsearch**.



# Slurm

Docker-based Slurm cluster for ARCH. Each subdirectory is a service image built from a shared base.

```
Bare-metal deployment (Rocky 9 + Ansible): see ANSIBLE.md.
Same files baked into the Docker images, materialized into a gitignored
`ansible/` tree via `./start.sh slurm ansible all`.
```

## Sub-Services

| Subdirectory     | Image                | Purpose                                                                                               |
|------------------|----------------------|-------------------------------------------------------------------------------------------------------|
| `base/`          | `arch/cn-rockylinux` | Base image (Rocky 9 + Slurm + Munge + SSSD + SPANK plugins)                                           |
| `slurmctld/`     | `arch/slurmctld`     | Controller daemon -- job scheduling, `job_submit.lua` enforcement                                     |
| `slurmd/`        | `arch/slurmd`        | Worker + login node -- also serves SSH when `NODE_TYPE=login`                                         |
| `slurmdbd/`      | `arch/slurmdbd`      | Accounting daemon -- stores job records in MariaDB                                                    |
| `slurmrestd/`    | `arch/slurmrestd`    | REST API -- used by bare-metal clusters (JWT + UNIX socket auth); Docker dev uses docker-exec instead |
| `munge/`         | `arch/slurm-munge`   | Key init (one-shot) -- generates munge.key, then exits                                                |
| `conf/`          | --                   | Shared config: `slurm.conf`, `cli_filter.lua`, `rates.lua`, `job_submit.lua`                          |
| `spank-plugins/` | --                   | SPANK plugin C sources (compiled into `.so` in base image)                                            |

## Build Order

All images inherit from **arch/cn-rockylinux**. Build context must be **slurm/**:

```
docker build -t arch/cn-rockylinux -f slurm/base/Dockerfile slurm/
docker build -t arch/slurm-munge -f slurm/munge/Dockerfile slurm/
docker build -t arch/slurmctld -f slurm/slurmctld/Dockerfile slurm/
docker build -t arch/slurmd -f slurm/slurmd/Dockerfile slurm/
docker build -t arch/slurmdbd -f slurm/slurmdbd/Dockerfile slurm/
docker build -t arch/slurmrestd -f slurm/slurmrestd/Dockerfile slurm/
```

Or use the Build Images button in the ColdFront cluster UI.

## Ports


| Service        | Port        | Purpose                                |
|----------------|-------------|----------------------------------------|
| slurmctld      | 6817        | Controller communication               |
| slurmd         | 6818        | Worker communication                   |
| slurmd (login) | 22          | SSH user access                        |
| slurmrestd     | 6820        | REST API (HTTP)                        |
| slurmrestd     | UNIX socket | Local auth (ColdFront via SO_PEERCRED) |

## Communication Architecture




ColdFront uses two backends depending on deployment mode:

### Docker Dev Mode (default)



Docker clusters (**use\_docker\_exec=True**) bypass slurmrestd entirely. ColdFront runs Slurm CLI commands (sacctmgr, sinfo, sacct, sreport) inside the slurmctld container via Docker SDK.

### Bare-Metal / Production Mode



Bare-metal clusters use slurmrestd REST API with UNIX socket (**SO\_PEERCRE**) or JWT auth.

### Communication Channels

| Channel        | From -> To                       | Protocol                                | Auth                                            | When Used                                                                    |
|----------------|----------------------------------|-----------------------------------------|-------------------------------------------------|------------------------------------------------------------------------------|
| Docker exec    | ColdFront -> slurmctld container | `docker exec` CLI                       | Docker socket (Peer UID)                        | Docker dev mode -- all Slurm queries (accounts, jobs, nodes, partitions)     |
| REST API       | ColdFront -> slurmrestd          | HTTP<br>(`http://slurmrestd:682`)       | `rest_auth/jwt` (TCP +<br>`X-SLURM-USER-TOKEN`) | Bare-metal / production clusters                                             |
| Slurm protocol | slurmrestd -> slurmctld/slurmdbd | Slurm RPC                               | Munge (shared key via<br>`munge-key` volume)    | Translates REST calls to native Slurm operations                             |
| Docker API     | ColdFront -> Docker daemon       | UNIX socket<br>(`/var/run/docker.sock`) | Peer UID                                        | Build images, deploy/restart/stop Slurm containers, read logs, health checks |

Key volumes shared between services:

- **munge-key** munge authentication key (all Slurm services)



- **slurm-state** slurmctld state directory (persistent across restarts)
- **slurm-db-data** MariaDB data (slurmdbd accounting)
- **slurm-archive** slurmdbd archive dumps (/var/lib/slurm/archive) written before purge; survives container recreate

## Daemons in detail

Each Slurm container runs one daemon. They communicate over the channels above; this section explains what each daemon does internally process model, state, configs, failure modes, debugging entry points.

### slurmctld -- controller daemon

The brain. One per cluster, runs in **slurmctld** container, listens on TCP 6817.

Responsibilities:

- Accept job submissions (sbatch/srun/salloc -> RPC) and place into queue
- Run the scheduler periodically: backfill + main scheduler decide which pending job gets which nodes
- Spawn **slurmstepd** on compute nodes (via slurmd) when a job starts
- Dispatch prolog/epilog hooks to **slurmd** for each job
- Push job/step/usage records to **slurmdbd** (via Slurm protocol)
- Maintain authoritative node state (**idle/alloc/mix/down/drain/fail/comp**) slurmd reports back, slurmctld decides
- Load + execute **job\_submit.lua** (server-side cost gate) on every submission
- Load + execute SPANK plugins listed in **plugstack.conf**
- Periodically write state to **/var/spool/slurm/state/** (jobs, steps, assoc cache, fed state) so the daemon can recover after restart

Key configs read:

- **/etc/slurm/slurm.conf** partition list, node list, scheduling params, plugin directives, prolog/epilog paths
- **/etc/slurm/plugstack.conf** SPANK plugins to load
- **/etc/slurm/job\_submit.lua** server-side enforcement
- **/etc/slurm/qos\_config.lua** billing toggles (loaded by job\_submit.lua)
- **/etc/slurm/gres.conf** GPU declarations per node (loaded only when nodes carry gres)

State persistence: **/var/spool/slurm/state/** (bind-mounted from **slurm-state** volume). Lost only on full **docker volume rm**. Slurm reconstructs the queue + assoc cache from this on restart; loss = lost queue + need to re-fetch assocs from slurmdbd.

Fail-over: Slurm supports primary/backup slurmctld via **SlurmctldHost=primary,backup** in slurm.conf. Not configured in CLOACK today single slurmctld per cluster. Production would set this up.

Debugging entry points:



```
docker exec dev-slurmctld scontrol show config # effective slurm.conf
docker exec dev-slurmctld scontrol show daemons # who's connected
docker exec dev-slurmctld scontrol show nodes # node state from controller's view
docker exec dev-slurmctld sdiag # scheduler stats (backfill cycle count, queue depth)
docker exec dev-slurmctld tail -f /var/log/slurm/slurmctld.log
```

**scontrol reconfigure** causes slurmctld to re-read **slurm.conf**, **plugstack.conf**, **job\_submit.lua**. Not sufficient to add/remove plugins from the plugin stack that needs a full daemon restart. See § Reconfigure / reload matrix below (commit 3).

## slurmd -- compute / login node daemon

The arms and legs. One per compute/login node. Runs in **cn-\*\*\* containers (compute)** and **login\*\*\* containers (login)**, with extra **NODE\_TYPE=login** env that enables ssh + cli\_filter, listens on TCP 6818.

Responsibilities (compute):

- Receive **launch\_tasks** RPC from slurmctld for each job step
- Set up cgroup constraints (**cgroup/v2**, per **cgroup.conf**)
- Run **prolog.sh** (which dispatches **prolog.d/\*.sh**)
- Spawn **slurmstepd** per job step (the actual user-process supervisor)
- Stream stdout/stderr back over the Slurm protocol when needed
- Run **epilog.sh** (which dispatches **epilog.d/\*.sh**) on job exit
- Report node state back to slurmctld every **SlurmdTimeout/2** seconds (default ~60s)
- Handle container-image via pyxis SPANK plugin (when present)

Responsibilities (login):

- Same as compute but does NOT register in slurm.conf **NodeName=** slurmd -D exits early on hosts not listed
- Wraps OpenSSH (**sshd**) on port 22 with PAM/SSSD for user authentication
- Optional OTP overlay (when **OTP\_MODE != disabled**) see slurmd/otp/README.md
- Mounts ARCH CLI tools (**sbalance**, **stree**, **sqos**, etc.) into the user's PATH
- Loads **CliFilterPlugins=cli\_filter/lua** so **sbatch/srun** runs the cost-preview gate
- Login hardening: single-instance wrappers, PAM limits, systemd user-slice caps, VS Code detection cron

Process tree on an active compute node (one running job):

```
slurmd (PID 1 inside container, tini-wrapped)
L-- slurmstepd: [12345.batch] (job's batch step)
 L-- /usr/bin/bash (user's submit script)
 L-- srun ...
 L-- slurmstepd: [12345.0] (srun-launched step)
 L-- <user binary>
```

Key configs read:

- **/etc/slurm/slurm.conf** partition + node membership (must match slurmctld)
- **/etc/slurm/cgroup.conf** (or **cgroup.conf.baremetal** rendered identically) resource constraints



- `/etc/slurm/gres.conf` GPU device files
- `/etc/slurm/pluginstack.conf` SPANK plugins (myspank, spank\_reemit\_cost, pyxis)
- `/etc/slurm/prolog.sh + epilog.sh` referenced via `slurm.conf` `Prolog=`, `Epilog=`

Node states (as reported via `sinfo`):

| State                                | Meaning                                                                                         |
|--------------------------------------|-------------------------------------------------------------------------------------------------|
| <code>`idle`</code>                  | Up, no jobs running, accepting new jobs                                                         |
| <code>`mix`</code>                   | Some allocated CPUs, free CPUs available for more jobs                                          |
| <code>`alloc`</code>                 | Fully allocated                                                                                 |
| <code>`comp`</code>                  | Job(s) completing -- running epilog                                                             |
| <code>`drain` / `drng`</code>        | Admin marked drain -- finish current jobs, accept no new ones                                   |
| <code>`down` / `down*`</code>        | slurmctld can't reach slurmd, OR slurmd reported an error (e.g., epilog failed -> drain reason) |
| <code>`fail` / `failing`</code>      | slurmd reported hardware failure                                                                |
| <code>`reserved`</code>              | Allocated to a Reservation                                                                      |
| <code>`idle*` etc. (* suffix)</code> | "not responding" -- slurmctld hasn't heard back in <code>`SlurmdTimeout`</code> window          |

Failure modes:

- Node "down" after container recreate slurmctld marks it down when slurmd disconnects. Auto-resumes via `scontrol update state=resume` in slurmd entrypoint background task (5s after start). If that fails, manually: `docker exec dev-slurmctld scontrol update nodename=cn-01 state=resume reason='manual'`.
- Epilog failed -> drain `clean-xdg-runtime.sh` or any drop-in returning non-zero drains the node. Memory project epilog xdg race covers this; all epilog scripts must swallow errors with `|| true; exit 0`.
- Munge auth failure slurmd can't talk to slurmctld. Symptoms: **Authentication failure** in slurmd log. See § Munge below.

Debugging entry points:

```
docker exec cn-01 scontrol show node cn-01 # local view from slurmd
docker exec cn-01 tail -f /var/log/slurm/slurmd.log
docker exec cn-01 ps auxf # process tree, slurmstepds
```

## slurmdbd -- accounting daemon

The memory. One per federation (shared across clusters) or one per cluster. Runs in **slurmdbd** container, listens on TCP 6819, backed by MariaDB.

Responsibilities:

- Persist account/association/QOS hierarchy (**sacctmgr** commands all hit here)
- Persist job/step/usage records (**sacct**, **sreport** read from here)
- Enforce associations + QOS limits when slurmctld submits a job, it asks slurmdbd "is this user+account+QOS valid? Within **GrpTRESMins**?"
- Periodic rolup: hourly/daily/monthly aggregates for fast **sreport** queries
- Archive + purge see § slurmdbd archive + purge below (already documented)

Schema overview (MariaDB tables, names prefixed with `_` per-cluster, plus shared tables):

| Table family | Purpose | Used by |
|--------------|---------|---------|
|--------------|---------|---------|



|                                               |                                                             |                                 |
|-----------------------------------------------|-------------------------------------------------------------|---------------------------------|
| `acct_table`, `assoc_table`                   | Account hierarchy + user-account associations               | sacctmgr, slurmctld at submit   |
| `qos_table`                                   | QOS definitions (Priority, GrpTRESMins, MaxWall, etc.)      | slurmctld for limit enforcement |
| `user_table`                                  | Slurm user records (NOT the LDAP users -- separate concept) | sacctmgr                        |
| `<cluster>_job_table`, `<cluster>_step_table` | Per-cluster job + step history                              | sacct                           |
| `<cluster>_assoc_usage_*_table`               | Per-cluster usage rollup (hour/day/month)                   | sreport                         |
| `<cluster>_event_table`                       | Node state change history                                   | sacct -X                        |
| `txn_table`                                   | Transaction log (who ran what `sacctmgr` command, when)     | audit                           |

How slurmctld talks to slurmdbd: native Slurm protocol over TCP 6819, munge-authenticated. slurmctld batches usage updates and pushes every **AccountingStorageEnforce**-triggered event. On startup, slurmctld pulls the assoc + qos cache from slurmdbd into RAM that's why slurmdbd MUST be reachable at slurmctld startup.

Role in federation: a single slurmdbd can back multiple slurmctld instances (one per federation cluster). The shared **acct\_table** / **assoc\_table** / **qos\_table** means a user defined once on the federation sees their associations on every member cluster. Per-cluster tables (**\_job\_table** etc.) keep job history segregated.

State persistence: MariaDB data lives in **slurm-db-data** volume. Archive dumps in **slurm-archive** (see § slurmdbd archive + purge). Lose **slurm-db-data** = lose all accounting; restore from **slurm-archive/\*.gz** via **sacctmgr archive load**.

Debugging entry points:

```
docker exec slurm-db mysql -uslurm -p"$STORAGEPASS" slurm_acct_db -e "show tables;"
docker exec slurmdbd tail -f /var/log/slurm/slurmdbd.log
docker exec dev-slurmctld sacctmgr list assoc format=Account,User,Partition,QOS # via slurmctld -> dbd
```

**sacctmgr show stats** shows slurmdbd's own RPC counters.

## slurmrestd -- REST API daemon

The front door for bare-metal. Runs in **slurmrestd** container, listens on TCP 6820 (Docker dev) or UNIX socket (production option).

Responsibilities:

- Translate HTTP requests into Slurm protocol RPCs to slurmctld + slurmdbd
- Provide OpenAPI-spec'd endpoints under **/slurm/v0.0.42/** (jobs, nodes, partitions, reservations) and **/slurmdb/v0.0.42/** (assoc, qos, users, jobs, usage)
- Validate auth on every request before forwarding

Auth:

- JWT (preferred): client sends **X-SLURM-USER-TOKEN:** header. Token signed by HS256 key at **/var/spool/slurm/jwt\_hs256.key** (must match between slurmctld + slurmdbd + slurmrestd single key, shared). Token has claim **sun=** (Slurm-user-name) and expires per **auth alt parameters max token lifespan**.
- UNIX socket peer cred (**SO\_PEERCREDS**): only when slurmrestd binds to a UNIX socket (not TCP). The connecting process's UID is read from the socket; slurmrestd assumes that user identity. Convenient on the same host but unusable across the network.
- Munge is not used for slurmrestd auth (it's daemon-to-daemon). slurmrestd **USES** munge internally to talk to



slurmctld/slurmdbd after the user is authenticated.

API version: v0.0.42 (Slurm 25.11). The Slurm REST API version is independent of the Slurm version Slurm 24.x to 25.x went from v0.0.40 to v0.0.42, breaking some endpoint paths. ColdFront's **SlurmClient** pins the version it uses; bumping requires syncing both sides.

Where ColdFront uses it: bare-metal **SlurmCluster** rows with **use\_docker\_exec=False** route all sacctmgr/sinfo/sacct equivalents through slurmrestd instead of **docker exec**. See § Slurm REST API Endpoints Used below for the specific endpoints.

Issuing a token (for manual testing or ad-hoc scripts):

```
docker exec slurmrestd scontrol token username=root lifespan=3600
Returns: SLURM_JWT=eyJhbGc...
```

Debugging entry points:

```
curl -H "X-SLURM-USER-TOKEN: $TOKEN" http://localhost:6820/slurm/v0.0.42/diag
curl -H "X-SLURM-USER-TOKEN: $TOKEN" http://localhost:6820/openapi/v3 | jq '.paths | keys'
docker exec slurmrestd tail -f /var/log/slurm/slurmrestd.log
```

OpenAPI spec: **GET /openapi/v3** returns the full machine-readable schema useful for generating clients or auditing endpoint coverage.

## munge -- daemon-to-daemon auth

The handshake. Not a Slurm daemon per se but Slurm's required auth layer. Tiny stateless service (process: **munged**); generates and verifies credentials over a UNIX socket (**/var/run/munge/munge.socket.2**).

What it does: when slurmctld wants to send an RPC to slurmd, slurmctld first calls **munge\_encode()** locally -> opaque credential string -> sends it inside the RPC. slurmd receives, calls **munge\_decode()** locally -> gets back the requester's UID/GID + the embedded payload. Both ends need the same shared key (**/etc/munge/munge.key**, 1024 random bytes).

Key distribution:

- Docker dev one-shot **arch/slurm-munge** init container generates the key into the **munge-key** shared volume. Every Slurm container mounts that volume.
- Bare-metal Ansible **roles/munge/tasks/main.yml** copies a key from **munge key src** (operator's vault) to **/etc/munge/munge.key** on every host. The same key file must be on every node before slurmd can authenticate.

Per-host requirements:

- File: **/etc/munge/munge.key**, mode **0400**, owner **munge**
- User: **munge** UID 449 (consistent across all hosts Docker base image + Ansible **slurm\_users** role both enforce this)
- Service: **munged** running and listening on the UNIX socket
- Time sync: clock skew > 60s between hosts breaks munge auth (credentials carry timestamps). Hence **chronyd** in the **common** role.



## Common failure modes:

| Symptom                                             | Cause                                                    | Fix                                                                    |
|-----------------------------------------------------|----------------------------------------------------------|------------------------------------------------------------------------|
| `Authentication failure` in<br>slurmctld/slurmd log | Different munge.key on each side                         | Re-distribute, restart `munged`                                        |
| `Credential is expired`                             | Clock skew > 60s                                         | Run `chronyc tracking` on both sides; sync NTP                         |
| `Invalid credential format`                         | UID mismatch (munge user has different UID on each host) | Ensure UID 449 everywhere                                              |
| `socket connection failed`                          | `munged` not running                                     | `systemctl status munge` / `docker exec <host> munged --foreground -v` |

## Round-trip self-test on any host:

```
munge -n | unmunge
STATUS: Success (0)
ENCODE_HOST: ...
DECODE_HOST: ...
UID: 449 (munge)
GID: 449 (munge)
```

If that fails, no Slurm daemon on this host can talk to any other.

## Authentication layers

CLOACK uses five distinct authentication mechanisms, each for a different relationship. Understanding which one is in play makes debugging "permission denied" errors much faster.

| Layer                                           | Protects                                                                   | Mechanism                                                                      | Where configured                                                                    | Touched by                                                                           |
|-------------------------------------------------|----------------------------------------------------------------------------|--------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|
| <b>**Munge**</b>                                | Slurm daemon <-> daemon (slurmctld <-> slurmd <-> slurmdbd <-> slurmrestd) | Shared HMAC key + timestamp                                                    | `/etc/munge/munge.key` on every host                                                | Every Slurm daemon; ensure key + clock + UID 449 are uniform                         |
| <b>**JWT**</b>                                  | External clients <-> slurmrestd (ColdFront, custom scripts)                | HS256 signed token, `X-SLURM-USER-TOKEN` header                                | `/var/spool/slurm/jwt_hs256.key` on slurmctld + slurmdbd + slurmrestd               | Issue with `scontrol token`; expiry per `auth_alt_parameters_max_tok`                |
| <b>**SSSD / PAM**</b>                           | SSH user identity on login nodes (POSIX uid/gid resolution, password auth) | LDAP bind + PAM stack with `pam_sss.so`                                        | ( <del>to be fixed</del> ) `etc/sss.conf` + `etc/pam.d/{sshd,password-auth}`        | Ansible role (Ansible); slurmd endpoint (Docker login branch)                        |
| <b>**OTP via Keycloak**</b>                     | Optional 2nd factor for SSH on login nodes (`OTP_MODE != disabled`)        | Custom PAM module `pam_kc_ssh.so` -> Keycloak ROPC or Device                   | `/etc/keycloak-pam.conf` + `/etc/ssh/sshd_config.d/50-otp.conf` + `/etc/pam.d/sshd` | `login_node_control/tasks/otp.yml`; see [slurmd/otp/README.md](slurmd/otp/README.md) |
| <b>**Slurm internal auth**</b> (sacctmgr admin) | Who can run sacctmgr / scontrol admin operations                           | Slurm `AdminLevel` field on user records + slurmctld's own ACL on certain RPCs | `sacctmgr list user format=user,adminlevel`                                         | Set by README; modify user <name> set adminlevel=Operator` (or `Admin`)              |

## Auth interactions worth knowing:

- A user can **ssh login01** (SSSD -> PAM authenticates) but still fail **sbatch** if their LDAP entry has no Slurm association (sacctmgr never seen them). That's the gap **coldfront sync\_slurm** closes.



- A user with valid Slurm association still can't **sbatch** to a partition where their QOS isn't allowed that's a partition-level **AllowQos/DenyQos** rule, enforced server-side regardless of how the job was submitted.
- ColdFront's REST calls to slurmrestd carry a JWT issued for the **slurm** user (**sun=slurm** claim). Inside slurmrestd, that token grants admin privileges to query any user's data which is why the **arch\_sync** plugin can list every account.
- The munge UID (449) and slurm UID (450) are intentional choices baked into ../base/base.config and the **slurm\_users** Ansible role they MUST match across all hosts in a federation, or munge auth and Slurm internal user matching break.

## slurmdbd archive + purge

slurmdbd archives every category to disk before MariaDB purges expired rows, so historical accounting can be restored long after the live DB has been trimmed.

| Category                                | Retention (PurgeXAfter) | Archive |
|-----------------------------------------|-------------------------|---------|
| Jobs                                    | 24 months               | yes     |
| Steps / Suspend / Reservations / Events | 12 months each          | yes     |
| Usage                                   | 36 months               | yes     |
| Transactions                            | not purged              | yes     |

**ArchiveDir=/var/lib/slurm/archive** in every environment (dev container + bare-metal).

- Dev (Docker): named volume **arch\_shared\_slurm\_archive** mounted into the slurmdbd container. The slurmdbd entrypoint creates + chowns the directory at startup; slurmdbd refuses to start if the dir is missing or unwritable.
- Bare-metal (Ansible): the **slurm\_config** role provisions **/var/lib/slurm/archive** (mode 0750, owner **slurm**) on the slurmdbd host and renders the same retention values via **slurmdbd.conf.j2**.

To restore historical jobs from a dump:

```
docker exec slurmdbd sacctmgr archive load file=/var/lib/slurm/archive/<dump>.archive
```

To tune retention, edit both **slurm/conf/slurmdbd.conf** (committed reference) and **slurm/conf/slurmdbd.conf.tpl** (rendered into **#{BASE\_DATA}/slurm/shared/slurmdbd.conf** by the entrypoint). For bare-metal parity, also update the embedded **slurmdbd.conf.j2** in **coldfront/custom/plugins/arch\_sync/ansible\_scaffold\_data.STATIC\_FILES** and run **python3 scripts/regen\_ansible\_scaffold.py**.

## Account hierarchy & associations

Slurm's identity model is account-based, not user-based. A user can submit jobs only if they have an association a row tying together (**user, account, partition, qos**). Three layers form the practical model:

### The hierarchy

CLOACK creates a five-level tree per cluster (configured via **cluster\_manager.compute\_slurm\_account\_for(...)**):



```

root <- cluster anchor (slurm.conf RootAccount)
L-- <affiliation> <- jhu, schmidt (per Affiliation.code)
 L-- <school> <- engineering, asas, soph, ssci (from Department.school)
 L-- <department> <- me, cs, bme, ... (Department.code)
 L-- <pi_username> <- rdesouz4, jbiondo2, ssci-rodriago
 |-- (default -- same name as PI) <- scavenger only
 L-- <pi_username>_<project_abbrev> <- per paid Project (with IO Number)
 L-- <pi_username>_<project_abbrev>_<subproject> <- optional sub-projects

```

Why a hierarchy: **sacctmgr** rollups (**sreport**) and quota inheritance work top-down. Setting **GrpTRESMins** on a school applies to every account under it. A PI's default account stays free (scavenger only) so they always have a way to run preemptible jobs; their paid project accounts get IO-Number-backed funding.

## What an "association" actually is

An entry in slurmdbd's **assoc\_table**:

```
(cluster, user, account, partition, qos_set, default_qos, ...)
```

Each user has one association per (account, partition). **AccountingStorageEnforce=associations,limits,qos** in **slurm.conf** makes slurmd refuse jobs whose (**user, account, partition**) triple has no matching row. This is the "user has no Slurm account" check that bites a user who's in LDAP but never went through **coldfront sync\_slurm**.

When a user submits **sbatch partition=a100 account=ssci-jdoe\_proj1**, slurmd checks:

1. Does an assoc exist for (user, ssci-jdoe\_proj1, a100)? If no -> reject.
2. Is the requested QOS in the assoc's **qos\_set**? If no -> reject.
3. Does the QOS's **GrpTRESMins** have remaining headroom (per usage rollup)? If no -> reject.
4. Does the partition's **AllowQos/DenyQos** admit this QOS? If no -> reject.

Only after all 4 pass does the job land in the queue.

## How ColdFront populates this

**coldfront sync\_slurm** (and the real-time **provision\_allocation\_account** signal) executes the equivalent of:

```

Hierarchy (idempotent -- sacctmgr "add" is no-op if exists)
sacctmgr -i add account jhu cluster=arch parent=root description='JHU'
sacctmgr -i add account engineering cluster=arch parent=jhu description='School of Engineering'
sacctmgr -i add account me cluster=arch parent=engineering description='Mechanical Engineering'
sacctmgr -i add account rdesouz4 cluster=arch parent=me description='PI: Ricardo Souza'
sacctmgr -i add account rdesouz4_proj1 cluster=arch parent=rdesouz4 description='Paid project'

User assoc onto the PI's default + paid accounts
sacctmgr -i add user rdesouz4 cluster=arch account=rdesouz4 qos=scavenger defaultqos=scavenger
sacctmgr -i add user rdesouz4 cluster=arch account=rdesouz4_proj1 qos=normal,premium,scavenger defaultqos=...

```

Memory **project\_slurm\_hierarchy** covers the JHU vs Schmidt branching (Schmidt PIs use a flat **ssci-\*\*\* namespace**, not school -> dept).



## System users -- never in the hierarchy

**SYSTEM\_USERNAMES** = {'admin', 'coldfront', 'root'} (in **billing\_queries.py**) are excluded from the Slurm hierarchy. **sync\_slurm** actively scrubs them from any account they might have been added to. Memory **project\_arch\_superuser**s for full rationale.

## QOS model

QOS = Quality of Service = a named bundle of limits + priority. Every job runs under exactly one QOS. CLOACK uses six QOS names with distinct roles:

| QOS                | Priority | Preemptible?                                                   | Max wall      | Billing                     | Used by                                                                                                                            |
|--------------------|----------|----------------------------------------------------------------|---------------|-----------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| `scavenger`        | 0        | <b>**Yes**</b> (preempted by `normal`/`premium`/`interactive`) | 7d            | Free (always)               | Default partition; any user's default-QOS for their PI default account                                                             |
| `normal`           | 100      | No                                                             | 7d            | Paid (per `rates.lu`)       | Standard paid QOS on `low`/`med`/`high`/`premium`/`a100`/`l40s`/`h100`/`h200`                                                      |
| `premium`          | 200      | No                                                             | 7d            | Paid (higher)               | `premium` partition + hardware-credit-holder allocations                                                                           |
| `interactive`      | 1000     | No                                                             | 4h (hard cap) | Paid                        | `interact` wrapper / Open OnDemand / tuning sessions -- high priority, short time                                                  |
| `ssci`             | 100      | No                                                             | 7d            | <b>**Free**</b> (zero-rate) | Schmidt Sciences `ssci-*` accounts -- same priority as `normal` but never charged                                                  |
| `wk_allocation_id` | inherits | No                                                             | 7d            | Paid                        | Per-allocation weekly cap. `GrpTRESMins=billing=<dollar_value*1000*60>`. One per `AllocationBilling` row that has `weekly_cap` set |

## How Slurm picks a QOS

Order of resolution at submit time:

1. **qos=** flag (highest priority) if user explicitly asks, that's what's used (assuming it's in their assoc's **qos\_set**)
2. **SLURM\_QOS** env var (second-highest)
3. The assoc's **DefaultQOS** (third)
4. Slurm's global **PriorityWeightQOS** as fallback

CLOACK sets **DefaultQOS=scavenger** on every user assoc (memory **project\_slurm\_hierarchy**). Reason: a plain **sbatch script.sh** without flags falls into scavenger (free, preemptible) rather than accidentally hitting **normal** and getting billed. To use paid: explicit **qos=normal** (or **qos=ssci** for Schmidt users) AND **partition=** AND **account=**.

## Partition x QOS interaction

**SlurmPartition.allow\_qos** (Django field) -> **AllowQos=...** in **slurm.conf**. The partition acts as a filter:

- **scavenger** partition: **AllowQos=scavenger** -> only the scavenger QOS accepted
- Paid CPU partitions (**low/med/high**): **AllowQos=normal,premium,interactive,ssci**



- **premium** partition: **AllowQos=normal,premium,ssci** (no interactive premium nodes are batch-only)

Even if a user has the right QOS in their assoc, the partition rejects if **AllowQos** doesn't list it. That's why **cli\_filter.lua** does a pre-check: it inspects both the user's QOS membership AND the partition's **AllowQos** before submitting, returning a friendly error like "Your QOS 'scavenger' is not allowed on partition 'premium' request a project with IO Number".

## Weekly cap (wk\_\*)

Special QOS class one per paid AllocationBilling row. Generated by **slurm sync. ensure weekly cap qos(allocation)** from the **weekly cap** dollar value. Format: **wk** (e.g., **wk 142**). Added to the user's assoc's **qos set** alongside **normal/premium**. When user picks **qos=wk 142**, slurmctld enforces **GrpTRESMins=billing=X** where  $X = \text{weekly\_cap} \times 1000 \times 60$ . See § Weekly Spend Cap below for the full mechanics.

## Partition x Account x QOS matrix

The concrete rules for who can submit what where. Decision tree for a typical **sbatch**:

```
Can rdesouz4 submit to partition=a100, account=rdesouz4_proj1, qos=normal?

1. Does assoc (rdesouz4, rdesouz4_proj1, a100) exist?
|-- No -> reject "Invalid account or account/partition combination"
L-- Yes -> continue

2. Is qos=normal in the assoc's qos_set?
|-- No -> reject "Invalid qos specification"
L-- Yes -> continue

3. Does partition 'a100' have AllowQos=...normal...?
|-- No -> reject "Job's QOS not permitted to use this partition"
L-- Yes -> continue

4. Is qos=normal's GrpTRESMins headroom > 0?
|-- No (used up) -> reject "Job violates accounting/QOS policy"
L-- Yes -> continue

5. Is this a paid partition AND no comment=accept_cost AND no ACCEPT_COST=1?
|-- (cli_filter.lua prompts on login; job_submit.lua rejects on controller)
L-- Accepted -> job lands in queue
```

Common matrix (CLOACK conventions):

| User type             | Default account              | Default QOS | Paid partitions reachable                                                  | Free always                                 |
|-----------------------|------------------------------|-------------|----------------------------------------------------------------------------|---------------------------------------------|
| JHU regular           | `<pi_username>`              | `scavenger` | `low`/`med`/`high`/`premium` via `<pi>_<abbrev>` accounts + `--qos=normal` | `scavenger` partition                       |
| Schmidt<br>(`ssci-*`) | `ssci-<pi>`                  | `scavenger` | Same paid partitions via `ssci-<pi>_<abbrev>` + `--qos=ssci` (zero-rate)   | `scavenger` + `ssci` QOS on paid partitions |
| External<br>(`ext-*`) | `ext-<user>`                 | `scavenger` | Restricted -- usually `low` only with `--qos=normal` (per allocation)      | `scavenger`                                 |
| Admin<br>(`cn=admin`) | All accounts (assoc on root) | `scavenger` | All partitions                                                             | All -- admin assocs bypass scoping          |

For a complete reference, query at runtime:



```
docker exec dev-slurmctld sacctmgr show assoc user=<username> format=Account,Partition,QOS,DefaultQOS,GrpTR...
```

## TRES + TRESBillingWeights

TRES = Trackable RESources. The unit Slurm uses for accounting and quota enforcement. CLOACK uses four:

| TRES       | What                                                          | Source                                      |
|------------|---------------------------------------------------------------|---------------------------------------------|
| `cpu`      | CPU cores allocated x seconds                                 | Hardware (slurm.conf NodeName CPUs=N)       |
| `mem`      | Memory allocated x seconds (MB-seconds)                       | Hardware (slurm.conf NodeName RealMemory=M) |
| `gres/gpu` | GPUs allocated x seconds                                      | `gres.conf` per cluster                     |
| `billing`  | <b>**Derived**</b> -- weighted combination of the above three | `SlurmPartition.tres_billing_weights`       |

### How billing is computed

**TRESBillingWeights="CPU=1.0,Mem=0.25G,GRES/gpu=2.0"** in `slurm.conf` tells `slurmctld`: for every CPU-second a job consumes, count 1.0 billing units; for every GB-second of memory, count 0.25; for every GPU-second, count 2.0.

A job on **`partition=a100 cpus-per-task=4 mem=16G gres=gpu:1 -t 60:00`** (1 hour) computes:

```
cpu_billing = 4 cores x 3600 s x 1.0 = 14,400
mem_billing = 16 GB x 3600 s x 0.25 = 14,400
gpu_billing = 1 GPU x 3600 s x 2.0 = 7,200

total billing = 36,000 units (= 600 billing-minutes = 10 billing-hours)
```

That number is what **`GrpTRESMins=billing=`** quotas track. The weekly cap (**`wk_*** QOS`**) enforces against this same billing value, multiplied by `dollars_to_grptres_mins()` (= dollars x 1000 x 60).

### Per-partition tuning

Each partition has its own **TRESBillingWeights**, set via **`SlurmPartition.tres_billing_weights`** in Django Admin. Drives the price differential between tiers:

- low: CPU=0.5,Mem=0.125G (half-rate)
- med: CPU=1.0,Mem=0.25G (1x)
- high: CPU=2.0,Mem=0.5G (2x)
- premium: CPU=3.0,Mem=0.75G (3x)
- a100: CPU=1.0,Mem=0.25G,GRES/gpu=2.0 (GPU bills heavily)

**`cli_filter.lua`** reads **`rates.lua`** (which mirrors these weights as \$/unit) to compute the \$ estimate shown to users.

## gres / GPU model



GRES = Generic RESource. Slurm's mechanism for GPUs (and other accelerators like RDMA fabrics). Two files declare them:

### gres.conf -- per-cluster, per-node

Generated by `cluster_manager.write_gres_conf(cluster)` from `SlurmNode.gres` rows. The `gres` field is the Slurm GRES spec `Name[:Type]:Count` (e.g. `gpu:2` typeless, or `gpu:a100:8` typed). The device `File=` differs per node depending on where it runs (`SlurmNode.deploy_type`):

```
gres.conf (real hardware -- bare_metal nodes; range File=, Count omitted so
slurmd derives it from the device range; a comment header per node)
ga1[38-43]: 8x a100
NodeName=ga1[38-43] Name=gpu Type=a100 File=/dev/nvidia[0-7]
gpu-02: 4x h100
NodeName=gpu-02 Name=gpu Type=h100 File=/dev/nvidia[0-3]
```

```
gres.conf (dev/Docker -- container nodes; fake GPUs, explicit Count + files)
NodeName=cn-01 Name=gpu Type=a100 Count=2 File=/tmp/fakegpu0,/tmp/fakegpu1
```

A `NodeName` holding a range (e.g. hostname `ga1[38-43]`) emits one line for the whole group; individual `SlurmNode` rows emit one line each. Dev fake GPUs are symlinks to `/dev/null` created by the slurmd entrypoint (`ln -sf /dev/null /tmp/fakegpu0` etc.). Slurm 25.11 requires distinct `File=` paths per GPU unit (rejects duplicates with `gpu duplicate device file name`), so the entrypoint creates `fakegpu0..fakegpu15`.

`export_cluster_ansible` regenerates `gres.conf` (alongside `slurm.conf`) on every bundle export, so a bare-metal cluster never ships a stale/missing copy; the `slurm_client` role's "Deploy gres.conf from bundle" task copies it to `/etc/slurm/gres.conf` and notifies `restart slurmd`.

### slurm.conf -- node + partition gres declarations

`NodeName` lines must declare `Gres=gpu:N` to match `gres.conf`. `PartitionName` lines include nodes with GRES.

```
NodeName=gpu-01 CPUs=64 RealMemory=512000 Gres=gpu:a100:2
NodeName=cn-01 CPUs=8 RealMemory=2048 Gres=gpu:a100:2 # dev fake
PartitionName=a100 Nodes=gpu-01,cn-01 ...
```

### User submission

```
sbatch --gres=gpu:1 # any GPU type, 1 unit
sbatch --gres=gpu:a100:2 # specifically 2 A100s
sbatch --constraint=h100 # alternative: node feature constraint
```

`gres=gpu:N` contributes to the `gres/gpu` TRES count, which `TRESBillingWeights` multiplies into `billing`, which `GrpTRESMins` quotas enforce. So GPU jobs always cost more than equivalent-wall CPU jobs.

### NVML autodetect



When **slurmd** runs on real GPU hardware with NVIDIA drivers, it auto-detects GPUs via NVML (set **slurm\_enable\_nvml: true** in the cluster's **group\_vars/all.yml** so **slurm\_build** adds **with-nvml** to the **./configure** line the flag is opt-in so minimal Rocky hosts without **cuda-nvml-devel** still build cleanly). The cluster\_manager-generated **gres.conf** is authoritative autodetect is a sanity check, but it confirms the gres.conf node matches the actual hardware.

## libnvidia-container-tools (GPU containers)

For container-image jobs on GPU hardware (pyxis + enroot), **libnvidia-container-tools** must be installed. The **enable\_gpu\_containers=true** flag in the per-cluster Ansible bundle triggers **roles/slurm\_client/tasks/gpu\_containers.yml** which installs the NVIDIA repo + toolkit. See **ANSIBLE.md** for details.

## Partitions & Allocation Tiers

Standard naming used across all clusters:

| Kind                      | Names                          | Notes                                   |
|---------------------------|--------------------------------|-----------------------------------------|
| **CPU (by billing tier)** | `low`, `med`, `high`           | TRES billing weight increases with tier |
| **GPU (by hardware)**     | `a100`, `l40s`, `h100`, `h200` | One partition per GPU model             |

QoS:

| Name                 | Purpose                                                                     |
|----------------------|-----------------------------------------------------------------------------|
| `scavenger`          | Free, preemptible. Long max wall, no cost.                                  |
| `interactive`        | Higher priority, max wall 4h, capped resources -- for tuning/debug/notebook |
| `normal` / `premium` | Standard paid QoS (backward compat)                                         |

Seeded per **SlurmCluster** via **./start.sh populate dev partitions a100,l40s,h100,h200,low,med,high,scavenger** (idempotent). QoS records are created automatically by **slurm sync** from the **SlurmQOS** model.

See § QoS model and § Partition x Account x QoS matrix above for the full enforcement logic.

## Job lifecycle end-to-end

A job goes through six phases from **sbatch** to the slurmdbd record. Understanding this end-to-end is the fastest way to debug "where did my job get stuck?" or "why did billing miss this run?".

## Sequence (happy path)



```

user@login01 slurmctld slurmd@cn-01
sbatch script.sh	
--> cli_filter.lua (login):	
* read rates.lua	
* read qos_config.lua	
* estimate cost	
* prompt [y/N] or	
check accept_cost	
submit_job RPC ----->	job_submit.lua:
(munge auth)	* re-compute cost
	* reject if !accept_cost
	* enforce QOS x partition
	* check assoc + GrpTRESMins
<-- JobID = 12345	INSERT into queue (state=PD)
	write state to /var/spool/...
	push job record to slurmdbd
	(scheduler cycle, ~30s default)
	* find candidate nodes
	* allocate cn-01
	* state PD -> R
launch_tasks RPC ----->	run prolog.sh:
(munge auth)	* prolog.d/create-xdg-runtime
	* (any other drop-ins)
	spawn slurmdstepd
	* set up cgroup (cgroup.conf)
	* exec script.sh as user
	* capture stdout/stderr
step exit ?----->	user script exits
	run epilog.sh:
	* epilog.d/clean-xdg-runtime
	* epilog.d/clean-enroot-data
<-- completion + accounting ----	state R -> CG -> CD
push final job record to dbd	

```

Total typical wall: prolog + script + epilog. Slurm bills for the entire window between launch\_tasks RPC and step exit, including prolog/epilog so a long-running epilog inflates the bill.

## Job states (sinfo/queue --states=)

| State        | Symbol | Meaning                                                                                                         |
|--------------|--------|-----------------------------------------------------------------------------------------------------------------|
| `PENDING`    | `PD`   | Queued, waiting for resources. Reason in `squeue -o %R` (e.g., `Resources`, `Priority`, `QOSGrpBillingMinutes`) |
| `RUNNING`    | `R`    | Allocated nodes + slurmdstepd spawned + user script executing                                                   |
| `COMPLETING` | `CG`   | Step exited; slurmd running epilog. Brief -- usually < 10s unless an epilog hangs                               |
| `COMPLETED`  | `CD`   | Final state, exit 0. Bills as success.                                                                          |
| `FAILED`     | `F`    | Final state, non-zero exit. Still bills for actual wall time consumed.                                          |
| `TIMEOUT`    | `TO`   | Hit `MaxWall`. Bills for wall up to the timeout.                                                                |
| `CANCELLED`  | `CA`   | User ran `scancel` or admin killed. Bills for wall up to cancel point.                                          |
| `NODE_FAIL`  | `NF`   | slurmd died mid-job. Slurm tries to requeue (`Requeue=1` per partition); otherwise bills + ends.                |
| `PREEMPTED`  | `PR`   | Job evicted by a higher-priority job (only for `scavenger` QOS). Requeued automatically.                        |

## Where state lives

| Phase            | State persisted in                                                                                                                                             |
|------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Pending in queue | slurmctld RAM + `/var/spool/slurm/state/job_state` (binary, recovered on restart)                                                                              |
| Running          | slurmctld RAM (job_state on disk) + slurmd RAM + cgroup hierarchy on the compute node                                                                          |
| Completed        | Slurm protocol -> slurmdbd -> MariaDB ` <cluster>_job_table` (queryable via `sacct`)</cluster>                                                                 |
| Old (purged)     | `slurm-archive` volume, ` <cluster&gt;.archive-2025-xx.dump` (loadable="" `sacctmgr="" archive="" load`)<="" td="" via=""></cluster&gt;.archive-2025-xx.dump`> |

## How to "trace" a stuck job



```
Where is it now?
squeue -j 12345 -o "%i %T %R %P %a %q"
JobID State Reason Partition Account QOS

Why is it pending?
squeue -j 12345 -o "%R" # Reason field
scontrol show job 12345 # full job dict -- look at "JobState=", "Reason=", "Nodes="

If running: where?
scontrol show job 12345 | grep NodeList
docker exec cn-01 ps auxf # see slurmstepd + user procs

If completed: did slurmdbd record it?
sacct -j 12345 -o JobID,State,Start,End,Elapsed,AllocTRES,Account,QOS,ReqTRES

If prolog/epilog drained the node:
scontrol show node cn-01 | grep -i reason
Look for "Reason=Prolog error" or "Reason=Epilog error" -- clear the drain after fixing:
scontrol update nodename=cn-01 state=resume reason='cleared'
```

## Where each cost-enforcement hook fires

```
sbatch -> cli_filter.lua <- LOGIN side (line 1)
?
submit_job RPC -> job_submit.lua <- CONTROLLER side (line 2)
?
queue -> scheduler -> launch_tasks
?
prolog.sh -> prolog.d/*.sh <- COMPUTE side, root context (line 3)
?
slurmstepd -> SPANK callbacks <- COMPUTE side, user context (line 4: myspank, spank_reemit_cost, pyxis)
? user script ?
slurmstepd exit
?
epilog.sh -> epilog.d/*.sh <- COMPUTE side, root context (line 5)
?
job completion record -> slurmdbd <- AUDIT (line 6: sacct sees this)
```

Each hook can reject (lines 1-2) or fail (lines 3-5, which drain the node). Audit (line 6) is read-only by the time it fires, billing is settled.

## Cost Enforcement & UX Plugins

Configuration and plugin code for cost estimation, user confirmation, and runtime cost handling in ARCH Slurm clusters.

## Components

### 1. conf/cli\_filter.lua -- Login node, UX only

#### What it does



- Loads billing rates from **conf/rates.lua** at submit time.
- Estimates job cost based on **partition**, CPUs, memory, GPUs, and wall time.
- Weekly budget: if the account has a **wk\_\*\*\* QOS cap**, shows remaining budget and warns if job cost exceeds it.
- Free partitions/accounts: shows \$0 cost estimate with resource details, then proceeds without confirmation. Applies to scavenger partition and **ssci-\*\*\* accounts**.
- Billable partitions: shows estimated cost and prompts **[y/N]** before submitting.
- QOS pre-check: verifies user's QOS is allowed on the target partition before submitting. Blocks with a clear error message (e.g., "Your QOS 'scavenger' is not allowed on partition 'premium'") and suggests creating a project with IO Number.
- Allocation-aware: skips billing display when running inside an existing allocation (**SLURM\_JOB\_ID** set), preventing duplicate output from **srun** within **salloc** (e.g., **interact** sessions).
- Non-interactive (**sbatch**): auto-accept with **export=ALL,ACCEPT\_COST=1** or **comment=accept\_cost**.
- Slurm's native **AccountingStorageEnforce=associations,limits,qos** provides server-side enforcement as a second layer.

### Who invokes it

The Slurm client binaries (**sbatch**, **srun**, **salloc**) they read **slurm.conf**, see **CliFilterPlugins=cli\_filter.lua**, then load **/etc/slurm/cli\_filter.lua** and call the Lua callbacks (**slurm cli setup defaults**, **slurm cli pre submit**) before the RPC to slurmctld. If **pre\_submit** returns an error or calls **os.exit(1)**, the submit aborts client-side.

### Where it's active

Login nodes only. Compute nodes mount **/etc/slurm/cli\_filter.lua** (via Docker volume or copied by Ansible) but their **slurm.conf** does NOT set **CliFilterPlugins=**, so the file is never loaded. Adding **CliFilterPlugins** globally to **slurm.conf** would break **sbatch** invoked from inside slurmctld (jobs submitted by the controller for prolog/epilog).

### Dependencies (loaded by dofile())

- **/etc/slurm/rates.lua** \$/core-hour table per partition, exported from ColdFront's Billing Rates page
- **/etc/slurm/qos\_config.lua** enforce billing + show billing on free part toggles per cluster, generated by **coldfront export qos config**

Both are re-read on every **sbatch/srun** invocation no daemon reload, no Slurm restart. Edit rates -> next submit picks them up.

### Deploy paths

| Env                | Source                      | Stage                                                                              | Mount/install                                                                       | Activated by                                                                                                                                                                                                                          |
|--------------------|-----------------------------|------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Dev/Docker         | `slurm/conf/cli_filter.lua` | `start.sh`<br>stage_slurm_shared_conf<br>(line 824) -><br>`\${BASE_DATA}/slurm/sh` | volume<br>`:/etc/slurm/cli_filter.lua:ro`<br>(`start.sh` lines 3922/3958/3988/4018) | login-node entrypoint appends `CliFilterPlugins=cli_filter.lua` to `<br>`/run/slurm.conf` and exports `SLURM_CONF=/run/slurm.conf`<br>via<br>`/etc/profile.d/slurm_conf.sh`<br>([slurmd/entrypoint.sh:107-127](slurmd/entrypoint.sh)) |
| Bare-metal/Ansible | `slurm/conf/cli_filter.lua` | role_slurm at<br>playbook_dir<br>}}../slurm/conf/cli_filter.lua`                   | 3922/3958/3988/4018<br>`:/etc/slurm/cli_filter.lua`<br>on login nodes only          | `lineinfile` in same task adds `CliFilterPlugins=cli_filter.lua` to<br>`/etc/slurm/slurm.conf`<br>([ansible/roles/login_node_control/tasks/lua_client.yml](../ansible/roles/login_node_control/tasks/lua_client.yml))                 |

### ColdFront controls behavior

**cli\_filter.lua** itself is static (last touched only for new features). What changes per cluster + over time are the two files it



## dofiles:

- `rates.lua` \_regenerated when staff edits `BillingRate` rows. `signals.py` watches `BillingRate` saves and triggers `export rates lua`. Output lands in `$(BASE_DATA)/slurm/shared/rates.lua` (dev) or per-cluster Ansible bundle (bare-metal).
- `qos_config.lua` \_regenerated when staff flips `SlurmCluster.enforce_billing` or `.show_billing_on_free_part`. `signals.py` watches `SlurmCluster` saves and triggers `export qos config`. Without this file, `cli filter.lua` falls back to hardcoded defaults (billing enabled, no display on free partitions).

Reload is passive: no `scontrol reconfigure`, no restart. Next `sbatch` sees the new values.

## CLI Examples

```
rdesouz4@login01:~$ srun -p scavenger --pty bash
```

```
[BILLING] Partition : scavenger
 CPU hours : 1.00 @ $0.1000/SU
 Total : ~$0.0000
[BILLING] Partition 'scavenger' -- no charge.
rdesouz4@cn-03:~$
```

```
rdesouz4@login01:~$ srun -p premium -A ssci-jdoe --pty bash
```

```
[BILLING] Partition : premium
 CPU hours : 1.00 @ $0.3500/SU
 Total : ~$0.0000
[BILLING] Account 'ssci-jdoe' -- no charge.
rdesouz4@cn-01:~$
```

```
rdesouz4@login01:~$ srun -p premium -A rdesouz4_pfp --qos=normal -t 120 --pty bash
```

```
[BILLING] Partition : premium
 CPU hours : 2.00 @ $0.3500/SU
 Total : ~$0.7000
 Weekly cap : $100.00/week
[BILLING] Proceed? [y/N]: y
[BILLING] Job 31 submitted -- estimated ~$0.7000.
rdesouz4@cn-01:~$
```

```
rdesouz4@login01:~$ srun -p a100 -A rdesouz4_proj1 --qos=normal --gres=gpu:2 --pty bash
```

```
[BILLING] Partition : a100
 CPU hours : 1.00 @ $0.3500/SU
 GPU hours : 2.00 @ $1.2000/hr
 Total : ~$2.7500
 Weekly cap : $100.00/week
[BILLING] Proceed? [y/N]: y
[BILLING] Job 41 submitted -- estimated ~$2.7500.
rdesouz4@gpu-01:~$
```



```
DefaultQOS = scavenger for every account (commit `a611402`).
A plain `srun --pty bash` defaults to partition=scavenger qos=scavenger and
just works. Paid partitions (low/med/high/premium/a100) need an explicit
`--qos=normal` (or `--qos=ssci` for `ssci-`/`ext-` users) -- the partition
rejects scavenger QOS.
```

Non-interactive sbatch:

```
sbatch --export=ALL,ACCEPT_COST=1 myjob.sh
```

## 2. conf/rates.lua -- Billing rates table

- Exported from the ColdFront Billing Rates page.
- Keyed by partition (**cpu**, **gpu**) and tier (**full\_rate**, **om\_only**).
- Loaded by **cli\_filter.lua** via **dofile("/etc/slurm/rates.lua")**.
- Update this file when rates change no image rebuild needed (bind-mounted).

## 3. conf/job\_submit.lua -- Controller, enforcement

- Recalculates and enforces cost acceptance on the slurmctld controller.
- Rejects jobs that lack explicit cost acceptance.
- Injects **SLURM\_JOB\_COST\_ESTIMATE** into the job environment.

## 4. spank-plugins/myspank.c -- SPANK plugin (sample)

- Minimal SPANK plugin that logs a message when loaded on compute nodes.
- Compiled by **base/Dockerfile** -> **/usr/lib64/slurm/myspank.so**.
- Loaded via **conf/pluginstack.conf**.

## 5. conf/spank\_reemit\_cost.c -- SPANK plugin (optional, runtime)

- Logs/processes **SLURM\_JOB\_COST\_ESTIMATE** at job start on compute nodes.
- Compiled by **base/Dockerfile** -> **/usr/lib64/slurm/spank\_reemit\_cost.so**.
- Loaded via **conf/pluginstack.conf**.
- Can be extended for telemetry or external integrations.

## How It Works

| Component           | Where?         | Purpose                                        | Enforces? |
|---------------------|----------------|------------------------------------------------|-----------|
| `cli_filter.lua`    | Login node     | Estimate & display cost, weekly budget, UX     | No        |
| `rates.lua`         | Login node     | Billing rates table (loaded by cli_filter)     | No        |
| `job_submit.lua`    | Controller     | Enforce acceptance, inject cost env var        | Yes       |
| `myspank.so`        | Compute node   | Sample SPANK plugin (logs on load)             | No        |
| `spank_reemit_cost` | Compute node   | Log/process cost at job start                  | No        |
| `spank_pyxis.so`    | All Slurm svcs | NVIDIA Pyxis -- `--container-image` via enroot | No        |



## Login Node Hardening

Login nodes (**NODE\_TYPE=login**) apply a small set of per-user controls to keep the shared host responsive:

| Mechanism                | File (source)                                                                                        | What it does                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|--------------------------|------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Single-instance wrapper  | <code>`slurm/base/login-wrapers/arch-single-instance`</code>                                         | <code>`flock`</code> lock per user; symlinks for <code>`code`</code> , <code>`code-server`</code> , <code>`jupyter`</code> , <code>`jupyter-lab`</code> , <code>`matlab`</code> , <code>`rstudio`</code> , <code>`windsurf`</code>                                                                                                                                                                                                                                                          |
| PAM limits               | <code>`slurm/base/login-limits/90-arch.conf`</code>                                                  | <code>`maxlogins=2`</code> , <code>`maxsyslogins=500`</code> , <code>`nproc`</code> soft/hard caps                                                                                                                                                                                                                                                                                                                                                                                          |
| systemd slice caps       | <code>`ansible/roles/login_node_control/templates/user-slice.conf.j2`</code>                         | <code>`CPUAccounting=yes`</code> , <code>`MemoryAccounting=yes`</code> , <code>`CPUQuota=100%`</code> (one core per user), <code>`MemoryMax=&lt;N&gt;M`</code> pulled from <code>`SlurmCluster.login_slice_mem_per_cpu`</code> in the ColdFront DB via <code>`cluster_manager.export_cluster_ansible()`</code> (dev default 2048 MB, prod 5120 MB). Bare metal only -- no-op without systemd. Editable per cluster in Django Admin -> Slurm Clusters -> "Login Node Limits".                |
| Detection cron           | <code>`slurm/base/list_vscode_users_email.sh`</code> + <code>`/etc/cron.d/arch-vscode-detect`</code> | Every 15 min: find VS Code Server / VS Code Remote Tunnels (code-tunnel / code-tunnel-server) / Windsurf / long-lived IDE processes; resolve user email via LDAP subtree (covers JHU + Schmidt OUs); warn by email when <code>`DRY_RUN=0`</code>                                                                                                                                                                                                                                            |
| OTP SSH (per login node) | <code>`slurm/slurmd/otp/{kc-ssh-auth,pam_kc_ssh.c,pam-sshd,sshd_config.d/50-otp.conf}`</code>        | Controlled by <code>`SlurmNode.otp_mode`</code> in Django admin: <code>`disabled`</code> (no OTP), <code>`mixed`</code> (ROPC terminal for ssci-*/ext-*/JHU admins via <code>`cn=admin`</code> , Device Flow browser for JHED via Entra MFA), <code>`device_only`</code> (browser only), <code>`ropc_only`</code> (terminal only). Two-prompt PAM (password + OTP code) via compiled <code>`pam_kc_ssh.so`</code> . Also disables <code>`PubkeyAuthentication`</code> on OTP-enabled nodes. |

Activation lives in **slurm/slurmd/entrypoint.sh** (login branch) and is configured via env vars **LDAP\_URI**, **LDAP\_BASE**, **VSCODE\_DETECT\_DRY\_RUN** (injected by **cluster\_manager.compose\_slurmd\_service**).

Bare-metal login nodes use the same source files via the Ansible role **ansible/roles/login\_node\_control** (playbook **ansible/login-node-control/playbook.yml**, inventory group **login\_nodes**). Flip **login\_detect\_dry\_run: 0** in **group\_vars/login\_nodes.yml** after a 1-week soak to enable outbound warning emails.

## Build

SPANK plugins are compiled automatically during the Docker image build. Build context must be **slurm/** (not **slurm/base/**) so both source paths resolve:



```
docker build \
 --file slurm/base/Dockerfile \
 -t arch/cn-rockylinux \
 slurm/
```

**cluster\_manager.py** (**build\_images()**) does this automatically.

---

## slurm.conf directives

```
AccountingStorageEnforce=associations,limits,qos
PlugStackConfig=/etc/slurm/plugstack.conf
```

- **AccountingStorageEnforce**: Users without a Slurm account association cannot submit jobs. Also enforces GrpTRESMin limits and QOS constraints.
  - Note: **ClFilterPlugins=cli\_filter/lua** is NOT in **slurm.conf**. It is injected at runtime by the login node entrypoint only (see **slurmd/entrypoint.sh**). Adding it to **slurm.conf** would break **sbatch** inside slurmd.
- 

## plugstack.conf

Plugins are loaded directly (no **plugstack.conf.d** subdirectory):

```
optional /usr/lib64/slurm/myspank.so
optional /usr/lib64/slurm/spank_reemit_cost.so
optional /usr/lib64/slurm/spank_pyxis.so
```

**optional** means the job proceeds even if the plugin fails to load.

Use **required** to abort the job on load failure.

---

## Volume mounts (per node container)

Slurm config files are staged under **\$BASE\_DATA/slurm/** (not the repo path) and bind-mounted into **/etc/slurm/** at runtime. Layout is per-cluster with a **shared/** subtree for federation-wide files:



```
$BASE_DATA/slurm/
|-- shared/ # seeded from repo slurm/conf/ on init
| |-- slurmdbd.conf
| |-- cgroup.conf
| |-- cli_filter.lua
| |-- rates.lua
| |-- plugstack.conf
| |-- prolog.d/
| |-- epilog.d/
L-- <cluster>/ # written by cluster_manager.write_slurm_conf
 |-- slurm.conf
 |-- gres.conf
```

| Host path                                 | Container path              |
|-------------------------------------------|-----------------------------|
| `\$BASE_DATA/slurm/<cluster>/slurm.conf`  | `/etc/slurm/slurm.conf`     |
| `\$BASE_DATA/slurm/<cluster>/gres.conf`   | `/etc/slurm/gres.conf`      |
| `\$BASE_DATA/slurm/shared/cgroup.conf`    | `/etc/slurm/cgroup.conf`    |
| `\$BASE_DATA/slurm/shared/cli_filter.lua` | `/etc/slurm/cli_filter.lua` |
| `\$BASE_DATA/slurm/shared/rates.lua`      | `/etc/slurm/rates.lua`      |
| `\$BASE_DATA/slurm/shared/plugstack.conf` | `/etc/slurm/plugstack.conf` |
| `\$BASE_DATA/slurm/shared/slurmdbd.conf`  | `/etc/slurm/slurmdbd.conf`  |

The repo **slurm/conf/** directory is the source of truth edit there. **./start.sh init** (via **stage slurm shared conf**) copies the shared subset into **\$BASE\_DATA/slurm/shared/**; **./start.sh slurm create** prepares the per-cluster dir and **cluster\_manager.write\_slurm\_conf(cluster)** writes **slurm.conf** / **gres.conf** straight into it. Mounts are generated by **cluster\_manager.py -> generate compose overlay()**.

Multi-cluster: each SlurmCluster (**dev, testing, arch, ...**) has its own subtree under **\$BASE\_DATA/slurm/** no file-name collisions, no shared state. Shared config lives in one place.

## CLI Utilities (installed in base image)

All ARCH CLI tools are installed at **/usr/local/bin/** and support **-M** for multi-cluster queries.

### User & PI Tools

| Command                     | Purpose                                                                                                               |
|-----------------------------|-----------------------------------------------------------------------------------------------------------------------|
| `interact -l`               | List available partitions, your accounts & QOS                                                                        |
| `interact -p <part> [opts]` | Start interactive session (validates walltime, shows weekly cap budget)                                               |
| `stree`                     | Show your own account tree (auto-detects `USER`)                                                                      |
| `stree --full`              | Your tree with user associations                                                                                      |
| `sbalance`                  | Show account core-hour utilization (current period)                                                                   |
| `sbalance --week`           | Show weekly spend cap status per account (cap vs usage in \$)                                                         |
| `sbalance --by-partition`   | Per-partition x per-user breakdown (CPU-hours + GPU-hours separately). Inherits `--start-date` / `--end-date` window. |
| `sqos all`                  | Show all QOS definitions (priority, limits)                                                                           |
| `sqos <user>`               | Show QOS for a specific user                                                                                          |
| `sqme`                      | Show your active jobs (formatted `squeue`)                                                                            |
| `seff <jobid>`              | Show CPU and memory efficiency for a completed job                                                                    |
| `sfeatures`                 | Show node features, GRES, and partition info                                                                          |



|                       |                                                     |
|-----------------------|-----------------------------------------------------|
| <code>`sbrief`</code> | Job state summary (count by Running, Pending, etc.) |
|-----------------------|-----------------------------------------------------|

## Staff / Admin Tools

| Command                                                  | Purpose                                          |
|----------------------------------------------------------|--------------------------------------------------|
| <code>`stree -r root`</code>                             | Full hierarchy from root (admin only)            |
| <code>`sassoc`</code>                                    | Show your associations (account, QOS, fairshare) |
| <code>`sassoc &lt;user&gt;`</code>                       | Show associations for a specific user            |
| <code>`sassoc --clusters`</code>                         | List all registered Slurm clusters               |
| <code>`sadmin priority &lt;jobid&gt;`</code>             | Set high priority (admin only)                   |
| <code>`sadmin addtime &lt;jobid&gt; &lt;time&gt;`</code> | Add time to running job (admin only)             |
| <code>`sadmin settime &lt;jobid&gt; &lt;time&gt;`</code> | Replace time limit (admin only)                  |

Note: Old tool names (``slurmtree``, ``slurmqs``, ``slurmassoc``) still work as symlinks.

## Account Hierarchy

The Slurm account tree mirrors the ColdFront organizational structure:

```
root
|-- jhu [Fairshare=20]
| |-- <school> (e.g. it)
| |-- <department> (e.g. arch)
| |-- <pi_username>
| |-- users
|-- schmidt [Fairshare=80]
| |-- ss (Schmidt Sciences)
| |-- ssci-<pi_username>
| |-- users
```

Department placement is driven by **AllocationBilling.department** (not **Project.department**).

Priority: non-default allocation billing dept > default allocation billing dept > Project dept > fallback.

Default departments ensure no PI is ever flat under root:

- JHU PIs without department -> **noschool/noddep**
- Schmidt PIs -> **ss** (Schmidt Sciences)
- Created automatically on startup via **seed\_initial\_data.py**

System accounts (**admin**, **coldfront**, **root**) are excluded from the Slurm hierarchy.

## Sudo (login & compute nodes)

Members of the LDAP group **cn=admin,ou=Groups,ou=jhu,dc=arch,dc=cluster** get passwordless sudo on all nodes. This is written by the entrypoint at container start (SSSD-independent):



```
/etc/sudoers.d/admin:
%admin ALL=(ALL) NOPASSWD: ALL
```

SSSD resolves the LDAP group as Unix group **admin** via **coldfront/sssd.conf**

(mounted as **/etc/sssd/sssd.conf.template** and started by the endpoint).

## Portal Setup (Before First Use)

Before PIs can create projects and submit jobs, staff must complete these one-time setup steps in ColdFront:

| Step                            | Where                           | What                                                                                 |
|---------------------------------|---------------------------------|--------------------------------------------------------------------------------------|
| 1. Import Schools & Departments | Admin Tools > Import            | CSV: `name, code, school, school_code, affiliation`                                  |
| 2. Create Slurm Cluster         | Django Admin > Slurm Clusters   | Name must match Slurm cluster (e.g. `arch`), set TRES types                          |
| 3. Create Partitions & Nodes    | Django Admin > Slurm Partitions | `scavenger`, `premium`, etc. Set AllowQos, TRESBillingWeights                        |
| 4. Configure QOS                | Django Admin > Slurm QOS        | `scavenger` (is_free=True), `normal`, `premium` (auto-created on deploy)             |
| 5. Set Billing Rates            | Django Admin > Billing Rates    | `cpu_full`, `cpu_om`, `gpu_full`, `gpu_om`, `storage`, `backup` with effective dates |
| 6. Activate Terms of Service    | Admin Tools > Import ToS        | Upload ToS CSV or create in Django Admin. Set `is_active=True`                       |
| 7. Create ColdFront Resource    | Django Admin > Resources        | Resource PIs select when requesting allocations                                      |

## ToS Enforcement

- Without an active ToS, all users can create projects freely.
- With an active ToS, users must accept all required sections before creating a project.
- Users who have not accepted the ToS will not get a Slurm account even if they can log in via Keycloak/LDAP, they cannot submit jobs (**slurm sync** skips them).

## Default vs Paid Allocations

| Type    | Project Name                  | Slurm Account       | QOS             | Billed? |
|---------|-------------------------------|---------------------|-----------------|---------|
| Default | PI username (no abbreviation) | `pi_username`       | scavenger       | No      |
| Paid    | Custom (with abbreviation)    | `pi_username_proj1` | normal, premium | Yes     |

## PI Promotion Pipeline

When a staff member promotes a user to PI via the ColdFront staff toggle (**/user/staff-manage-users//toggle-pi/**), the following steps happen automatically:

1. Project created with status **Active** no separate approval step required.
2. Allocation created with status **Active** on every allocatable Cluster resource immediately usable.
3. **run\_slurm\_sync** is queued (via django-q or inline) scoped to the new PI's username propagates the Slurm account immediately.



The Slurm account name follows these rules:

| Project abbreviation  | Slurm account name      |
|-----------------------|-------------------------|
| `default` (none set)  | `<pi_username>`         |
| custom (e.g. `proj1`) | `<pi_username>_proj1`   |
| custom + SubProject   | `<pi_username>_proj1_A` |

The same pipeline runs when using the legacy **promote.py** view (*/user/promote/*).

## Usage Gauges

Each allocation has its own Core Usage gauge. The gauge data source depends on the allocation type:

| Gauge                 | Allocation Type            | Data Source                                                            | What It Shows                            |
|-----------------------|----------------------------|------------------------------------------------------------------------|------------------------------------------|
| Core Usage (Hours)    | Default                    | `slurm2coldfront()` -> SlurmJob table                                  | Free core-hours used (scavenger QOS)     |
| Core Usage (Hours)    | Scavenger (custom project) | `slurm2coldfront()` -> SlurmJob table                                  | QOS core-hours used (normal/premium QOS) |
| Weekly Cap Usage (\$) | Paid with cap set          | `GET /api/v1/allocations/<pk>/weekly-usage/` -> slurmrestd (live AJAX) | Weekly dollar spend vs cap               |

- Default allocations (project = PI username, no abbreviation) get scavenger QOS and are never billed.
- Paid allocations (custom project with abbreviation) get normal/premium QOS and are billed.
- Each allocation has its own Slurm account (**pi\_username** vs **pi\_username\_proj1**), so jobs never mix between gauges.
- The Weekly Cap gauge only appears when **Weekly Cap (\$)** or **Weekly Cap (Hours)** is set on the allocation.

## Weekly Spend Cap

Allocations can have a Weekly Cap (\$) or Weekly Cap (Hours) attribute. Dollar cap takes precedence. The sync pipeline creates a per-allocation QOS (**wk\_{alloc\_id}**) with **GrpTRESMins=billing=** limits that work with TRESBillingWeights.

### Conversion formula

- 1 SU (Service Unit) = \$0.001
- **GrpTRESMins = \$X \* 1,000 \* 60** (dollars -> SU -> SU-minutes)
- Example: \$50 weekly cap -> **GrpTRESMins=billing=3000000**

### How it works

1. AllocationAttribute **Weekly Cap (\$)** or **Weekly Cap (Hours)** is set via ColdFront Allocation Detail.
2. **run\_slurm\_sync** reads the cap value and sets **GrpTRESMins** on the account association:
  - With cap: **GrpTRESMins=cpu={core\_hours\*60},billing={cap\_mins}** (billing differs from cpu)
  - Without cap: **GrpTRESMins=cpu={mins},billing={mins}** (billing = cpu, no separate cap)
3. The cap applies across all partitions automatically (no AllowQos changes needed).



4. When the billing-TRES cap is reached, Slurm rejects new jobs for that account.
5. Weekly cap gauge is served live by **GET /api/v1/allocations/weekly-usage/** (AJAX on allocation detail page). The **sync\_weekly\_cap\_usage** 5-min polling task was removed in Phase 2 the endpoint UPSERTs **Weekly Cap Usage (\$)** AllocationAttribute as a debounced side-effect (5 min).
6. **cli\_filter.lua** detects weekly cap (billing != cpu) and shows cap after estimated cost.
7. **sbalance week** shows weekly cap status per account (color-coded CLI report).

Note: The old **wk\_\*\*\* QOS approach was abandoned because QOS needed AllowQos on each partition and was not applied automatically as defaultqos.**

## GrpTRESMins practical examples

Given: Core Usage (Hours) = 20,000 -> **mins = 20000 x 60 = 1,200,000**

Weekly Cap (\$) = \$30 -> **cap\_mins = \$30 x 1000 x 60 = 1,800,000**

| `tres_types` (admin)   | Weekly Cap | GrpTRESMins result                             |
|------------------------|------------|------------------------------------------------|
| `cpu,billing`          | none       | `cpu=1200000,billing=1200000`                  |
| `cpu,billing`          | \$30       | `cpu=1200000,billing=1800000`                  |
| `cpu,gres/gpu,billing` | none       | `cpu=1200000,gres/gpu=1200000,billing=1200000` |
| `cpu,gres/gpu,billing` | \$30       | `cpu=1200000,gres/gpu=1200000,billing=1800000` |

- All non-billing TRES get the same value: **core\_hours x 60**
- **billing** TRES gets the weekly cap value when set; otherwise same as **cpu**
- Cap detection: **billing != cpu** means a weekly cap is active
- The cap applies across ALL partitions (no per-partition cap)

## QOS bootstrap

**slurm\_sync** automatically creates **premium** and **scavenger** QOS in Slurm if they don't exist (only **normal** exists by default). Default-project accounts get **qos=scavenger**; custom-project accounts get **qos=normal,premium**.

## Partition access control

**SlurmPartition.allow\_qos** generates the **AllowQos=...** directive in **slurm.conf**. Partitions use **AllowQOS** to restrict which accounts can submit:

| Partition   | AllowQOS                   | Effect                                       |
|-------------|----------------------------|----------------------------------------------|
| `cpu`       | `normal,premium`           | Only accounts with `normal` or `premium` QOS |
| `cpu-dev`   | `normal,premium`           | Same -- excludes scavenger-only accounts     |
| `cpu-dev2`  | `normal,premium,scavenger` | Open to all                                  |
| `scavenger` | `scavenger,normal,premium` | Open to all                                  |

Default-project accounts (QOS=**scavenger** only) are restricted to **cpu-dev2** and **scavenger** partitions.

## TRES (Trackable RESources)

TRES are the resource types that Slurm tracks, limits, and bills per job. Available TRES are configured in **slurm.conf**



via **AccountingStorageTRES**.

| TRES       | What it is                                                  |
|------------|-------------------------------------------------------------|
| `cpu`      | CPU cores x time                                            |
| `mem`      | Memory allocated                                            |
| `billing`  | Weighted cost metric (calculated from TRESBillingWeights)   |
| `node`     | Whole nodes used                                            |
| `gres/gpu` | GPU devices x time (requires `GresTypes=gpu` + `gres.conf`) |

GrpTRESMins on Slurm accounts caps the maximum each group can consume. **slurm\_sync** sets this from the allocation's Core Hours. The TRES types included in the limit are configured via **SlurmCluster.tres\_types** in the admin (default: **cpu,billing**; production with GPUs: **cpu,gres/gpu,billing**).

TRESBillingWeights on partitions defines how resources are weighted for billing. Configurable per partition via **SlurmPartition.tres\_billing\_weights** in the admin. Example:

```
CPU=1.0, Mem=0.25G, GRES/gpu=2.0
```

This means 1 GPU counts as 2 "billing units". A job using 1 GPU for 1 hour costs 2 billing-hours. The **billing** TRES in **GrpTRESMins** caps the weighted total, not just raw CPU.

To check available TRES: **sacctmgr show tres format=Type,Name -P -n**

## tres\_types vs TRESBillingWeights vs BillingRate

These three concepts work together but serve different purposes:

| Concept              | Where defined           | What it does                                            | Example                                      |
|----------------------|-------------------------|---------------------------------------------------------|----------------------------------------------|
| `tres_types`         | `SlurmCluster` admin    | Which TRES are included in `GrpTRESMins` limits         | `cpu,billing` or `cpu,gres/gpu,billing`      |
| `TRESBillingWeights` | `SlurmPartition` admin  | How resources are weighted for the Slurm `billing` TRES | `CPU=1.0,Mem=0.25G,GRES/gpu=2.0`             |
| `BillingRate`        | ColdFront Billing Rates | Price in dollars charged to PIs for invoicing           | `cpu_full: \$0.35/SU`, `gpu_full: \$2.50/hr` |

**tres\_types** controls what goes into **GrpTRESMins** (Slurm enforcement). Staff sets it once per cluster. Adding **gres/gpu** here makes Slurm track GPU limits.

**TRESBillingWeights** controls how Slurm calculates the **billing** TRES per job. Set per partition. A GPU partition with **GRES/gpu=2.0** means 1 GPU-hour = 2 billing-units. This determines how fast a job consumes the **GrpTRESMins=billing=** limit. Partition type (CPU vs GPU) is inferred from this field: if it contains **GRES/gpu**, it's a GPU partition.

**BillingRate** is purely financial the dollar amount shown on invoices and cost estimates. These are independent of Slurm weights and can change each semester without affecting Slurm limits.

They are deliberately kept separate because they change at different rates and serve different audiences (Slurm admin vs finance).

## Admin reset



Staff can reset a cap's usage counter via Django Admin -> AllocationAttribute -> select "Weekly Cap (\$)" or "Weekly Cap (Hours)" attributes -> action "Reset weekly cap QOS usage counter in Slurm". This deletes and re-creates the QOS to zero the accumulated usage. Supports both dollar and hours cap types.

## Full Detail Invoice (per-job breakdown)

The billing report table (**/billing/report/**) has three action buttons per row: Admin, PDF, and TXT. A fourth button Detail will be added to provide a full per-job breakdown while keeping the existing PDF and TXT invoices unchanged.

### What the Detail download includes

| Section               | Content                                                                                     |
|-----------------------|---------------------------------------------------------------------------------------------|
| BILLING DETAILS       | Project, PI, IO Number, Department, Billing Type                                            |
| USAGE SUMMARY         | Jobs, CPU Hours, GPU Hours, Storage (TB), Backup (TB)                                       |
| <b>**JOB DETAIL**</b> | One row per `SlurmJob` record: JobID, Partition, Start, CPUs, GPUs, CPU-Hrs, GPU-Hrs, State |
| CHARGES               | CPU, GPU, Storage, Backup charges + TOTAL                                                   |

- Output format: plain text (**.txt** download), 80-char width.
- Capped at 500 job rows; shows "... and N more jobs" if truncated.
- Storage and Backup rows appear only when > 0.

### Access control

| Role               | Can see Detail button?             |
|--------------------|------------------------------------|
| Superuser / Staff  | Yes -- all periods                 |
| PI (project owner) | Yes -- own allocation periods only |
| Regular user       | No                                 |

### Implementation plan

- New view **PIInvoicePeriodDetailView** at **/my-period/detail/** (**pi-period-detail**)
  - Queries **SlurmJob.objects.filter(account=period.account, start\_timestegte=period.start\_date, start\_timestelte=period.end\_date)**
- Detail button added in **billing\_report.html** ACTION column (between TXT and Del)
- No changes to existing PDF or TXT views

## Job Usage Metrics (SlurmJob)

Per-job tracking is done via the **SlurmJob** model, populated by **collect\_sacct**:



```
Collect last 7 days of job data from sacct
python manage.py collect_sacct --current-week

Collect current month
python manage.py collect_sacct --current-month

Import from a pipe-delimited sacct dump file
python manage.py collect_sacct --from-file coldfront/templates/jobs_usage.csv
```

**collect\_sacct** is also called automatically at the end of **run\_slurm\_sync()** (which runs daily via cron at 3:30 AM).

The Job Usage Metrics dashboard (**/billing/jobs/**) provides XDMoD-style analytics:

- Group by: User, Account, Partition, QOS, State, Department, School
- Sortable columns: CPU Hours, GPU Hours, Wall Hours, Jobs, Users, Avg CPUs, etc.
- Summary cards: Total CPU/GPU/Wall hours, job count, active users, avg wait time
- CSV export and raw job list view
- PI sees own jobs + jobs in their accounts; admin sees all

## Submit Job UI

Any authenticated ColdFront user can submit a batch job at **/billing/submit/**.

### What it does

- Paste or write a shell script (with **#SBATCH** directives)
- Choose a partition (loaded live from the active Slurm cluster)
- Choose a Slurm account (scoped to the user's allocations)
- Live cost estimate shown as you type based on **#SBATCH ntasks, cpus-per-task, time, gres=gpu**
- Scavenger / free partitions: submits directly with an "? No charge" badge no modal
- Billable partitions (e.g. **premium**): shows a cost confirmation modal (Resource / Hours / Rate / Charge table) before submitting. user must click Agree & Submit to proceed

### Account scoping

| Role              | Accounts shown                                            |
|-------------------|-----------------------------------------------------------|
| Staff / superuser | All `slurm_account_name` values from all allocations      |
| PI                | Only accounts from their own PI allocations               |
| Regular user      | Only accounts from allocations they are active members of |

### Billable partition detection

Partition billing is fully dynamic no hardcoded names. A partition is billable if it has an active **BillingRate** record in ColdFront (Admin -> Billing Rates). Adding a new billable partition via the admin automatically enables cost confirmation in the Submit Job UI.



## Multi-year rate preview (?as\_of\_date=)

**Submit Job** accepts `?as_of_date=YYYY-MM-DD` to preview rates active on any past or future date. The Options panel exposes a date picker that reloads the page with the new parameter; the cost estimate then uses **BillingRate.objects.filter(effective\_from/te=as\_of\_date, effective\_togte=as\_of\_date)**. Invalid input silently falls back to today. Submission always uses today's rates regardless of the preview date the picker is for budget planning, not back-dated billing.

## Job Submission via SlurmClient

**SlurmClient** (in `coldfront/custom/plugins/arch_sync/slurm_client.py`) supports job submission and cancellation with two backends:

| Backend     | When used                                             | Mechanism                                              |
|-------------|-------------------------------------------------------|--------------------------------------------------------|
| docker-exec | <code>`use_docker_exec=True`</code> (Docker clusters) | <code>`docker exec slurmctld sbatch --wrap ...`</code> |
| REST API    | bare-metal / UNIX socket                              | <code>POST `/slurm/{version}/job/submit`</code>        |

### Submit a job

```
from coldfront.plugins.arch_sync.slurm_client import SlurmClient

client = SlurmClient.from_cluster(cluster) # or construct directly
job_id = client.submit_job(
 script="#!/bin/bash\nsrun hostname",
 partition="cpu-dev",
 account="myaccount",
 job_name="test-job",
)
print(f"Submitted: {job_id}")
```

### Cancel a job

```
client.cancel_job(job_id)
```

## Management operations (existing)

Nodes, partitions, accounts, and job queries continue to use docker-exec or REST as before.

## PySlurm (future option)

PySlurm provides Cython bindings to **libslurm.so** for direct Python-level Slurm calls. It is not currently installed in this stack because:

- PySlurm bindings must be built against the exact same Slurm version as the running daemons.



- ColdFront runs on **python:3.10-slim** (Debian); Slurm daemons run on Rocky Linux 9.6.
- Bridging these images requires either migrating ColdFront to Rocky Linux or running PySlurm inside the **slurmctld** container.

The current REST API + docker-exec **sbatch** approach covers all job submission needs without this complexity.

If PySlurm is needed in the future:

1. Pin Slurm version in **slurm/base/Dockerfile** (e.g. **slurm-25.11.2**).
2. Switch ColdFront base image to **rockylinux:9.6**.
3. Install matching **pyslurm** in the ColdFront container:

```
pip install pyslurm==22.5.9
```

4. Add a **PySlurm** backend to **SlurmClient**.

## Pyxis / Enroot (Container Jobs)

Pyxis is a SPANK plugin that lets users run container images (via NVIDIA Enroot) as Slurm jobs. Installed in this stack, the plugin ships in every Slurm image, the runtime ships on compute and login nodes.

### What Pyxis provides

- **srun** container-image=nvcr.io/nvidia/pytorch:24.01-py3 ... syntax
- Rootless container execution via enroot (no Docker daemon required on compute nodes)
- Native GPU passthrough inside container jobs (requires **libnvidia-container-tools** installed via Ansible on bare-metal GPU nodes, omitted from the Docker dev image)
- Works with **sbatch** and interactive **srun**

### How it's wired

| Layer                                    | Image                       | Path                             | Source                                                                                                        |
|------------------------------------------|-----------------------------|----------------------------------|---------------------------------------------------------------------------------------------------------------|
| Pyxis SPANK plugin<br>(`spank_pyxis.so`) | `arch/cn-rockylinux` (base) | `usr/lib64/slurm/spank_pyxis.so` | `base/Dockerfile` builds [pyxis](https://github.com/NVIDIA/pyxis) v0.20.0 from source                         |
| Enroot runtime + helper binaries         | `arch/slurmd`               | `usr/bin/enroot*`                | `slurmd/Dockerfile` clones + builds [enroot](https://github.com/NVIDIA/enroot) v3.5.0, runs `make install &&` |
| Plug-in registration                     | shared conf                 | `etc/slurm/plugstack.conf`       | <b>optional</b> <code>pluginsdir`<br/>`usr/lib64/slurm/spank_pyxis.so`</code>                                 |

The SPANK plugin lives in the base image so every Slurm service (slurmctld / slurmd / slurmdbd) can load **plugstack.conf** without warnings. The enroot runtime is only installed on **arch/slurmd** (compute and login nodes) because only those nodes actually launch containers. **optional** (not **required**) means a missing **.so** warns but doesn't kill job submission safer for downstream clusters that disable Pyxis.

### Picking up changes

A full rebuild is required because the base + slurmd images change:



```
./start.sh slurm destroy && ./start.sh rebuild dev && ./start.sh slurm start
```

## Example job with Pyxis

```
srun \
 --partition=a100 \
 --account=myaccount \
 --container-image=nvcr.io/nvidia/pytorch:24.01-py3 \
 --container-mounts=/scratch:/scratch \
 python train.py
```

## Bumping versions

Both pins live in their respective Dockerfiles as **ARGS** override at build time:

```
docker build --build-arg PYXIS_VERSION=0.21.0 -t arch/cn-rockylinux -f slurm/base/Dockerfile slurm/
docker build --build-arg ENROOT_VERSION=3.6.0 -t arch/slurmd -f slurm/slurmd/Dockerfile slurm/
```

## References

- Pyxis: <https://github.com/NVIDIA/pyxis>
- Enroot: <https://github.com/NVIDIA/enroot>
- Enroot + Pyxis setup guide: <https://github.com/NVIDIA/pyxis/wiki/Setup>

## Slurm REST API Endpoints Used

ColdFront communicates with Slurm via **slurmrestd** (REST API). The **SlurmClient** class supports both Docker-exec and REST backends.

| Action                  | Endpoint                                 |
|-------------------------|------------------------------------------|
| List partitions (sinfo) | GET /slurm/v0.0.42/partitions            |
| List nodes              | GET /slurm/v0.0.42/nodes                 |
| List jobs (squeue)      | GET /slurm/v0.0.42/jobs                  |
| Submit job              | POST /slurm/v0.0.42/job/submit           |
| Get job detail          | GET /slurm/v0.0.42/job/{id}              |
| Cancel job              | DELETE /slurm/v0.0.42/job/{id}           |
| Create account          | POST /slurm/v0.0.42/accounts             |
| Delete account          | DELETE /slurm/v0.0.42/accounts/{name}    |
| Create reservation      | POST /slurm/v0.0.42/reservations         |
| Delete reservation      | DELETE /slurm/v0.0.42/reservation/{name} |

## Auth

Both JWT and munge supported chosen per SlurmCluster record. JWT token stored encrypted in the **SlurmCluster** model. UNIX socket auth uses **SO PEERCRED** (peer-UID based).



## Sync Pipeline

### LDAPSync (ldap\_sync.py)

Entry point: `run_ldap_sync(pi_username=None)`

| Function                                                     | What it does                                                                                                                                                                                                                                                                                                                    |
|--------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>`sync_ldap_users(</code>                               | Create/update users in ou=People (routes to jhu or schmidt tree), sync admin/staff group, email, passwords                                                                                                                                                                                                                      |
| <code>)sync_ldap_group<br/>s()``</code>                      | Sync PI groups from ColdFront DB, create PI scratch directories                                                                                                                                                                                                                                                                 |
| <code>`sync_ldap_memb<br/>ers()``</code>                     | Add/remove members from PI LDAP groups, create home dirs, manage scratch symlinks                                                                                                                                                                                                                                               |
| <code>`_flush_sssd_cac<br/>he_on_login_node<br/>s()``</code> | Post-sync hook: <code>`sss_cache -E`</code> on every Docker login-node container; <code>`ansible-playbook ansible/sssflush/playbook.yml`</code> per bare-metal cluster. Forces <code>`id`</code> / <code>`sudo`</code> to see new group membership within ~60 s instead of <code>`entry_cache_timeout`</code> (5400 s default). |

Tree routing: **ssci-\*\*\* username prefix** -> ou=schmidt, @jhu.edu **email** -> ou=jhu, **default** -> ou=jhu

Login-node cache TTL: set **entry\_cache\_timeout = entry\_cache\_user\_timeout = entry\_cache\_group\_timeout = 60** in every login node's `/etc/sss/sss.conf`. `ansible/slurm-verify/playbook.yml` emits a WARN (not failure) when any of these exceed 60 on a login node. Open shells still need re-SSH after a group change **getgroups(2)** is set at PAM login, not refreshed on demand.

### SlurmSync (slurm\_sync.py)

Entry point: `run_slurm_sync(pi_username=None)` bidirectional sync via sacctmgr/sreport CLI.

`coldfront2slurm()` PUSH:

- Ensure JHU root account (shares=20) and Schmidt root account (shares=80)
- Ensure PI account hierarchy under correct root (school -> dept -> PI -> project)
- Add active users to Slurm accounts, zero out inactive users (MaxJobs=0)
- Sync GrpTRESMins + fairshare from Core Usage Hours attribute
- Set weekly cap on billing TRES when **Weekly Cap (\$)** is set

`slurm2coldfront()` PULL:

- Read actual CPU/billing hours per quarter via sreport
- Write back to AllocationAttribute: Core Usage (Hours)

## How Syncs Are Triggered

| Event                   | What runs                                                     |
|-------------------------|---------------------------------------------------------------|
| Allocation approved     | LDAPSync (scratch dir + group) + SlurmSync (account + QOS)    |
| Member added to project | LDAPSync (home dir + group) + SlurmSync (add user to account) |
| Password changed        | LDAPSync (password hash to LDAP)                              |
| Daily cron (2 AM)       | Full LDAP sync                                                |
| Daily cron (3 AM)       | Full Slurm sync                                               |



All async tasks dispatched via django-q with **TASK\_HOOK** for admin LogEntry audit trail.

## Cluster Model

### Shared Cluster (Tier 1)

- Partitions are pre-defined by admins (not created per PI request)
- PI gets allocation approved -> Slurm account created under JHU or Schmidt root
- Account gets access to partitions based on QOS (scavenger for default, normal+premium for paid)
- Partition is chosen at job submission time (not at allocation time)

### Reserved Tier (Tier 2 -- optional, future)

- PI requests reservation -> specifies nodes + time window
- Admin approves + optional budget/payment gate
- Plugin creates Slurm reservation via REST API
- Priority QOS applied (bypasses fair-share queue)
- On expiry: nodes return to shared pool automatically

## Reconfigure / reload matrix

Slurm has three reload tiers. Picking the wrong one for a given change either does nothing or causes job-killing downtime. Use this matrix:

| Change                                                                                                                                   | Reload mechanism                                                               | Daemons affected                                 | Job impact                                                              |
|------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------|--------------------------------------------------|-------------------------------------------------------------------------|
| <b>**slurm.conf` content**</b> (most fields -- partition tweaks, scheduler params, plugin directives)                                    | <code>`scontrol reconfigure`</code>                                            | slurmctld + slurmd + slurmdbd reread the file    | None (running jobs unaffected)                                          |
| <b>**plugstack.conf**</b> add/remove an <code>`optional /.../*.so`</code> line                                                           | <code>`scontrol reconfigure`</code>                                            | slurmctld + slurmd reload the plugin stack       | None -- plugins reload between job submissions                          |
| <b>**Lua plugin file content**</b> ( <code>`cli_filter.lua`</code> , <code>`rates.lua`</code> , <code>`qos_config.lua`</code> )          | None -- re-read on every CLI invocation via <code>`dofile()`</code>            | None                                             | None                                                                    |
| <b>**Lua plugin file content**</b> ( <code>`job_submit.lua`</code> )                                                                     | <code>`scontrol reconfigure`</code> (slurmctld caches loaded Lua               | slurmctld                                        | None                                                                    |
| <b>**`gres.conf`**</b> changes                                                                                                           | <code>`scontrol reconfigure`</code>                                            | slurmd (reads at startup; reconfigure re-parses) | None                                                                    |
| <b>**`cgroup.conf`**</b> changes                                                                                                         | Daemon restart of slurmd                                                       | slurmd on affected nodes                         | New jobs use new config; running jobs keep old cgroup                   |
| <b>**slurmdbd.conf**</b> changes                                                                                                         | Restart slurmdbd                                                               | slurmdbd                                         | None for jobs; slurmctld reconnects when slurmdbd returns               |
| <b>**JobSubmitPlugins=**</b> add/remove directive in <code>slurm.conf</code>                                                             | Daemon restart of slurmctld                                                    | slurmctld                                        | Brief queue stall (~seconds); already-running jobs unaffected           |
| <b>**CliFilterPlugins=**</b> add/remove on login <code>slurm.conf</code>                                                                 | None -- <code>`sbatch`</code> re-reads <code>slurm.conf</code> each invocation | n/a                                              | None                                                                    |
| <b>**SPANK`so` file change**</b> (recompiled <code>myspank.so</code> / <code>spank_reemit_cost.so</code> / <code>spank_pyxis.so</code> ) | Daemon restart of EVERY Slurm daemon that loads plugstack                      | slurmctld + slurmd + slurmdbd                    | Brief downtime; jobs keep running but no new submissions during restart |



|                                                          |                                                 |                                                               |                                                                        |
|----------------------------------------------------------|-------------------------------------------------|---------------------------------------------------------------|------------------------------------------------------------------------|
| <b>**Slurm version bump**</b> (e.g., 25.11.2 -> 25.11.3) | Full image rebuild + container recreate         | All daemons                                                   | Drain nodes first; jobs lose unless slurmctld state survives           |
| <b>**`munge.key` rotation**</b>                          | Restart `munged` on every host (must be atomic) | Every Slurm daemon (auth breaks during <del>slurmctld</del> ) | New auth fails until all hosts have new key; sync via Ansible playbook |
| <b>**`SlurmctldHost` primary/backup change</b>           | Full slurmctld restart (both primary + backup)  | <del>slurmctld</del>                                          | Brief queue stall                                                      |

Rules of thumb:

- Lua files = **dofile()** per invocation = no reload needed for cli\_filter/rates/qos\_config; **scontrol reconfigure** for job\_submit.lua
- Anything touching the C-level plugin stack (\*\*\*.so **change**, JobSubmitPlugins= **directive add/remove**) = **daemon restart**
- Anything touching slurmdbd.conf or the MariaDB connection = slurmdbd restart
- Anything touching node hardware semantics (NodeName CPUs/Mem) or partition membership = **scontrol reconfigure** (live), but drain the node first if you're shrinking
- Munge key = atomic key rotation via Ansible (otherwise auth breaks mid-rotation)

Handlers in Ansible automate this:

- **notify: reconfigure slurm** (in roles/slurm\_config, slurm\_server, slurm\_client) fires scontrol reconfigure
- **notify: restart slurmd** (in roles/slurm\_client) systemctl restart
- **notify: restart slurmctld** (in roles/slurm\_server) systemctl restart, brief queue stall
- **notify: restart sssd** (in roles/common) systemctl restart sssd

When you edit a file in **slurm/conf/**, the corresponding Ansible task that copies it carries the right notify, so the matching reload happens automatically.

## Troubleshooting

### SSSD not resolving users after docker restart

Symptom: SSH to login node fails with **Invalid user** or **id** returns "no such user", but the user exists in LDAP.

Cause: SSSD is started as a background process by the entrypoint script, not by a service manager. When you run **docker restart login01**, the container PID 1 (entrypoint) restarts and re-launches SSSD. However, if you restart the container without compose (e.g., **docker restart** directly), the **LDAP\_ADMIN\_PASS** env var may not be set, causing the SSSD startup block to be skipped.

Fix:

```
Option 1: Restart via start.sh (recommended -- preserves env vars from compose)
./start.sh slurm restart

Option 2: Manual SSSD start inside the container
docker exec login01 sssd -D

Option 3: Recreate via compose (guaranteed clean state)
./start.sh slurm stop && ./start.sh slurm start
```



Prevention: Always use **.start.sh slurm restart** or **docker compose up** instead of **docker restart** for Slurm containers that depend on SSSD.

### slurmrestd returns 404 on API calls

Symptom: ColdFront cluster detail page shows "Cannot reach slurmrestd: 404".

Cause: The cluster's API version (**api\_version** field) does not match what slurmrestd serves. Slurm 25.11 serves **v0.0.41** and **v0.0.42**, not **v0.0.40**.

Fix: Update the cluster API version in Django Admin > Slurm Clusters > edit the cluster > set API Version to **v0.0.42**.

### slurmrestd returns 502 but data is valid

Symptom: Partition queries return HTTP 502 with error "Slurmdb query failed: Header lengths are longer than data received" but the JSON response body contains valid partition data.

Cause: Known Slurm 22.05 issue the TRES metadata query to slurmdbd fails, but partition data is still returned. The **slurm client.py** handles this gracefully by accepting valid JSON responses even on non-2xx status codes.

No action needed this is handled automatically.

### Nodes stuck in "down" state after container recreate

Symptom: **sinfo** shows nodes as **down** after recreating slurmd containers.

Cause: slurmctld marks nodes as down when slurmd disconnects during container recreation.

Fix: The slurmd entrypoint includes a background **scontrol update state=resume** that runs automatically. If nodes remain down:

```
docker exec slurmctld scontrol update nodename=cn-01 state=resume reason="container recreated"
```

### slurmctld crashes with "No Assoc usage file" on fresh cluster (Slurm 25.11)

Symptom: slurmctld exits with **fatal: No Assoc usage file (/var/spool/slurm/slurmctld/assoc\_usage)** to recover on a brand-new cluster with no previous state.

Cause: Slurm 25.11 introduced mandatory association usage state recovery. On a fresh cluster the **assoc\_usage** binary file does not exist in **StateSaveLocation**. Unlike previous versions, 25.11 treats this as a fatal error instead of creating the file on first start. An empty placeholder file also fails because slurmctld validates the binary header (requires version 10496-11264). See SchedMD Bug 20125.

Fix: The slurmctld entrypoint uses the **-i** flag ("ignore state errors") on startup. This tells slurmctld to skip incompatible/missing state files and start fresh. Once running, slurmctld creates valid state files in **StateSaveLocation** so subsequent restarts work normally without **-i**.



```
In slurmctld/entrypoint.sh:
exec slurmctld -D -v -i
```

Note: The **-i** flag discards any previous job/association usage data. This is safe for Docker dev clusters (no real usage to lose). For production bare-metal upgrades, use **scontrol reconfigure** after migrating state files from the previous version.

## AllowQos on partitions causes fatal on fresh cluster

Symptom: slurmctld exits with **fatal: Partition premium has an invalid AllowQOS (normal,premium)** on first start.

Cause: The generated **slurm.conf** includes **AllowQos=scavenger,normal,premium** on partition lines, but QOS records do not exist yet in slurmdbd (fresh cluster). slurmctld validates QOS names at startup.

Fix: The **start.sh slurm start** sequence runs **sync\_slurm** after slurmctld starts, which creates the QOS entries. On the very first boot, partitions start without **AllowQos**; the next **scontrol reconfigure** picks up the QOS constraints after sync creates them.

## Multi-Cluster Federation

CLOACK supports federated multi-cluster deployments. The federation model is what's used in production at JHU. This section describes the current state (what works today) plus the few items still pending.

### Federation -- multiple slurmctld, shared slurmdbd

```
+-----+
| slurmdbd | <- 1 unique (shared accounting DB)
| slurmrestd | <- 1 unique (REST API for all clusters)
L-----+-----+
+-----+-----+
? ? ?
+-----+ +-----+ +-----+
|slurmctld| |slurmctld| |slurmctld|
| "arch" | |"discovery"| | "dev" |
L-----+ L-----+ L-----+
? ? ?
+-----+ +-----+ +-----+
|cn-01..N| |gpu-01..N| |dev-01|
L-----+ L-----+ L-----+
? ? ?
L-----+-----+
+-----+
| login01 | <- shared login node
L-----+
```

Benefits:

- Unified accounting database billing across all clusters
- **sacct -M arch**, **sacct -M discovery** same commands, -M selects cluster
- **sbalance**, **slurmtree**, **slurm\_qos** work with **-M** flag
- 1 slurmrestd serves all clusters via **/slurmdb/v0.0.42/jobs?cluster=X**



- Shared login node users submit with **sbatch -M discovery**
- Same users/accounts across clusters (associations per cluster)

## Independent clusters -- separate slurmdbd per cluster

```
+-----+ +-----+
| slurmdbd-1 | | slurmdbd-2 |
| restd-1 | | restd-2 |
L-----+-----+
 ? ?
 arch + disc dev
```

Each slurmdbd manages N clusters. Each slurmrestd connects to its own slurmdbd. More complex but isolates failures.

## ColdFront integration

The **SlurmCluster** model already supports both architectures:

```
Federation: same slurmrestd for all
SlurmCluster(name="arch", slurmrestd_url="http://slurmrestd:6820")
SlurmCluster(name="discovery", slurmrestd_url="http://slurmrestd:6820")

Independent: separate endpoints
SlurmCluster(name="arch", slurmrestd_url="http://restd-1:6820")
SlurmCluster(name="dev", slurmrestd_url="http://restd-2:6820")
```

## Topology fields (cluster\_type + core\_services + slurmdbd\_host)

**SlurmCluster** carries two independent enums that together encode all four federation layouts:

- **cluster\_type** = **container** | **bare\_metal** (where **slurmctld** + compute nodes run)
- **core\_services** = **shared\_container** | **dedicated** (where **slurmdbd** + **slurmrestd** run)
- **slurmdbd\_host** = **host:port** (only used when **core\_services=dedicated**, for the Service Health TCP probe)

| # | cluster_type | core_services    | Scenario                                                                                                |
|---|--------------|------------------|---------------------------------------------------------------------------------------------------------|
| A | container    | shared_container | Dev docker default (current dev stack)                                                                  |
| B | bare_metal   | dedicated        | Prod bare-metal -- own slurmdbd / slurmrestd per deploy                                                 |
| C | container    | dedicated        | Docker compute federated against an external DBD                                                        |
| D | bare_metal   | shared_container | Hybrid -- bare-metal compute federated against the ColdFront deploy's shared Docker slurmdbd/slurmrestd |

**./start.sh** presets cover A (**dev\_docker**), B (**prod\_jhu/prod\_schmidt/prod\_discovery/prod\_arch**), and D (**hybrid\_bare\_shared**). Data migration **0086\_slurmcluster\_core\_services** backfills existing container records to **shared\_container** and bare-metal records to **dedicated**.

The Service Health page (**/user/service-health/**) uses these fields to decide how to verify each daemon: container clusters via **docker inspect**, bare-metal via **SlurmClient.ping()** + TCP probe of **slurmdbd\_host**, and hybrid clusters via the shared Docker containers already running in the ColdFront deploy.

Current code uses **-M cluster\_name** (auto-injected by **\_slurm\_command\_docker()**) which works with Federation out



of the box.

## Implemented today

- Resource <-> SlurmCluster link every Allocation resolves to a Resource of type Cluster, which carries an FK to SlurmCluster. Auto-created via post\_save signal on SlurmCluster.
- Per-cluster sync slurm\_sync and collect sacct iterate all active SlurmCluster rows, routing sacctmgr/sacct calls per cluster via -M.
- Multi-org on single cluster SlurmCluster.root account picks the deployment model. root or arch means both JHU + Schmidt under one root; jhu or schmidt filters allocations to that affiliation only.
- Real-time provisioning on allocation activated signal calls provision\_allocation\_account() via transaction.on\_commit(), creating the full hierarchy (root -> school -> dept -> PI -> project) on the target cluster's slurmctld via REST API immediately when an allocation is approved.
- Service Health topology probing the /user/service-health/ page reads cluster\_type + core\_services + slurmdbd\_host and probes each cluster appropriately (Docker inspect for container, SlurmClient.ping() + TCP for bare-metal).

## Pending

- End-to-end multi-cluster dev infra current dev compose runs one Docker cluster at a time. Running two simultaneously (with two slurmctld + shared slurmdbd) for federation testing is supported in code but not wired into start.sh.
- Cross-cluster job submission (**sbatch -M other\_cluster**) works at the Slurm level once federation is registered via sacctmgr add federation, not yet automated by cluster manager.

## Bare-metal Deployment (Ansible)

For bare-metal clusters the ARCH Portal emits a per-cluster scaffold at **ansible/clusters/** (inventory, **slurm.conf**, **group\_vars/all.yml**) on every Deploy/Reconfigure. A full set of Phase 6 playbooks lives next to it:

| Path                                             | Role                                                                                                                                                                                                                                                                                                   |
|--------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>`ansible/slurm-verify/playbook.yml`</code> | Read-only pre-flight checks (OS, time sync, UIDs 449/450, munge key, slurm version, SPANK support, ports).                                                                                                                                                                                             |
| <code>`ansible/slurm-server/playbook.yml`</code> | Compile Slurm 25.11.2 from source (same flags as <code>`slurm/base/Dockerfile`</code> ), install <code>slurmctld</code> / <code>slurmdbd</code> / <code>slurmrestd</code> on the hosts in the <code>`&lt;cluster&gt;_slurmctld`</code> / <code>`_slurmdbd`</code> / <code>`_slurmrestd`</code> groups. |
| <code>`ansible/slurm-client/playbook.yml`</code> | Same build stack + <code>slurmd</code> on the hosts in <code>`&lt;cluster&gt;_baremetal`</code> ; deploys <code>`cgroup.conf`</code> and optional <code>`gres.conf`</code> .                                                                                                                           |

Shared roles under **ansible/roles/**: **common**, **slurm\_users**, **slurm\_build**, **munge**, **slurm\_config** (deploys **slurm.conf** + **plugstack.conf** + compiles SPANK `.so` from **slurm/spank-plugins/** and **slurm/conf/**), **slurm\_server**, **slurm\_client**.

Typical flow:

```
ansible-galaxy install -r ansible/requirements.yml
ansible-playbook -i ansible/clusters/<name>/inventory.ini ansible/slurm-verify/playbook.yml
ansible-playbook -i ansible/clusters/<name>/inventory.ini ansible/slurm-server/playbook.yml
ansible-playbook -i ansible/clusters/<name>/inventory.ini ansible/slurm-client/playbook.yml
```

Shortcut: **./start.sh slurm prod-sim [cluster]** runs the verify playbook in **check diff** mode (read-only pre-flight) against the per-cluster inventory. Requires **ansible-playbook** on PATH.



`./start.sh slurm status` lists all registered clusters from the ColdFront DB (both `cluster_type=container` and `cluster_type=bare_metal`) with type, root account, node count, partition count, and `slurmrestd_url`. Live `sinfo/squeue` is printed only for the current container cluster; bare-metal clusters use `prod-sim` or direct SSH for live status.

Scope is Slurm software only not OS install, networking, storage, or drivers. See [ansible/README.md](#) for the full layout and reconfig matrix.

## Import partitions, QOS, and nodes from a live Slurm cluster

For bare-metal cluster registration, `_ensure_default_nodes()` does NOT auto-seed (it runs only for `cluster_type=container`, see [signals.py:679](#)). Instead, after creating the `SlurmCluster` row in Django Admin, snapshot the live cluster into ColdFront via:

```
Create-only mode (idempotent -- never overwrites existing rows)
docker exec coldfront python -m django import_slurm_cluster --cluster <name>

Drift reconciliation (overwrites Slurm-owned mutable fields)
docker exec coldfront python -m django import_slurm_cluster --cluster <name> --update

Dry-run / diff preview (no DB writes)
docker exec coldfront python -m django import_slurm_cluster --cluster <name> --update --check
```

Read-only on the Slurm side calls `client.nodes()`, `client.partitions()`, `client.list_qos()` (all `GET /slurmdb/{v}` or `GET /slurm/{v}` REST endpoints). Idempotent via `get_or_create((cluster, name))` / `get_or_create((cluster, hostname))`.

update whitelist (Slurm-owned, overwritable):

| Model                         | Updatable                                                                                                    | CF-only (never touched)                                                                                                                          |
|-------------------------------|--------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>`SlurmNode`</code>      | <code>`cpus`</code> , <code>`real_memory`</code> , <code>`gres`</code> , <code>`state`</code>                | <code>`node_type`</code> , <code>`otp_mode`</code>                                                                                               |
| <code>`SlurmPartition`</code> | <code>`max_time`</code> , <code>`default_time`</code> , <code>`allow_qos`</code> , <code>`is_default`</code> | <code>`kind`</code> , <code>`hardware`</code> , <code>`tier_code`</code> , <code>`billing_rate_key`</code> , <code>`tres_billing_weights`</code> |
| <code>`SlurmQOS`</code>       | <code>`priority`</code> , <code>`description`</code>                                                         | <code>`is_free`</code> , <code>`billing_rate_tier`</code> , <code>`ensure_in_slurm`</code>                                                       |

The same logic is exposed in Django Admin as 3 actions on `/admin/arch_sync/slurmcluster/` (Import / Update / Preview drift). Not auto-called from signals (explicit-only by design JWT bootstrap races + Admin form save shouldn't block on REST I/O).

## Phase 2 Data Models (added 2026-04-18)

Three models land with migration `arch_sync.0084_reservation_jobtemplate_deploylog`; admin is registered, CRUD UI lands in a separate pass.

| Model                             | Purpose                                                                                                                                                                                                                                                                                                                                                                                                                   |
|-----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>`ClusterReservation`</code> | PI-requested node reservation window. Status lifecycle <code>`pending` -&gt; <code>active` -&gt; <code>expired` / <code>denied` / <code>cancelled`</code>. Holds the <code>scontrol</code> inputs (nodes, partition, start/end, flags, reason, slurm_account). <code>`slurm_reservation_id`</code> stores the name returned by <code>`scontrol create</code></code></code></code></code>                                  |
| <code>`JobTemplate`</code>        | Resubmits batch script with <code>{placeholders}</code> . Per-owner uniqueness; <code>`is_public`</code> flag shares with lab members.                                                                                                                                                                                                                                                                                    |
| <code>`DeploymentLog`</code>      | Audit trail for cluster operational events: <code>`deploy`</code> , <code>`stop`</code> , <code>`restart`</code> , <code>`destroy`</code> , <code>`drain`</code> , <code>`resume`</code> , <code>`reserve`</code> , <code>`unreserve`</code> , <code>`sync`</code> , <code>`ansible`</code> , <code>`other`</code> . Written by <code>`cluster_manager.py`</code> hooks + <code>`slurm_sync._refresh_ctld_cache`</code> . |



## Defensive ctld reconfigure

**slurm\_sync.\_refresh\_ctld\_cache(cluster)** runs **scontrol reconfigure** at the end of every per-cluster sync. Added after an observed incident where scavenger submissions failed with "Invalid account or account/partition combination" because slurmctld's association cache lagged behind slurmdbd after a bulk sync. Cheap, idempotent, does not kill running jobs; errors are logged WARN, never raised.



# Slurm Restart Guide

---

When SLURM services show warning or error status in the Service Health dashboard, you can use the UI "Restart" buttons to restart them automatically, or run manual commands as shown below.

## SLURM Services Overview

- slurmdbd (SLURM Database Daemon) [Accounting and job history database](#)
- slurmctld (SLURM Controller Daemon) [Primary cluster controller](#)
- slurmrestd (SLURM REST API Daemon) [REST API for job submission and queries](#)
- cn-01, cn-02, cn-03 [Compute node daemons \(slurmd\)](#)
- login01 [Login node daemon \(slurmd\)](#)

## Restart Order (Important)

When restarting SLURM services, follow this order to avoid dependency issues:

1. slurmdbd [Start the database daemon first](#)
2. slurmctld [Then the controller daemon](#)
3. slurmrestd [Then the REST API](#)
4. Compute nodes [Finally the compute node daemons](#)
5. Login nodes [And login node daemons](#)

## Quick Restart Commands

### Using Docker Compose (Recommended)

The project uses **docker-compose-slurm-arch-dev.yml** to manage SLURM services.

Restart individual services:



```
Restart slurmdbd
docker compose -f docker-compose-slurm-arch-dev.yml restart slurmdbd

Restart slurmctld
docker compose -f docker-compose-slurm-arch-dev.yml restart slurmctld

Restart slurmrestd
docker compose -f docker-compose-slurm-arch-dev.yml restart slurmrestd

Restart a compute node
docker compose -f docker-compose-slurm-arch-dev.yml restart cn-01

Restart login node
docker compose -f docker-compose-slurm-arch-dev.yml restart login01
```

Restart all SLURM services (proper order):

```
docker compose -f docker-compose-slurm-arch-dev.yml restart slurmdbd
sleep 10
docker compose -f docker-compose-slurm-arch-dev.yml restart slurmctld
sleep 5
docker compose -f docker-compose-slurm-arch-dev.yml restart slurmrestd
docker compose -f docker-compose-slurm-arch-dev.yml restart cn-01 cn-02 cn-03 login01
```

## Using Docker Directly

Check service status:

```
docker ps | grep slurm
```

Restart a specific container:

```
Get the container ID first
docker ps -f name=slurmdbd --format "table {{.ID}}\t{{.Names}}\t{{.Status}}"

Then restart it
docker restart <container_id>
```

## UI-Based Restart

1. Go to Admin Menu -> Service Health (or **/user/service-health/**)
2. Look for SLURM services showing ?? Warn or ? Error status
3. Click the Restart button to restart the service
4. Monitor the Restart Results panel for status and logs

## Troubleshooting



## Services Won't Start

Check logs:

```
docker compose -f docker-compose-slurm-arch-dev.yml logs slurmdbd
docker compose -f docker-compose-slurm-arch-dev.yml logs slurmctld
```

Check dependencies:

- **slurmdbd** requires the MariaDB database to be healthy
- **slurmctld** requires **slurmdbd** to be healthy first
- Compute nodes require **slurmctld** to be running

## Test Service Health

```
Test slurmrestsd API connectivity
curl -X GET http://localhost:6820/slurm/v0.0.36/ping

Check slurmdbd database
docker exec slurmdbd mysqladmin -u slurm -p<password> ping

Check slurmctld status
docker exec slurmctld sinfo
```

## Health Check Indicators

The Service Health page shows:

- ? OK Service is running and healthy
- ? Warn Service is reachable but may have issues
- ? Error Service is unreachable or not responding

## Related Configuration

- Docker Compose config: **docker-compose-slurm-arch-dev.yml**
- SLURM config: **slurm/conf/slurm.conf**
- SLURM DB config: **slurm/conf/slurmdbd.conf**

For more details on SLURM administration, see the official SLURM documentation.



# OpenLDAP

Centralized POSIX-compatible identity directory for JHU and Schmidt clusters. Stores users, groups, and membership. Runs SSSD for local NSS/PAM caching.

## Build

```
docker build -t arch/openldap openldap/
```

Image: **osixia/openldap:1.5.0** (multi-platform: amd64 + arm64).

## Directory Structure

```

openldap/
|-- Dockerfile
|-- entrypoint-wrapper.sh -- Runs init scripts before osixia entrypoint
|-- healthcheck-wrapper.sh -- Runs replace_ldif.sh once when LDAP becomes healthy
|-- sssd.conf -- SSSD config (LDAP caching, NSS/PAM)
|-- ldap.conf -- OpenLDAP client defaults
|-- certs/ -- TLS certificates
|-- certs.sh -- Certificate generation
|-- local/
 |-- bin/ -- Custom scripts (LDIF generation, user/group management)

```

## LDAP Schema

- Base DN: **dc=arch,dc=cluster**
- Admin DN: **cn=admin,dc=arch,dc=cluster**
- JHU users: **ou=People,ou=jhu,dc=arch,dc=cluster**
- JHU groups: **ou=Groups,ou=jhu,dc=arch,dc=cluster**
- Schmidt users: **ou=People,ou=schmidt,dc=arch,dc=cluster**
- Schmidt groups: **ou=Groups,ou=schmidt,dc=arch,dc=cluster**
- Schema: RFC2307bis (posixAccount/posixGroup)

## Environment Variables

| Variable                           | Purpose                                    |
|------------------------------------|--------------------------------------------|
| <code>`LDAP_ADMIN_PASSWORD`</code> | Admin bind password                        |
| <code>`LDAP_ORGANISATION`</code>   | Organization name                          |
| <code>`LDAP_DOMAIN`</code>         | Domain (e.g. <code>`arch.cluster`</code> ) |

## Port



- 389 (LDAP)

## Debugging

```
Search all users
docker exec openldap ldapsearch -x -D "cn=admin,dc=arch,dc=cluster" -w $LDAP_ADMIN_PASS -b "dc=arch,dc=clus..."

Search groups
docker exec openldap ldapsearch -x -D "cn=admin,dc=arch,dc=cluster" -w $LDAP_ADMIN_PASS -b "dc=arch,dc=clus..."
```



## Keycloak

OAuth 2.0/OIDC identity broker. Manages two realms (JHU and Schmidt) with LDAP federation and optional Azure Entra ID integration. Provides SSO for ColdFront and Helpdesk.

### Build

```
docker build -t arch/keycloak keycloak/
docker build -t arch/keycloak-config -f keycloak/Dockerfile.config keycloak/
```

### Directory Structure

```
keycloak/
|-- Dockerfile -- Main Keycloak server image
|-- Dockerfile.config -- One-shot config container (creates realms, clients, providers)
|-- configure-keycloak.sh -- Auto-configures realms, LDAP providers, OIDC clients, Entra IdP
|-- healthcheck-wrapper.sh -- Healthcheck + trigger configure on first ready
|-- ldap-providers/
| |-- jhu-ldap-provider.json -- JHU LDAP federation config
| |-- schmidt-ldap-provider.json -- Schmidt LDAP federation config
|-- themes/
| |-- arch-jhu/
| | |-- login/ -- JHU login theme (custom CSS, JHU shield, ARCH footer)
| | |-- account/ -- JHU account console theme (ARCH footer, profile restrictions)
| |-- arch-schmidt/
| | |-- login/ -- Schmidt login theme
| | |-- account/ -- Schmidt account console theme
|-- certs/ -- TLS certificates (CN=localhost)
```

### Realms & LDAP Federation

| Realm     | Display Name          | LDAP Users Base        | LDAP Groups Base       | Login Theme    | Account Theme  |
|-----------|-----------------------|------------------------|------------------------|----------------|----------------|
| `jhu`     | ARCH Portal - JHU     | `ou=People,ou=jhu`     | `ou=Groups,ou=jhu`     | `arch-jhu`     | `arch-jhu`     |
| `schmidt` | ARCH Portal - Schmidt | `ou=People,ou=schmidt` | `ou=Groups,ou=schmidt` | `arch-schmidt` | `arch-schmidt` |

Both federations: **importEnabled=true**, daily full sync, connection pooling, WRITABLE mode.

### OIDC Clients

| Client ID        | Display Name   | Realms       | Purpose              |
|------------------|----------------|--------------|----------------------|
| `cloak-portal`   | ARCH Portal    | jhu, schmidt | ColdFront Portal SSO |
| `cloak-helpdesk` | ARCH Help Desk | jhu, schmidt | Helpdesk SSO         |

Clients are auto-created by **configure-keycloak.sh**. Secrets are exported to **/shared/coldfront.env** and **/shared/helpdesk.env**.



## SSO Flow

All authentication (ColdFront + Helpdesk) goes through Keycloak:

1. Unauthenticated users hitting the portal home page (/) are auto-redirected to `/oidc/jhu/authenticate/` (JHU Keycloak realm)
2. Keycloak authenticates via LDAP (username/password) or Azure Entra (JHED SSO)
3. On success, redirects back to the application with OIDC token
4. Logout from either app triggers Keycloak single sign-out (ends session for all apps)

Schmidt users access `/oidc/schmidt/authenticate/` directly. Portal <-> Helpdesk cross-links use OIDC authenticate endpoints for seamless SSO.

**SESSION\_COOKIE\_SAMESITE = 'Lax'** is required for OIDC flows **Strict** breaks session persistence on cross-site redirects from Keycloak.

## Azure Entra ID (JHED SSO)

Controlled by **JHU\_ENTRA\_ENABLED=1**. When enabled, creates an **entra-jhu** IdP in the JHU realm with claim mappers (email, first name, last name). The redirect URI for Azure App Registration:

```
https://<keycloak-hostname>:40443/realms/jhu/broker/entra-jhu/endpoint
```

Schmidt realm does not use Entra (LDAP only).

## Two-Factor Authentication (2FA)

2FA is mandatory on first login for both realms (passwordless rollout see project memory **auth-passwordless-model**). The **CONFIGURE TOTP** required action is enabled with **defaultAction=True** so every new user is prompted to enrol a TOTP credential immediately after their first successful authentication, with no way to skip the prompt mid-flow.

| Setting          | Value                                             | Set by                                                              |
|------------------|---------------------------------------------------|---------------------------------------------------------------------|
| `CONFIGURE_TOTP` | `enabled=True` (users can configure)              | `enable_passwordless_required_actions()` in `configure-keycloak.sh` |
| `defaultAction`  | **True** (forced on every new user's first login) | same                                                                |
| OTP type         | TOTP, 6 digits, 30-second period                  | realm OTP policy                                                    |

Users configure 2FA at the Keycloak Account Console (also reachable via the ColdFront banner link):

- JHU: **`https:///realms/jhu/account/account-security/signing-in`**
- Schmidt: **`https:///realms/schmidt/account/account-security/signing-in`**

Compatible apps: Google Authenticator, Authy, Microsoft Authenticator. OTP label in app shows realm display name (e.g., "ARCH Portal - JHU").

## Toggling Default Action via the Keycloak admin UI



To make TOTP optional (enrolment available in Account Console but not forced on first login), uncheck **Default Action** for **Configure OTP**:

```
Realms (top dropdown) -> jhu (or schmidt)
-> Authentication (left menu)
-> Required Actions (tab)
 -> "Default Action" column has a checkbox per row
 -> uncheck the row "Configure OTP"
```

Caveat drift gets reverted on next deploy. **configure-keycloak.sh** calls **enable\_passwordless\_required\_actions** on every **keycloak-config** start, which PUTs **defaultAction=true** back. To make the change durable, edit the function in **configure-keycloak.sh** (drop the **.defaultAction = true** from the jq payload) and re-deploy. Same caveat applies to the inverse (re-enabling via UI after editing the code to disable it).

## SSH Authentication (login nodes)

Two chained brokers Keycloak is the OIDC broker, Entra is the real IdP for JHED users:

1. SSH client -> Keycloak Device Flow (realm **jhu**, client **ssh-client-jhu**)
2. Browser opens the Keycloak URL -> Keycloak detects the federated user and redirects to Microsoft Entra (JHU tenant)
3. JHED credentials + MFA are entered in Entra Keycloak never sees the password
4. Entra returns the OIDC assertion to Keycloak (Identity Provider broker)
5. Keycloak issues its own access token for the Device Flow
6. **kc-ssh-auth** polls Keycloak, receives the token -> exit 0 -> PAM success

From SSH's perspective, the only endpoint the helper talks to is **https://keycloak:8443/realms/jhu/...** Strong authentication (JHED password + MFA) happens inside Entra.

Routing by username (mode **mixed**, slurm/slurmd/otp/kc-ssh-auth):

| User                                     | Real credential stored in     | Broker / Flow                               |
|------------------------------------------|-------------------------------|---------------------------------------------|
| `ssci-*`                                 | OpenLDAP (Schmidt federation) | Keycloak ROPC (realm `schmidt`)             |
| `ext-*`                                  | OpenLDAP (JHU federation)     | Keycloak ROPC (realm `jhu`)                 |
| `cn=admin` members                       | OpenLDAP (JHU federation)     | Keycloak ROPC (realm `jhu`)                 |
| JHED federated (e.g. `ricardo.jacomini`) | Azure Entra ID tenant JHU     | Keycloak Device Flow (realm `jhu`) -> Entra |

ROPC paths consume password + TOTP at the SSH terminal (2 PAM prompts). Device Flow shows a verification URL + user code and polls Keycloak until the user approves in a browser.

## Account Console Customization

Personal info fields are restricted:

- Username, Email read-only for users (admin-only edit). Tooltip: "To change this field, go to ARCH Portal > My Profile"
- First name, Last name editable by users



Account Console uses custom themes with:

- JHU shield logo linking to ARCH Portal
- ARCH attribution footer with Keycloak version link

## Helpdesk Role Mapping (LDAP Groups)

Helpdesk staff/admin roles are managed via LDAP groups, synced to Keycloak, and mapped to realm roles:

```
LDAP group (helpdesk-admin) -> Keycloak group -> realm role (helpdesk-admin) -> is_superuser
LDAP group (helpdesk-staff) -> Keycloak group -> realm role (helpdesk-staff) -> is_staff
```

To grant helpdesk admin access, add the user to the **helpdesk-admin** LDAP group.

## Environment Variables

| Variable                               | Purpose                                                           |
|----------------------------------------|-------------------------------------------------------------------|
| <code>`KEYCLOAK_ADMIN`</code>          | Admin username                                                    |
| <code>`KEYCLOAK_ADMIN_PASSWORD`</code> | Admin password                                                    |
| <code>`LDAP_ADMIN_PASS`</code>         | LDAP bind password                                                |
| <code>`JHU_DJANGO_CLIENT_ID`</code>    | Portal OIDC client ID (default: <code>`cloak-portal`</code> )     |
| <code>`HELPDESK_OIDC_CLIENT_ID`</code> | Helpdesk OIDC client ID (default: <code>`cloak-helpdesk`</code> ) |
| <code>`JHU_ENTRA_ENABLED`</code>       | <code>`1`</code> = enable Azure Entra ID for JHU realm            |
| <code>`JHU_ENTRA_TENANT_ID`</code>     | Azure tenant ID                                                   |
| <code>`JHU_ENTRA_CLIENT_ID`</code>     | Entra application ID                                              |
| <code>`JHU_ENTRA_CLIENT_SECRET`</code> | Entra client secret                                               |
| <code>`KEYCLOAK_JHU_REALM`</code>      | JHU realm name (default: <code>`jhu`</code> )                     |
| <code>`KEYCLOAK_SCHMIDT_REALM`</code>  | Schmidt realm name (default: <code>`schmidt`</code> )             |

## Port

- 8443 (HTTPS), mapped to 40443 on host

## Output

Auto-exports OIDC client secrets:

- `/shared/coldfront.env` ColdFront (JHU OIDC RP CLIENT SECRET, SSCI OIDC RP CLIENT SECRET)
- `/shared/helpdesk.env` Helpdesk (HELPDESK JHU OIDC CLIENT SECRET, HELPDESK SSCI OIDC CLIENT SECRET)

## Troubleshooting & Admin Recipes

The **keycloak** container has neither **curl** nor **python3**. For REST calls use a sidecar on **arch\_default**; for realm/user ops use **kcadm.sh** inside the container.

### kcadm login (reusable config file)



```
docker exec keycloak bash -c '
/opt/keycloak/bin/kcadm.sh config credentials \
 --config /tmp/kcadm.config \
 --server https://localhost:8443 \
 --realm master --user admin \
 --password "$KC_BOOTSTRAP_ADMIN_PASSWORD"
'
```

All subsequent **kcadm.sh** calls take **config /tmp/kcadm.config**. The TLS warning is expected (self-signed dev cert).

## REST sidecar (when kcadm is not enough)

```
TOKEN=$(docker run --rm --network arch_default curlimages/curl:latest sh -c \
"curl -sk -X POST https://keycloak:8443/realms/master/protocol/openid-connect/token \
-d grant_type=password -d client_id=admin-cli -d username=admin \
-d password=$(docker exec keycloak bash -c 'echo $KC_BOOTSTRAP_ADMIN_PASSWORD')" \
| sed -n 's/.*"access_token":\"([^\"]*)\".*$/\1/p')

docker run --rm --network arch_default curlimages/curl:latest \
sh -c "curl -sk -H 'Authorization: Bearer $TOKEN' \
'https://keycloak:8443/admin/realms/jhu'"

```

## Rename OTP issuer (label shown in authenticator apps)

TOTP issuer comes from the realm **displayName** not **otpPolicyIssuer** (removed from the REST schema in current Keycloak). Set by **configure\_realm\_themes()** in **configure-keycloak.sh**. To change at runtime without rebuild:

```
kcadm update realms/jhu --config /tmp/kcadm.config \
-s 'displayName=ARCH Portal - JHU'
```

Only new OTP enrollments pick up the new issuer existing entries in users' apps keep the old label until re-registered.

## Apply login/account themes to the master realm

```
kcadm update realms/master --config /tmp/kcadm.config \
-s loginTheme=arch-jhu \
-s accountTheme=arch-jhu \
-s registrationAllowed=false
```

**configure\_realm\_themes()** already loops over **master** + **jhu** + **schmidt**, so this only needs to be run manually on pre-existing deployments.

## Create a Keycloak admin user



```
Admin of the master realm (full Keycloak admin):
kcadm create users --config /tmp/kcadm.config -r master \
 -s username=rdesouz4 -s enabled=true -s email=rdesouz4@jhu.edu -s emailVerified=true
kcadm set-password --config /tmp/kcadm.config -r master \
 --username rdesouz4 --new-password 'ChangeMe!' --temporary=false
kcadm add-roles --config /tmp/kcadm.config -r master \
 --username rdesouz4 --rolename admin

Realm-scoped admin (only the jhu realm):
kcadm add-roles --config /tmp/kcadm.config -r jhu \
 --username rdesouz4 --cclientid realm-management --rolename realm-admin
```

## Clear a stale OTP credential

User deleted the entry from their authenticator app, but Keycloak still has the credential they can't re-enroll without removing it.

```
USER=rdesouz4; REALM=jhu
UID=$(kcadm get users --config /tmp/kcadm.config -r $REALM \
 -q username=$USER --fields id --format csv --noquotes | tail -1)

List credentials (find the id where type=otp):
kcadm get users/$UID/credentials --config /tmp/kcadm.config -r $REALM \
 --fields id,type,userLabel

Delete the OTP credential:
kcadm delete users/$UID/credentials/<CRED_ID> --config /tmp/kcadm.config -r $REALM
```

Alternative force re-enrollment without deleting (user is prompted to configure TOTP again on next login, old entry is replaced):

```
kcadm update users/$UID --config /tmp/kcadm.config -r jhu \
 -s 'requiredActions=["CONFIGURE_TOTP"]'
```

User self-service: Account Console -> <https://realms/jhu/account> -> Signing in -> Two-factor authentication -> Remove.

## "Already completed" marker -- re-run configuration

**configure-keycloak.sh** skips itself when **/opt/keycloak/data/.configured** exists. To force a re-run (e.g., after editing the script):

```
docker exec keycloak rm -f /opt/keycloak/data/.configured
docker cp configure-keycloak.sh keycloak:/keycloak/configure-keycloak.sh
docker exec keycloak bash /keycloak/configure-keycloak.sh
```

## Inspect realm settings (theme, OTP, registration)



```
for r in master jhu schmidt; do
 echo "=== $r ==="
 kcadm get realms/$r --config /tmp/kcadm.config \
 --fields displayName,loginTheme,accountTheme,registrationAllowed
done
```



# PostgreSQL

Primary data persistence layer. Initializes and maintains three databases for the CLOACK stack.

## Build

```
docker build -t arch/postgres postgres/
```

Image: **postgres:16**.

## Databases

| Database     | Owner       | Purpose                                                     |
|--------------|-------------|-------------------------------------------------------------|
| `keycloak`   | `keycloak`  | Keycloak realms, clients, providers                         |
| `coldfront`  | `coldfront` | ColdFront Django application data                           |
| `slurm_jobs` | `coldfront` | High-volume Slurm job metrics (separate DB for performance) |

Created idempotently by **init-databases.sh** via Docker's **/docker-entrypoint-initdb.d/** mechanism.

## Directory Structure

```
postgres/
|-- Dockerfile
L-- init-databases.sh -- Idempotent SQL initialization (3 DBs + 2 users)
```

## Environment Variables

| Variable                | Purpose                          |
|-------------------------|----------------------------------|
| `POSTGRES_USER`         | Admin user (default: `postgres`) |
| `POSTGRES_PASSWORD`     | Admin password                   |
| `KEYCLOAK_DB_NAME`      | Keycloak DB name                 |
| `KEYCLOAK_DB_USER`      | Keycloak DB user                 |
| `KEYCLOAK_DB_PASSWORD`  | Keycloak DB password             |
| `COLDFRONT_DB_NAME`     | ColdFront DB name                |
| `COLDFRONT_DB_USER`     | ColdFront DB user                |
| `COLDFRONT_DB_PASSWORD` | ColdFront DB password            |

## Port

- 5432 (PostgreSQL)



## Healthcheck

**pg\_isready** (10-second interval, 10 retries).



## Base Configuration

---

Centralized configuration shared across all services.

### Files

| File          | Purpose                                                         |
|---------------|-----------------------------------------------------------------|
| `base.config` | Static defaults: UID/GID ranges, paths, hostnames, domain names |
| `bin/`        | Shared utility scripts                                          |


### base.config

Sourced by multiple services (postgres, keycloak, openldap) for consistent UID/GID ranges and default credentials. Key values:

- SLURM\_UID / MUNGE\_UID System account UIDs (must match across all Slurm nodes)
- LDAP\_DOMAIN, LDAP\_BASE\_DN LDAP directory settings
- Default passwords Used only for initial dev bootstrap



# Architecture Overview





max\_wall=4h, capped) · normal/premium (legacy paid).



#### SLURM ACCOUNT HIERARCHY

```
root
├─ jhu [20] (organization)
│ └─ it (school)
│ └─ arch (department)
│ ├── rdesouz4 ← default (scavenger)
│ └─ rdesouz4_proj1 ← paid (premium)
└─ schmidt [80] (organization)
 └─ ssci (school)
 └─ ss (department)
 └─ ssci-gomes ← always free (ssci-*)
```



#### BILLING, STORAGE & FEATURES

##### Billing

Rates · Invoices ·  
Credits

##### Storage & Backup

Per-project TB  
allocation

##### Job Metrics

XDMoD-style  
dashboard

##### Terms of Service

Required for Slurm  
access

##### Weekly Spend Cap

GrpTRESMins enforcement



#### SCHEDULED TASKS (DJANGO-Q)

##### sync\_ldap

Daily 2 AM

##### sync\_slurm

Daily 3 AM

##### collect\_sacct

Every 15 min

##### weekly\_cap

Every 5 min

##### core\_reset

Configurable

##### cleanup

Hourly :30

Click any block to see details